

# Data Management in Machine Learning: Challenges, Techniques, and Systems

**Arun Kumar**

UC San Diego  
La Jolla, CA, USA

**Matthias Boehm**

IBM Research – Almaden  
San Jose, CA, USA

**Jun Yang**

Duke University  
Durham, NC, USA

**SIGMOD 2017**

# Who We Are



**Arun Kumar**

UC San Diego  
La Jolla, CA, USA



Bismarck



Columbus



Orion

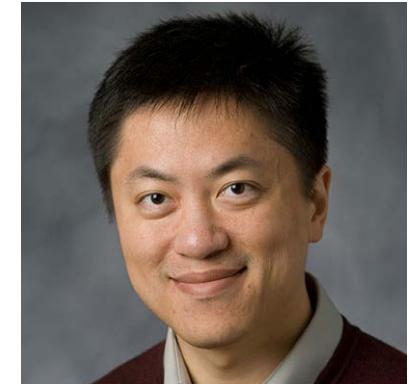


Hamlet



**Matthias Boehm**

IBM Research – Almaden  
San Jose, CA, USA



**Jun Yang**

Duke University  
Durham, NC, USA



Apache  
SystemML™



# Motivation: A Data-Centric View of ML

## ■ Application Perspective

- Machine learning / advanced analytics / deep analytics
- ➔ **Modern data-driven applications** (e.g., BI, e-commerce, healthcare)

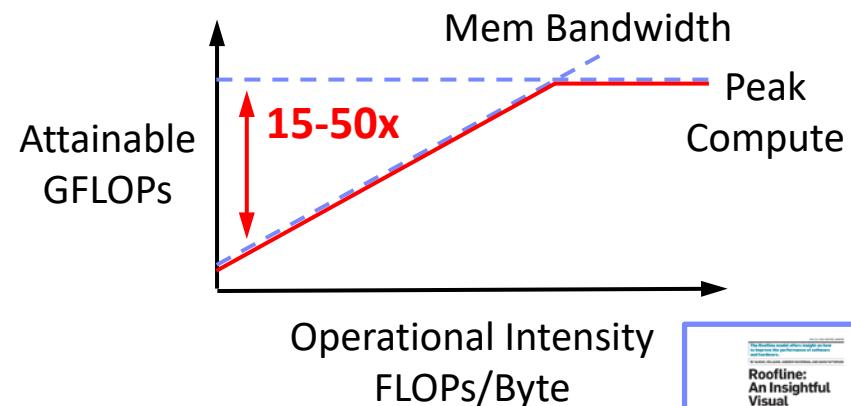
## ■ Workload Perspective

- Repetitive ML workflows
- Often **iterative ML algorithms**
- Often **I/O-bound operations**  
(e.g., matrix-vector multiplications)

## ■ Systems Perspective

- **ML in data systems**
- **DB-inspired ML systems**
- **ML Lifecycle Systems**

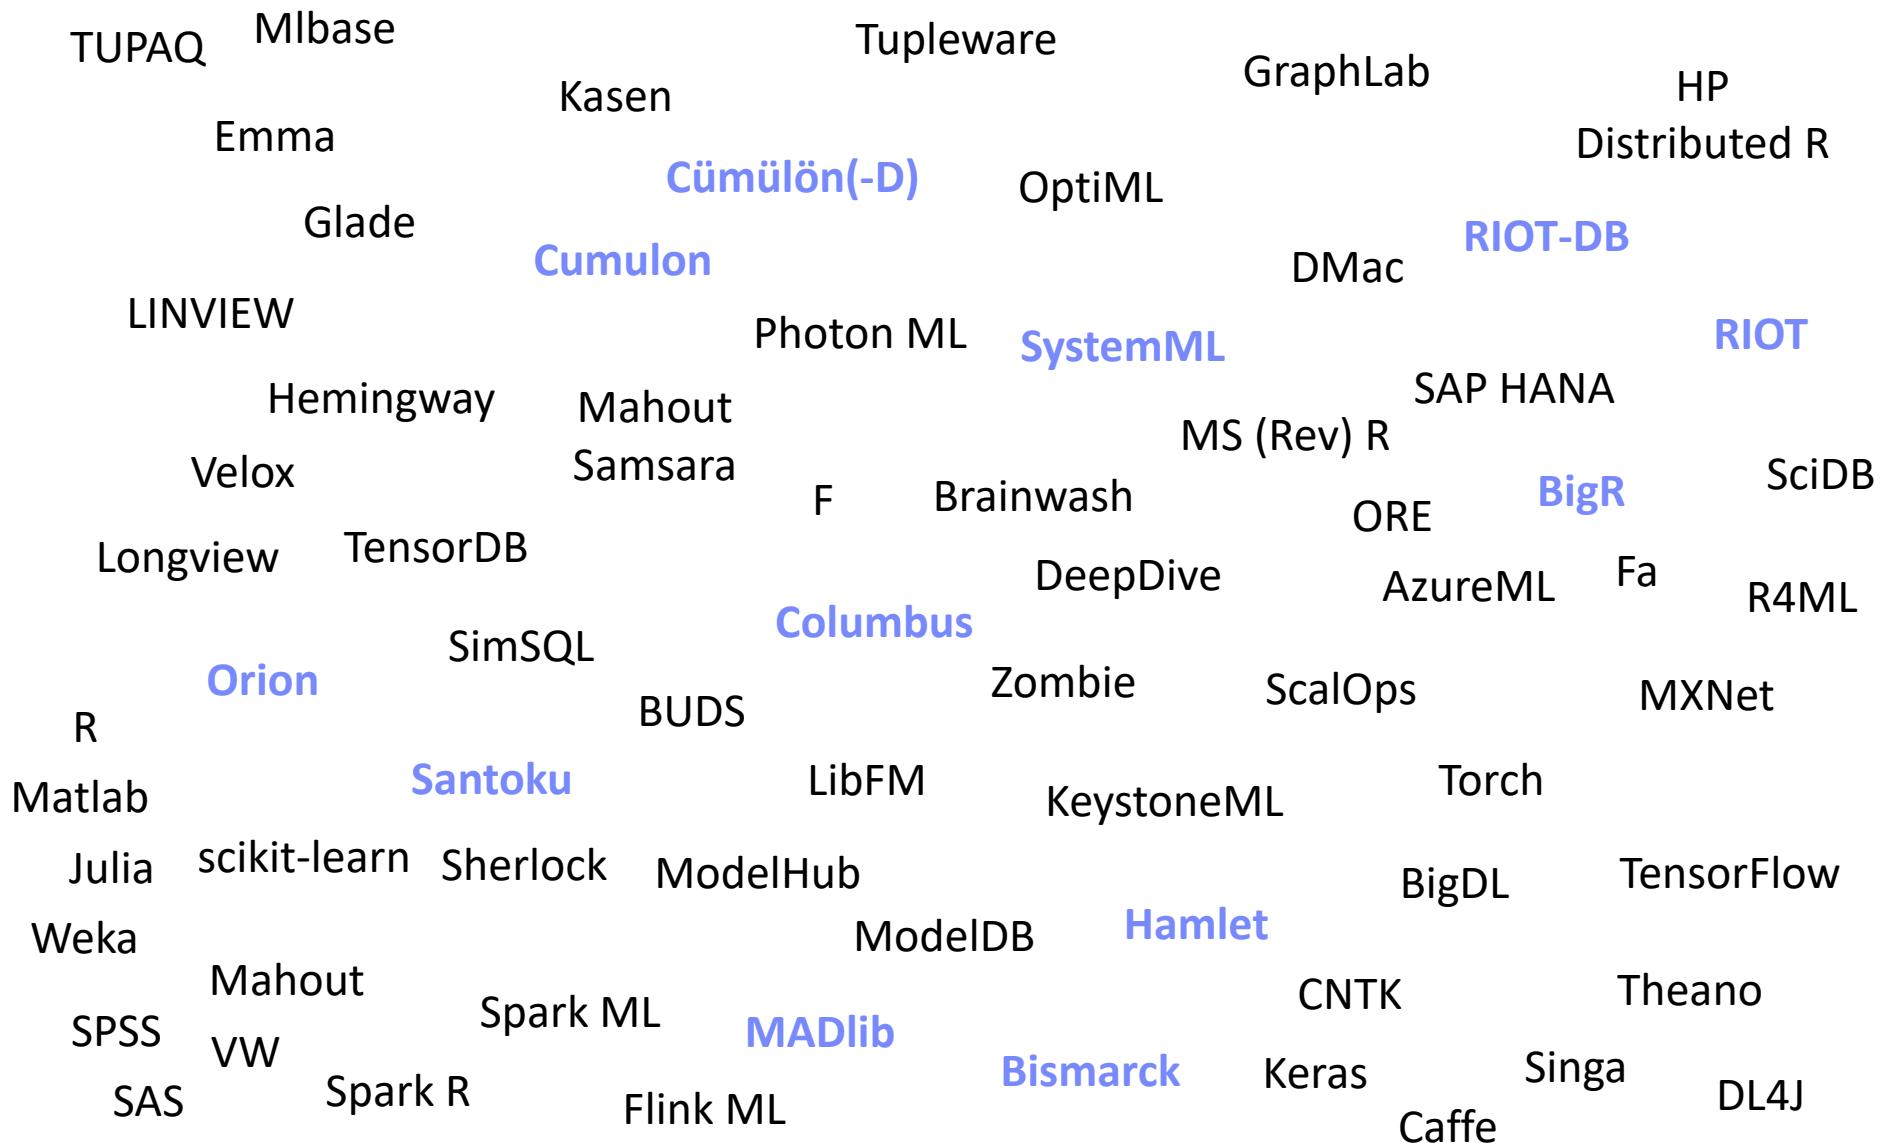
This  
Tutorial



CACM'09



# Motivation: Systems Landscape



# Motivation: Tutorial Goals

- **Overall Goal:** Comprehensive review of systems and techniques that tackle data management challenges in the context of ML workloads
- **#1 Categorize Existing Systems**
  - ML in data systems, DB-inspired ML systems, ML lifecycle systems
- **#2 Survey State-of-the-Art Techniques**
  - Query gen, UDFs, factorized learning, deep DBMS integration
  - Optimization and runtime techniques, incl. resource elasticity
  - Model selection and model management

## → Intended Takeaways

- Awareness of existing systems and techniques
- Survey of effective optimization and runtime techniques
- Overview of open research problems

# What this Tutorial is NOT

- **Introduction to Machine Learning**



[SIGMOD'13]

[SIGMOD'16]

- **Tutorial on General-Purpose Systems**

- Dataflow systems
- Graph-focused systems

[SIGMOD  
Record'16]

- **Tutorial on Deep Learning**

- Deep learning algorithms
- Deep learning systems (e.g., Torch, Theano, BigDL, TensorFlow, MXNet, CNTK, Singa, Keras, Caffe, DL4J)



[CIDR'17]

- **Tutorial on ML for RDBMS Internals**

- Cost models
- Workload prediction (e.g., in Peloton)



# Tutorial Outline

## ML in Data Systems

- **2 Query Generators and UDFs** 14min JY
- **3 Factorized Learning and Deep RDBMS Integration** 8min AK

## DB-Inspired ML Systems

- **4 Rewrites, Operator Selection, and Fusion** 14min MB
- **5 Compression, Scan Sharing, and Index Structures** 10min MB
- [▪ **6 Cloud Resource Elasticity** 10min JY ]

## ML Lifecycle Systems

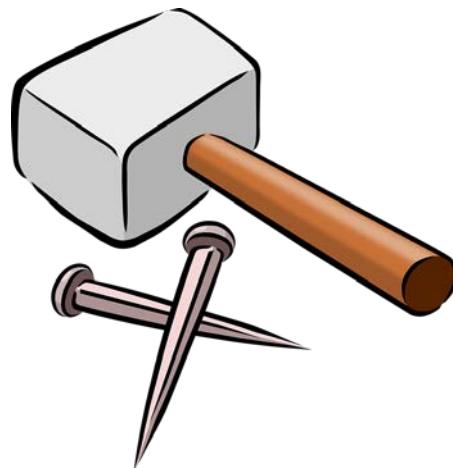
- **7 Feature Engineering, Model Selection/Management** 16min AK

## Open Problems and Q&A

# Part 2: ML with SQL & UDF

*“I suppose it is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail.”*

*Abraham Maslow, 1966*



**Jun Yang**

Duke University  
Durham, NC, USA

**SIGMOD 2017**

# ML in Database – Why?

- **Convenience**

- “Elephants” (octopi?) have shown remarkable flexibility
  - A single platform for not only data management, transformation, and querying, but also ML and application of insights

- **Efficiency**

- Move the analysis, not data
  - Can co-optimize various steps involved in the “big data pipeline”

- **Declarativeness**

- Simplifies development
  - Enables effective automatic optimization, which helps scalability/efficiency
  - One area where the DB community has plenty to offer

# Roadmap

- **First, examples of what SQL can do for ML, at various levels of abstraction:**
  - Matrix multiply
  - Ordinary least squares
  - Gradient descent
  - ( See backup slides for
    - $k$ -means
    - Markov-chain Monte-Carlo
- **Then, a brief discussion of approaches to using SQL for ML**

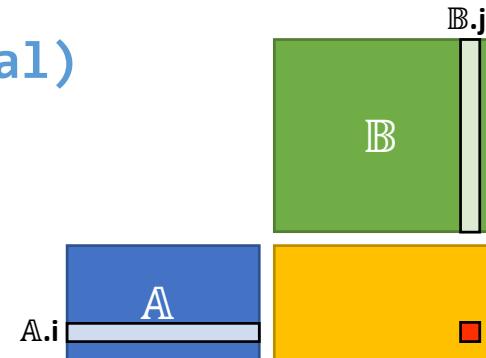
# Matrix Multiply: Take 1

- Data:  $A(i, j, val), B(i, j, val)$

– Basically a sparse representation

```
■ SELECT A.i, B.j, SUM(A.val*B.val)
  FROM A, B
 WHERE A.j = B.i
 GROUP BY A.i, B.j;
```

*MAD Skills [VLDB'09]*



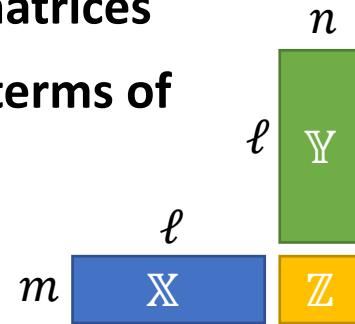
- Works pretty well for sparse matrices
- Not so good for dense matrices, but still beats “small-data” platforms when data doesn’t fit in memory

# Matrix Multiply: Take 2

- Data:  $A(i, \text{row}), B(j, \text{col})$  **MAD Skills [VLDB'09]**
  - `row` and `col` are `ARRAY` types or user-defined vector types
  - Basically a row-/column-major representation
- UDF (user-defined function):  $\text{dotproduct}(v_1, v_2)$  computes the dot product of two vectors

```
SELECT A.i, B.j, dotproduct(A.row, B.col)
FROM A, B;
```

- Works fine for dense matrices
- But still suboptimal in terms of compute-to-I/O ratio



Computation:  $O(\ell mn)$ , or volume  
I/O:  $O(m\ell + \ell n + nm)$ , or surface  
☞ Want instead “blocky” units to maximize compute-to-I/O ratio

- Also note the change in representation (from input to output)

# Matrix Multiply: Take 3

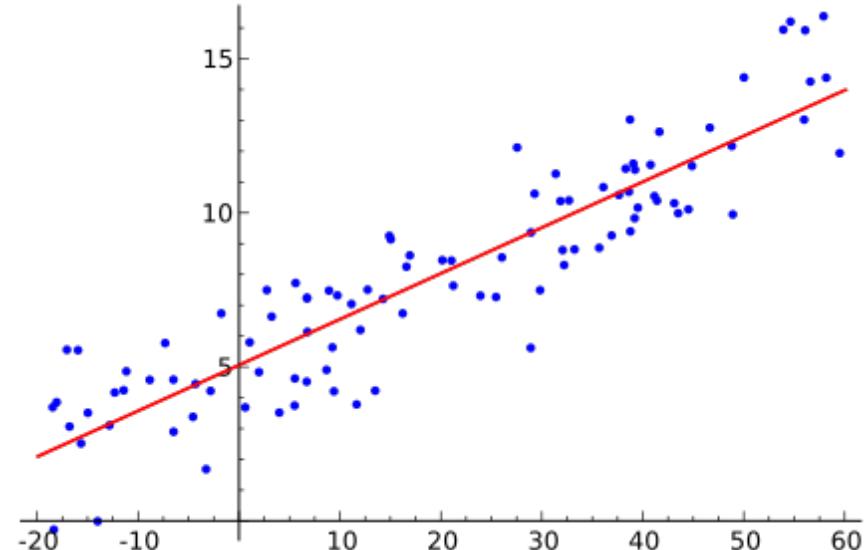
- **Data:  $A(\underline{i}, \underline{j}, V)$ ,  $B(\underline{i}, \underline{j}, V)$**  *RIOT-DB [CIDR'09]* *SimSQL [ICDE'17]*
  - $V$  represents a submatrix; assume the dimensions are compatible
  - Basically a blocked representation
- **UDFs**
  - $\text{matmult}(V_1, V_2)$  computes the product of two matrices
  - $\text{matsum}(V)$  is a UDA (user-defined aggregate) that sums up input matrices

```
SELECT A.i, B.j, matsum(matmult(A.V, B.V))
FROM A, B
WHERE A.j = B.i
GROUP BY A.i, B.j;
```
- **Choose a “big enough”  $V$  with good aspect ratio**
  - E.g., square  $V$ 's beat skinny  $V$ 's
- **UDFs can use optimized libraries like BLAS**

# Ordinary Least Squares

- To fit data  $(X, y)$  to a linear model  
 $y = X\beta + \epsilon$ :

$$\beta^* = (X^T X)^{-1} X^T y$$



- Computation involves basic matrix operators expressible in SQL with help of UDFs

- Inverse is tougher, but assuming the input matrix is small:
  - Code it as a UDF with memory-resident input
  - Processing won't benefit from DBMS though

*MAD* [VLDB'09, '12]  
*SimSQL* [ICDE'17]

# Observation

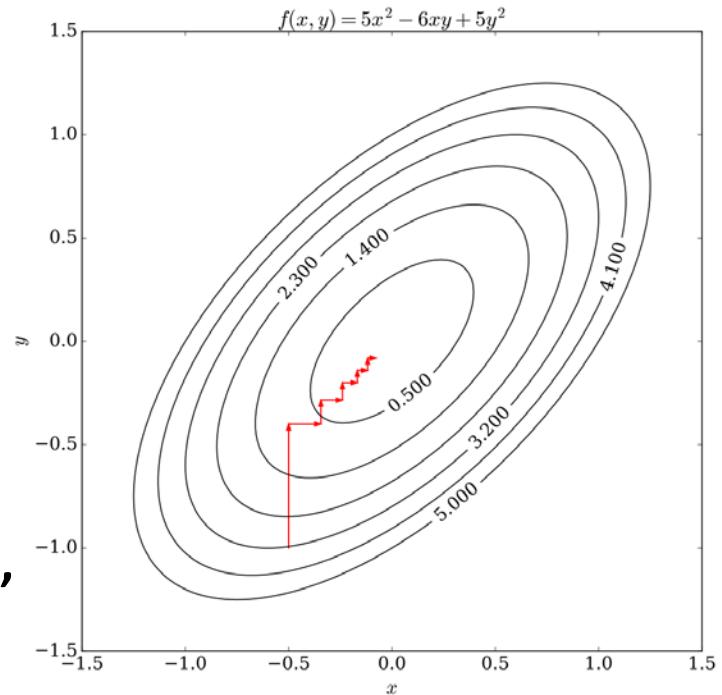
- **How far can UDF and UDA go? Surprisingly very!**
- **UDF (oftentimes coded in other languages, e.g., Python and R)**
  - Either on the tuple-level (invoked by SQL queries),
  - Or like an application program (invoking SQL queries)
- **UDA**
  - `Init(state)` initializes the state
  - `Accumulate(state, data)` computes updated state with new data
  - [optional] `Merge(state, state)` merges intermediate results computed over disjoint input subsets
  - `Finalize(state)` computes the final result from the state

☞ This pattern covers lots of iterative computation in ML, e.g.

  - $k$ -means (backup slides) ***GLADE*** [**LADIS'11, SIGMOD'12**], ***MADlib*** [**VLDB'12**]
  - Gradient descent (next)

# Gradient Descent (GD)

- Given a model with parameters  $w$ , we want to learn from data  $D$ , i.e., minimize a loss function  $F(w; D)$ 
  - E.g., sum of loss over all training data + a regularization term
  
- Start with some guess  $w_0$
- In each step  $t + 1$ , update  $w$  in the direction of the gradient of the loss function at  $w_t$ , i.e.,  $F'(w_t)$
- Rinse and repeat
  
- Under certain (commonly held) conditions, GD converges to a local minimum
  - If  $F$  is convex, that's its global minimum



# Stochastic GD (SGD)

- Suppose  $F(w; D)$  is linearly separable over  $D$ 
  - I.e.,  $F(w; D) = \sum_i f_i(w; d_i)$ ,  
where  $i$  iterates over the data points  $D = \{d_i\}_i$
- Instead of updating  $w$  using the “full gradient” computed over  $D$  in each GD step, just choose a single point in  $D$ 
  - I.e., use  $f'_i(w)$  to approximate  $F'(w)$
- Remarkably, for convex  $F(w)$ , SGD also converges to the global minimum, even if we pick points from  $D$  in a fixed, arbitrary order
  - Albeit at a slower rate

# GD/SGD in SQL

- **GD (full gradient)**
  - Computation of full gradient over  $D$  can be done by a query using UDA
  - Several options for driving outer loop
    - *MADlib* [VLDB'12] uses Python UDF
    - *ScalOps* [DeBull'12] uses Datalog
      - Underlying implementation is MapReduce instead of SQL
- **SGD *Bismarck* [SIGMOD'12]**
  - The entire procedure can be written as a query over  $D$  using UDA—each *Accumulate()* corresponds to one step

# MCMC in SQL

- **MCMC (Markov-Chain Monte-Carlo) is a key method in Bayesian ML**
- **Bayesian ML comes down to analyzing the “posterior” distribution**  
 $P(\text{parameters, hidden variables} \mid \text{observations})$
- **Direct analysis is often hard, so we use Monte-Carlo simulation**
  - Repeatedly sample from the posterior, and analyze the samples
- **But sampling directly from the posterior is often hard, so we use MCMC**
  - A sampler generates a Markov chain of samples, whose stationary distribution is the target posterior

- ☞ You can do Gibbs sampling (a form of MCMC) in *SimSQL [SIGMOD'13]*
- With user-defined “value-generating” functions that draw samples
  - See backup slides for details

# Approaches to SQL+ML

## Backend choices

- “On top of” (e.g., *RIOT-DB* [CIDR'09], *MAD* [VLDB'09, VLDB'12]) vs. “inside” DBMS (e.g., *SimSQL* [ICDE'17])
- Not DBMS, but still inspired by or rooted in DBMS
  - General-purpose “big-data” platform (e.g., *SystemML* [ICDE'11, VLDB'16], *Cumulon* [SIGMOD'13])
  - Specialized system from ground up (e.g., *RIOT* [ICDE'10], *SciDB* [CSE'13])

## Interface choices

- SQL + libraries or extensions (e.g., *MAD* [VLDB'09, VLDB'12], *SimSQL* [ICDE'17], *Oracle Data Mining*, ...)
- ML-oriented languages on top of SQL (e.g., *RIOT-DB* [CIDR'09], *BUDS/SimSQL* [SIGMOD'17], *Oracle R Enterprise*, ...)

# Interface: SQL + Libraries/Extensions

- Especially nice with integrated model management, e.g.,  
*Oracle Data Mining*

- Can create, store, update, and apply models in SQL

```
-- Create model settings:
CREATE TABLE svm_settings(
    setting_name VARCHAR2(30), setting_value VARCHAR2(30));
INSERT INTO svm_settings VALUES(
    dbms_data_mining.algo_name,
    dbms_data_mining.algo_support_vector_machines);

-- ...
-- Build model:
DBMS_DATA_MINING.CREATE_MODEL(
    model_name => 'svm_model',
    mining_function => dbms_data_mining.classification,
    data_table_name => 'mining_data_build_v',
    case_id_column_name => 'cust_id',
    target_column_name => 'affinity_card',
    settings_table_name => 'svm_settings');

-- Apply model:
DBMS_DATA_MINING.APPLY(
    model_name => 'svm_model',
    data_table_name => 'mining_data_apply_v',
    case_id_column_name => 'cust_id',
    result_table_name => 'svm_apply_result');
```

# Interface: no SQL

- Let user write w/o SQL
  - Provide a library that handles the underlying storage
  - SQL underneath (e.g., *Oracle R Enterprise*)
  - Other “big-data” engines (*Apache Mahout*, *Spark R*, *Mahout*)

Bayesian LASSO in *B*

```

invGamma = externalFunction(...)...  

invGaussian = externalFunction(...)...  

multiNormal = externalFunction(...)...  
  

X = read("test_data/xb.bin", format="binary")  

y = read("test_data/yb.bin", format="binary")  

y_avg = avg(y)  

y = y - y_avg  
  

# compute the matrix X'X, and X'Y  

XX = t(X) %*% X  

XY = t(X) %*% y  
  

# number of data points and number of features  

n = nrow(X)  

m = ncol(X)  
  

shape_prior = 1.0  

scale_prior = 1.0  

mean_prior = matrix(1.0, rows=1, cols=m)  

sigma2= invGamma(shape_prior, scale_prior)  

tau= invGaussian(mean_prior, shape_prior)  
  

niter = 5  
  

for (i in 1:niter) {  
  

  A = XX + diag(t(tau))  

  A_inv = inv(A)  

  mu = A_inv %*% XY  

  covariance = A_inv * sigma2  

  beta = multiNormal(t(mu), covariance)  

  remain_sum1 = (t(y) - beta %*% t(X))  

    (y - X %*% t(beta)) / 2.0  

  remain_sum2 = (beta * beta) %*% t(tau) / 2.0  

  scale_m = 1.0 + remain_sum1 + remain_sum2  

  scale = as.scalar(scale_m[1,1])  

  shape = 1.0 + (n-1.0)/2.0 + m/2.0  

  sigma2 = invGamma(shape, scale)  

  tau_mu = sqrt(sigma2 / (beta * beta))  

  tau = invGaussian(tau_mu, 1.0)
}

```

R, Python, etc.)  
can be implemented by  
SystemML [SIGMOD'17],  
SystemML [ICDE'11, VLDB'16],

---

```

externalFunction(...)...  

= externalFunction(...)...  

= externalFunction(...)...  
  

est_data/xb.bin", format="binary")  

est_data/yb.bin", format="binary")  

(y)  

vg  
  

he matrix X'X, and X'Y  

*% X  

*% Y  
  

data points and number of features  
  

= 1.0  

= 1.0  

= matrix(1.0, rows=1, cols=m)  

Gamma(shape_prior, scale_prior)  

ssian(mean_prior, shape_prior)  
  

:niter) {  
  

  diag(t(tau))  

  nv(A)  

  v %*% XY  

  ...ce = A_inv * sigma2  

  multiNormal(t(mu), covariance)  

  sum1 = (t(y) - beta %*% t(X))  

  X %*% t(beta)) / 2.0  

  sum2 = (beta * beta) %*% t(tau) / 2.0  

  = 1.0 + remain_sum1 + remain_sum2  

  as.scalar(scale_m[1,1])  

  1.0 + (n-1.0)/2.0 + m/2.0  

  = invGamma(shape, scale)  

  = sqrt(sigma2 / (beta * beta))  

  nvGaussian(tau_mu, 1.0)
}

```

---

... in SystemML

# Summary

- You can get a lot of mileage for machine learning with SQL+UDF (octopus + hammer)
- Deep roots in
  - DBMS extensibility research
  - Array DBMS, e.g., *SciDB* [CSE'13]; see *Rusu & Cheng* [arXiv 2013] for survey
- Next: more opportunities for deeper ML+DB integration



# References for Part 2: ML with SQL & UDF

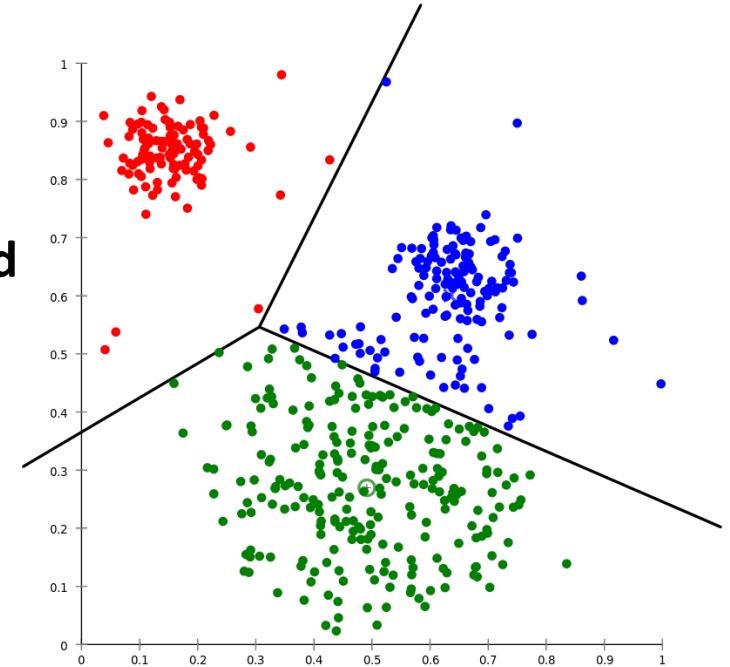
- **Bismarck [SIGMOD'12]** Feng et al. “Towards a Unified Architecture for in-RDBMS Analytics.” SIGMOD 2012
- **BUDS/SimSQL [SIGMOD'17]** Gao et al. “The BUDS Language for Distributed Bayesian Machine Learning.” SIGMOD 2017
- **Cumulon [SIGMOD'13]** Huang et al. “Cumulon: optimizing statistical data analysis in the cloud.” SIGMOD 2013
- **GLADE [LADIS'11]** Rusu & Dobra. “GLADE: A Scalable Framework for Efficient Analytics.” LADIS 2011
- **GLADE [SIGMOD'12]** Cheng et al. “GLADE: Big Data Analytics Made Easy.” SIGMOD 2012
- **MAD Skills [VLDB'09]** Cohen et al. “MAD skills: new analysis practices for big data.” PVLDB 2(2), 2009
- **MADlib [VLDB'12]** Hellerstein et al. “The MADlib Analytics Library or MAD Skills, the SQL.” PVLDB 5(12), 2012
- **RIOT-DB [CIDR'09]** Zhang et al. “RIOT: I/O-efficient numerical computing without SQL.” CIDR 2009
- **RIOT [ICDE'10]** Zhang et al. “I/O-efficient statistical computing with RIOT.” ICDE 2010
- **Rusu & Cheng [arXiv 2013]** Rusu & Cheng. “A Survey on Array Storage, Query Languages, and Systems.” <https://arxiv.org/abs/1302.0103>
- **ScalOps [DeBull'12]** Borkar et al. “Declarative systems for large-scale machine learning.” IEEE Data Eng. Bulletin, 35(2), 2012
- **SciDB [CSE'13]** “SciDB: A Database Management System for Applications with Complex Analytics.” Comp. Sci. Eng. 15(3), 2013
- **SimSQL [SIGMOD'13]** Cai et al. “Simulation of Database-Valued Markov Chains Using SimSQL.” SIGMOD 2013
- **SimSQL [ICDE'17]** Luo et al. “Scalable Linear Algebra on a Relational Database System.” ICDE 2017
- **SystemML [ICDE'11]** Ghoting et al. “SystemML: Declarative machine learning on MapReduce.” ICDE 2011
- **SystemML [VLDB'16]** Boehm et al. “SystemML: Declarative machine learning on Spark.” PVLDB 9(13), 2016

# Part 2 Backup/Extra Slides

# $k$ -Means Clustering

- Given  $n$  points, find  $k$  centroids to minimize sum of squared distances between each point and its closest centroid

- EM-style iterative algorithm:
  1. Pick initial  $k$  candidate centroid locations
  2. Assign each point to the closest candidate
  3. Reposition each candidate as the centroid of its assigned points
  4. Repeat 2-3 above until assignment changes no more (or very little)



# $k$ -Means as UDA

- **State:**  $k$  candidates with locations + cluster info  
 $\{\langle \text{loc}_i, \text{sum}_i, \text{cnt}_i \rangle\}_{1 \leq i \leq k}$
- **Init:** given centroid locations, with sum and count of 0
- **Accumulate:** given a data point  $p$ , find the candidate  $i$  closest to  $p$ ; increment  $\text{sum}_i$  by  $p$ 's coordinates and  $\text{cnt}_i$  by one
- **Merge:** merge  $\langle \text{loc}, \text{sum}, \text{cnt} \rangle$  records by loc; add sum and cnt
- **Finalize:** for each  $i$ , compute new  $\text{loc}_i$  as  $\text{sum}_i / \text{cnt}_i$
- **One SQL query with this UDA gives** *GLADE* [LADIS'11, SIGMOD'12]  
**one iteration of the EM algorithm** *MADlib* [VLDB'12]
  - For the next iteration, the UDA will be initialized with the  $k$  locations computed from the previous
  - Can use a UDF to drive overall iterations
  - Termination condition can be evaluated in SQL too (see *MADlib*)

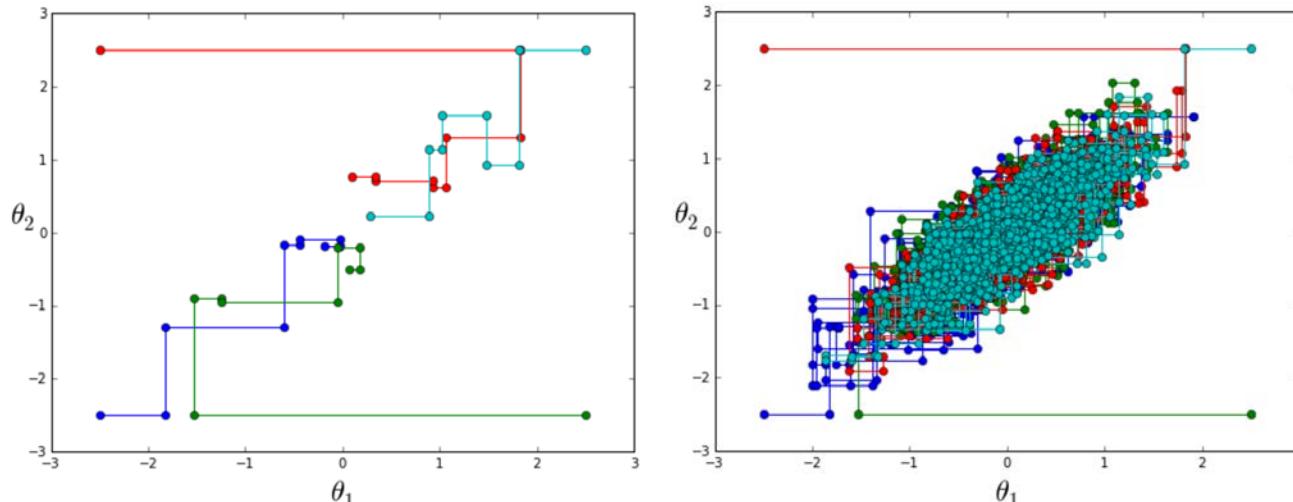
# Markov-Chain Monte-Carlo (MCMC)

- **Bayesian ML comes down to analyzing the “posterior” distribution**  
 $P(\text{parameters, hidden variables} \mid \text{observations})$
- **Direct analysis is often hard, so we use Monte-Carlo simulation**
  - Repeatedly sample from the posterior, and analyze the samples
- **But sampling directly from the posterior is often hard, so we use MCMC**
  - A sampler generates a Markov chain of samples, whose stationary distribution is the target posterior

# Example: Gibbs Sampling

- Suppose we have an  $n$ -variate distribution, but the conditional distributions are easier to sample from
- Begin with some initial sample  $\mathbf{z}^{(0)}$
- For the  $(t + 1)$ -th sample  $\mathbf{z}^{(t+1)}$ , sample each component  $z_i^{(t+1)}$  conditioned on all other components sampled most recently, i.e.,  

$$p\left(z_i^{(t+1)} \mid z_1^{(t+1)}, \dots, z_{i-1}^{(t+1)}, z_{i+1}^{(t)}, z_n^{(t)}\right)$$
- Rinse and repeat



# MCMC in SimSQL

*SimSQL [SIGMOD'13]*

- Think of each sample as a table (tables)
- Write UDF to define “VG” (value-generating) functions that draw samples
- Write SQL with VG functions to define how to generate  $T[t]$  (instance of table T in the  $t$ -th sample) from  $T[t - 1]$
- Write SQL to simulate multiple MCMC chains, and to compute distributional properties for variables of interest from  $T[t]$ ’s across T’s,  $t$ ’s, and chains

## ☞ An example of staying true to the declarative roots of databases

- But also need new techniques not in traditional DBMS, e.g.:
  - Plans are huge—cut them into “frames”; observe execution stats of last frame and to optimize the next
  - Use “tuple bundles” to instantiate/process multiple possible worlds simultaneously

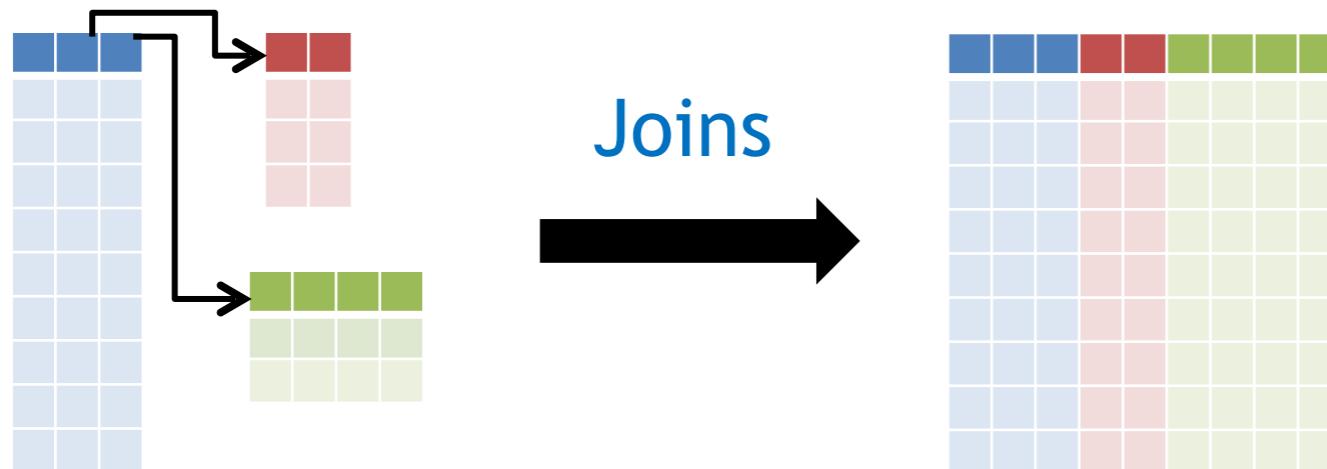
# Part 3: Learning Over Joins, SRL, and Deep RDBMS Integration

Arun Kumar  
UC San Diego  
La Jolla, CA, USA

SIGMOD 2017

# Overview: Learning Over Joins

**Problem:** Many datasets are multi-table      ↔      ML toolkits assume single-table inputs      →      ML after joining tables



**Overheads:**  
Extra storage  
Computational redundancy  
Join time  
Maintenance headaches

## Learning Over Joins: “Push Down” ML through joins

- 1) Over standard data systems: Orion, Santoku, Morpheus
- 2) Over a “factorized database” system: FDB-F
- 3) Special-purpose tools: libFM, TensorDB, Compressed ML

*Related but orthogonal:* Statistical relational learning (DeepDive, etc.)

# Learning Over Joins

Over standard data systems: Orion, Morpheus, Santoku

**Example:** GLMs with gradient descent (GD)

$$L(w) = \sum_{i=1}^n f(w'x_i, y_i) \quad \nabla L(w) = \sum_{i=1}^n g(w'x_i, y_i)x_i$$
$$w'x = w'_S x_S + w'_R x_R \quad x = [x_S \ x_R]$$
$$T = S \bowtie R$$

**Orion [SIGMOD'15]:**

Introduced the scalable “factorized learning” idea

Easy UDA implementation on existing data systems (RDBMS/Hive/Spark)

**Morpheus [VLDB'17]:**

Generalizes factorized learning to any ML algorithm in *linear algebra*

“Push down” rewrites for matrix-vector mult., gramian, ginv, etc.

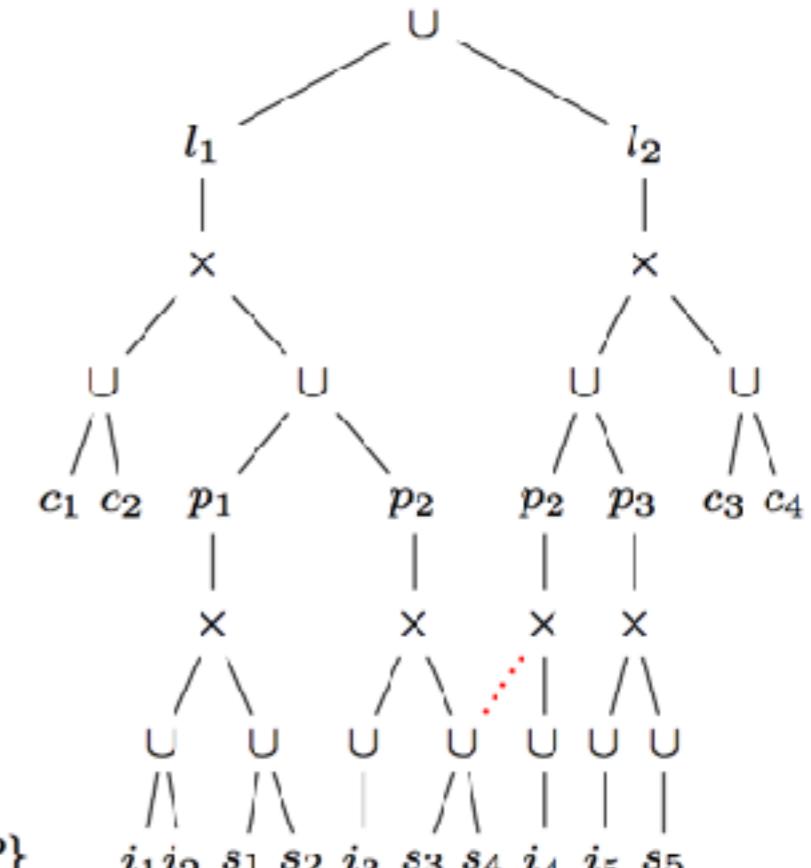
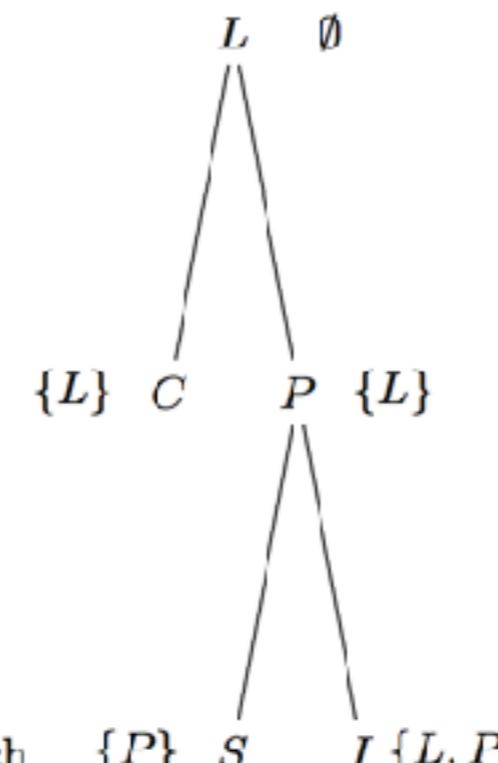
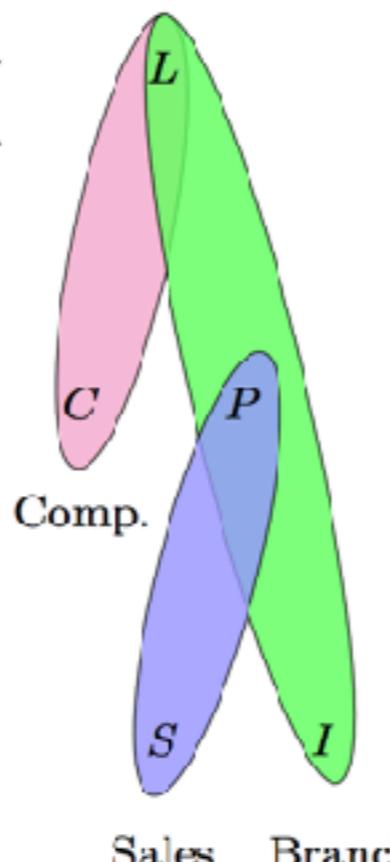
**Santoku [VLDB'15]:** Discrete features (Naive Bayes, trees, etc.)

# Learning Over Joins

Over a “factorized database” system: **FDB-F [SIGMOD’16]**

Generalized semiring-based aggregates over “factorized joins”

Sales	Branch	Natural Join				
<u>P S</u>	<u>L P I</u>	<u>L</u>	<u>C</u>	<u>P</u>	<u>I</u>	<u>S</u>
$p_1 s_1$	$l_1 p_1 i_1$	$l_1$	$c_1$	$p_1$	$i_1$	$s_1$
$p_1 s_2$	$l_1 p_1 i_2$	$l_1$	$c_1$	$p_1$	$i_1$	$s_2$
$p_2 s_3$	$l_1 p_2 i_3$	$l_1$	$c_1$	$p_1$	$i_2$	$s_1$
$p_2 s_4$	$l_2 p_2 i_4$	$l_1$	$c_1$	$p_1$	$i_2$	$s_2$
$p_3 s_5$	$l_2 p_3 i_5$	$l_1$	$c_1$	$p_2$	$i_3$	$s_3$
		$l_1$	$c_1$	$p_2$	$i_3$	$s_4$
		.....				
<u>Competition</u>		above block for $c_2$				
		.....				
		$l_2$	$c_3$	$p_2$	$i_4$	$s_3$
		$l_2$	$c_3$	$p_2$	$i_4$	$s_4$
		$l_2$	$c_3$	$p_3$	$i_5$	$s_5$
		.....				
		above block for $c_4$				



# SRL; Deep RDBMS Integration

SRL combines statistical learning with logic-based rules/constraints

“Non-IID” ML models  
(MVDs, EMVDs, JDs)      *NIPS’12 tutorial by Lise Getoor*  
                                 *Book with Ben Taskar*

Inference and learning often perform joins internally!

Scalable grounding using RDBMS: Tuffy [VLDB’10]

Incremental maintenance: IncrementalDeepDive [VLDB’15]

*Increasing interest in deeper integration of ML into DBMS kernel!*

SAP HANA SLACID: Linear algebra kernels in an RDBMS [SSDBM’14]

New compressed sparse row/col. representations

Integrated API for basic access patterns and lin. alg. ops

OpenMP-based shared memory parallelism in DBMS task scheduler

# References: Part 3

- Columbus [SIGMOD'14]: Materialization Optimizations for Feature Selection Workloads  
DeepDive [DataEng'14]: Feature Engineering for Knowledge Base Construction  
FDB-F [SIGMOD'16]: Learning Linear Regression Models over Factorized Joins  
IncrementalDeepDive [VLDB'15]: Incremental Knowledge Base Construction Using DeepDive  
Morpheus [VLDB'17]: Towards Linear Algebra over Normalized Data  
Orion [SIGMOD'15]: Learning Generalized Linear Models Over Normalized Data  
Santoku [VLDB'15]: Demonstration of Santoku: Optimizing Machine Learning over Normalized Data  
SLACID [SSDBM'14]: SLACID - Sparse Linear Algebra in a Column-Oriented In-Memory Database System  
Tuffy [VLDB'10]: Tuffy: Scaling up Statistical Inference in Markov Logic Networks using an RDBMS

# Backup Slides

# Statistical Relational Learning Systems

Captures logical dependencies between entities/variables

“Non-IID” ML models  
(MVDs, EMVDs, JDs)

*PODS tutorial by Lise Getoor on Tue!*  
*(also NIPS’12; book with Taskar)*

**Example:** Markov Logic Network (MLN); used by DeepDive

	weight	rule		
paper(PaperID, URL)	5	$\text{cat}(p, c1), \text{cat}(p, c2) \Rightarrow c1 = c2$	$(F_1)$	wrote(‘Joe’, ‘P1’)
wrote(Author, Paper)	1	$\text{wrote}(x, p1), \text{wrote}(x, p2), \text{cat}(p1, c) \Rightarrow \text{cat}(p2, c)$	$(F_2)$	wrote(‘Joe’, ‘P2’)
refers(Paper, Paper)	2	$\text{cat}(p1, c), \text{refers}(p1, p2) \Rightarrow \text{cat}(p2, c)$	$(F_3)$	wrote(‘Jake’, ‘P3’)
cat(Paper, Category)	$+\infty$	$\text{paper}(p, u) \Rightarrow \exists x. \text{wrote}(x, p)$	$(F_4)$	refers(‘P1’, ‘P3’)
	-1	$\text{cat}(p, ‘Networking’)$	$(F_5)$	cat(‘P2’, ‘DB’)
				...

Schema

Λ Markov Logic Program

Evidence

MLN inference (MAP) computes “most probable world” by plugging values of variables to predict

Grounding + Search  
Involves joins!

Scalable grounding using RDBMS: Tuffy [VLDB’10]

Scalable Gibbs sampling: Elementary [SIGMOD’13]

Incremental maintenance: IncrementalDeepDive [VLDB’15]

# Deep RDBMS Integration

**Integrating linear algebra kernels into an RDBMS: SAP HANA**

**SLACID [SSDBM'14]:** Mutable columnar layout for sparse matrices

- Compressed sparse row/col. representation + incr. delta

- Integrated API for basic access patterns and lin. alg. ops

OpenMP-based shared memory parallelism in DBMS task scheduler

**Time series-specific systems: Fa, F2DB**

**Fa [VLDB'07]:** “Declarative forecasting” queries for time series

- Projection and shift-based time series feature transformations

- Feature ranking and subset selection heuristics

- Lin. reg., Bayesian networks, SVM, CART, Random Forest

- Both one-time and continuous forecasting

# Part 4: Rewrites, Operator Selection, and Operator Fusion

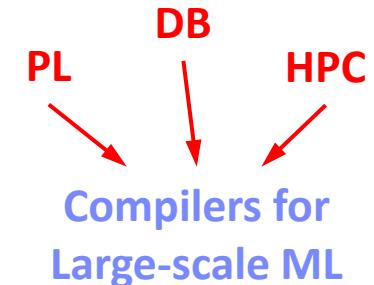
**Matthias Boehm**

IBM Research – Almaden  
San Jose, CA, USA

**SIGMOD 2017**

# Overview Optimizing Compilers for ML Algorithms

- **Comparison Query Optimization**
  - Rule- and cost-based rewrites and operator ordering
  - Physical operator selection and query compilation
  - Linear algebra / other ML operators, DAGs, control flow, sparse/dense formats
  
- **#1 Interpretation** (operation at-a-time)
  - Examples: [Morpheus](#) [PVLDB'17]
  
- **#2 Lazy Expression Compilation** (DAG at-a-time)
  - Examples: [RIOT](#) [CIDR'09], [Mahout Samsara](#) [MLSystems'16]
  - Examples w/ control structures: [Weld](#) [CIDR'17], [OptiML](#) [ICML'11], [Emma](#) [SIGMOD'15]
  
- **#3 Program Compilation** (entire program)
  - Examples: [SystemML](#) [PVLDB'16], [Cumulon](#) [SIGMOD'13], [Tupleware](#) [PVLDB'15]



## Optimization Scope

```

1: X = read($1); # n x m matrix
2: y = read($2); # n x 1 vector
3: maxi = 50; lambda = 0.001;
4: intercept = $3;
5: ...
6: r = -(t(X) %*% y);
7: norm_r2 = sum(r * r); p = -r;
8: w = matrix(0, ncol(X), 1); i = 0;
9: while(i<maxi & norm_r2>norm_r2_trgt)
10: {
11:   q = (t(X) %*% X %*% p) + lambda*p;
12:   alpha = norm_r2 / sum(p * q);
13:   w = w + alpha * p;
14:   old_norm_r2 = norm_r2;
15:   r = r + alpha * q;
16:   norm_r2 = sum(r * r);
17:   beta = norm_r2 / old_norm_r2;
18:   p = -r + beta * p; i = i + 1;
19: }
20: write(w, $4, format="text");
  
```

# Logical Simplification Rewrites

- **Traditional PL Rewrites** (e.g., TensorFlow, OptiML, SystemML)
  - CSE, constant folding, branch removal
- **Algebraic Simplification Rewrites** (e.g., SystemML, FAQ [PODS'16])
  - $t(X) \%*% y \rightarrow t(t(y) \%*% X)$
  - $\text{trace}(X \%*% Y) \rightarrow \text{sum}(X * t(Y))$
  - $\text{sum}(X + Y) \rightarrow \text{sum}(X) + \text{sum}(Y)$
  - $\text{sum}(X^2) \rightarrow t(X) \%*% X, \text{ iff } \text{ncol}(X)=1$
- **Loop Vectorization** (e.g., OptiML, SystemML)
 

```
for(i in a:b)
    X[i,1] = Y[i,2] + Z[i,1] → X[a:b,1] = Y[a:b,2] + Z[a:b,1]
```

- **Incremental Computations**
  - Delta update rules (e.g., **LINVIEW** [SIGMOD'14], factorized [CoRR'17])
  - Incremental iterations (e.g., Flink)       $A = t(X) \%*% X + t(\Delta X) \%*% \Delta X$
  - Update-in-place (e.g., SystemML)       $b = t(X) \%*% y + t(\Delta X) \%*% \Delta y$

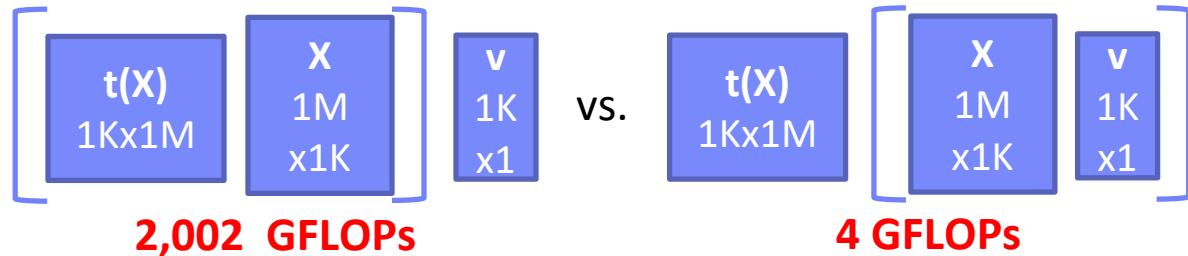
# Logical Simplification Rewrites

## Matrix Multiplication Chain Optimization

- **Optimization Problem**

- Matrix multiplication chain of n matrices  $M_1, M_2, \dots M_n$  (associative)
- Optimal parenthesization of the product  $M_1 M_2 \dots M_n$

**Example**  
 $t(X) \%*\% X \%*\% v$

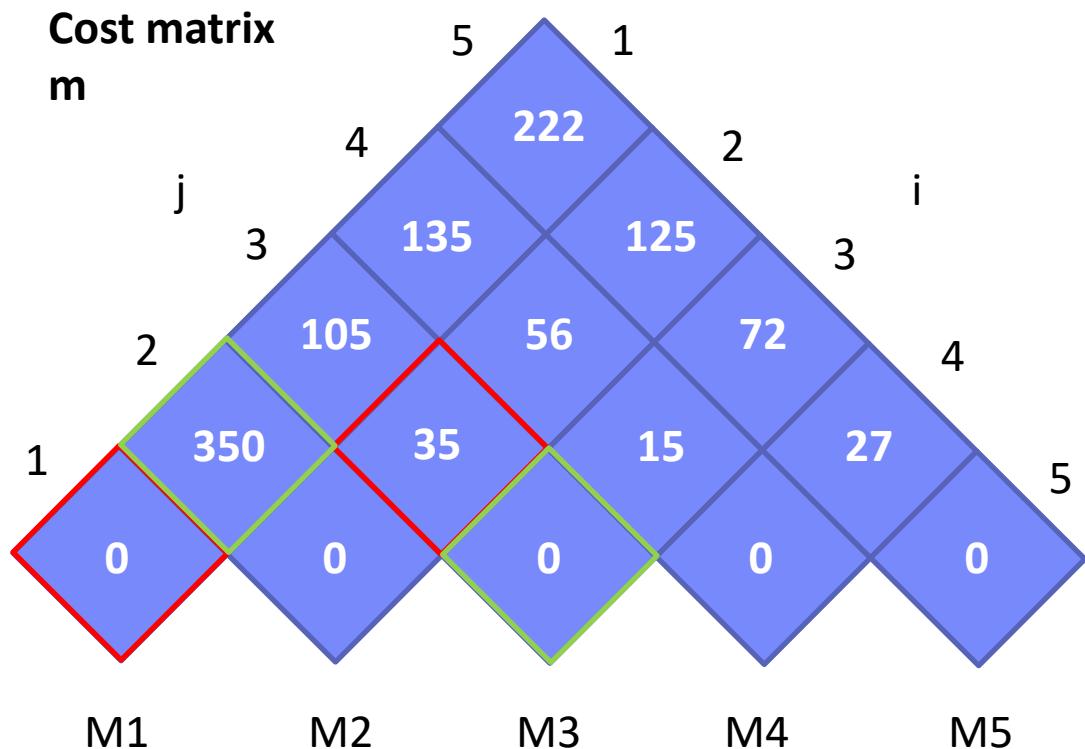


- **Search Space Characteristics**

- Naïve exhaustive: Catalan numbers  $\rightarrow \Omega(4^n / n^{3/2})$
- DP applies: (1) optimal substructure, (2) overlapping subproblems
- Textbook DP algorithm [MIT Press'09]:  $\Theta(n^3)$  time,  $\Theta(n^2)$  space
  - Examples: [SystemML](#) [Data Eng. Bull. '14], [RIOT](#) (including I/O costs), [SpMachO](#) (including sparsity for intermediates) [EDBT'15],
  - Best known algorithm:  $O(n \log n)$

# Matrix Multiplication Chain Optimization

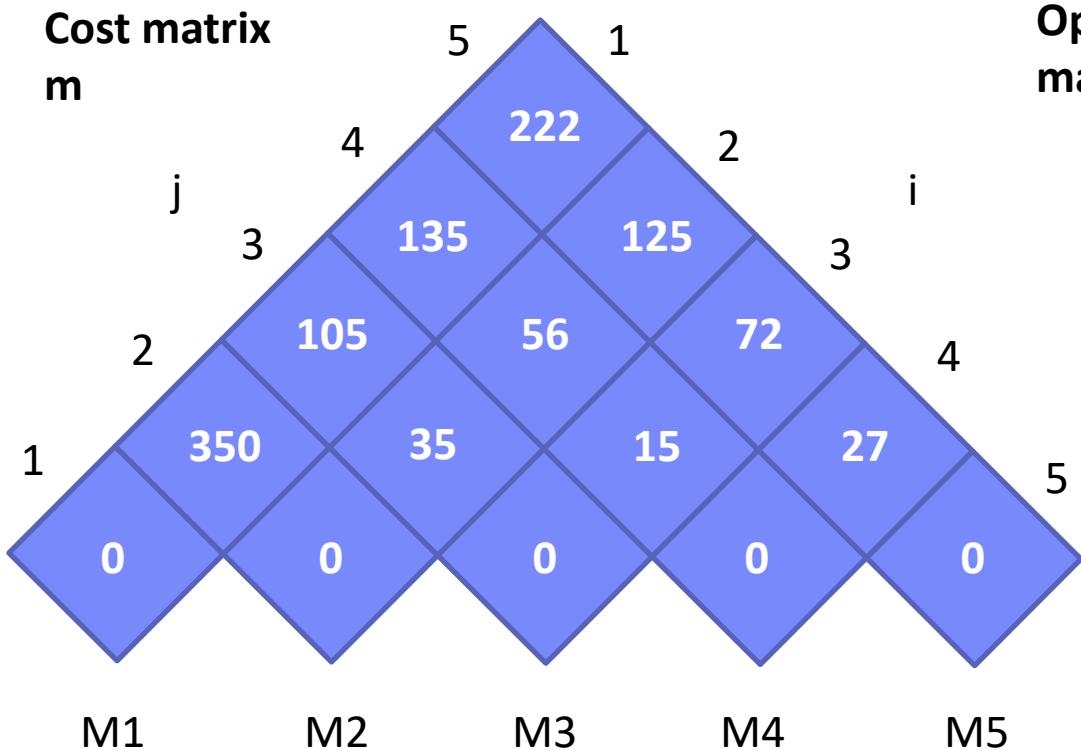
M1	M2	M3	M4	M5
10x7	7x5	5x1	1x3	3x9



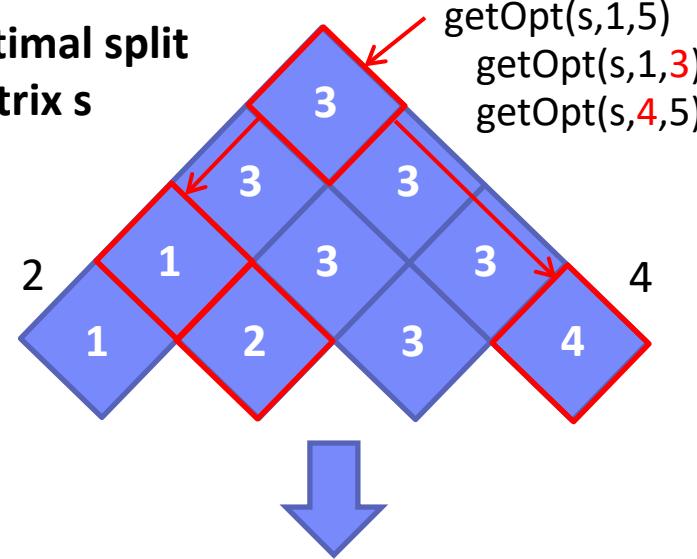
$$\begin{aligned}
 m[1,3] &= \min( \\
 &\quad m[1,1] + m[2,3] + p_1 p_2 p_4, \\
 &\quad m[1,2] + m[3,3] + p_1 p_3 p_4 ) \\
 &= \min( \\
 &\quad 0 + 35 + 10 * 7 * 1, \\
 &\quad 350 + 0 + 10 * 5 * 1 ) \\
 &= \min( \\
 &\quad 105, \\
 &\quad 400 )
 \end{aligned}$$

# Matrix Multiplication Chain Optimization

M1	M2	M3	M4	M5
10x7	7x5	5x1	1x3	3x9



**Optimal split matrix s**



( M1 M2 M3 M4 M5 )

( ( M1 M2 M3 ) ( M4 M5 ) )

( ( M1 ( M2 M3 ) ) ( M4 M5 ) )

→ ((M1 (M2 M3)) (M4 M5))

→ Open questions: DAGs; other operations,  
joint opt w/ rewrites, CSE, fusion, and physical operators

# Physical Rewrites and Optimizations

## ▪ Distributed Caching

- Redundant compute vs. memory consumption and I/O
- #1 Cache intermediates w/ multiple refs (Emma)
- #2 Cache initial read and read-only loop vars (SystemML)

## ▪ Partitioning

- Many frameworks exploit co-partitioning for efficient joins
- #1 Partitioning-exploiting operators (SystemML, Emma, Samsara)
- #2 Inject partitioning to avoid shuffle per iteration (SystemML)
- #3 Plan-specific data partitioning (SystemML, Dmac [SIGMOD'15], Kasen [VLDB'16])

## ▪ Other Data Flow Optimizations (Emma)

- #1 Exists unnesting (e.g., filter w/ broadcast → join)
- #2 Fold-group fusion (e.g., groupByKey → reduceByKey)

## ▪ Physical Operator Selection

# Physical Operator Selection

- **Common Selection Criteria**

- **Data and cluster characteristics** (e.g., data size/shape, memory, parallelism)
- **Matrix/operation properties** (e.g., diagonal/symmetric, sparse-safe ops)
- **Data flow properties** (e.g., co-partitioning, co-location, data locality)

- **#0 Local Operators**

- SystemML mm, tsmm, mmchain; Samsara/Mllib local linalg

Selection  
Preference

- **#1 Special Operators** (often fused operators)

- Special patterns (SystemML **tsmm**, tsmm2, mapmmchain, pmm; Samsara AtA)
- Sparsity exploiting (SystemML wdivmm, **wsloss**, wcemm; Cumulon maskMult)

- **#2 Broadcast-Based Operators** (aka broadcast join)

- SystemML **mapmm**, **mapmmchain**

- **#3 Co-Partitioning-Based Operators** (aka improved repartition join)

- SystemML zipmm; Emma, Samsara OpAtB

- **#4 Shuffle-Based Operators** (aka repartition join)

- SystemML cpmm, rmm; Samsara OpAB

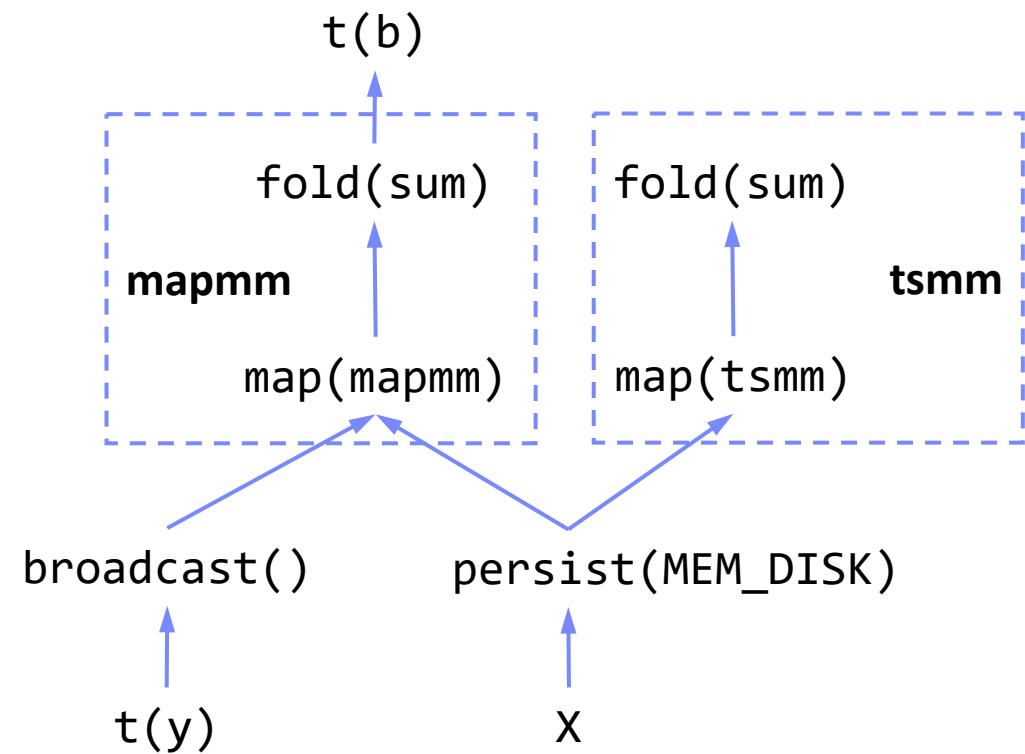
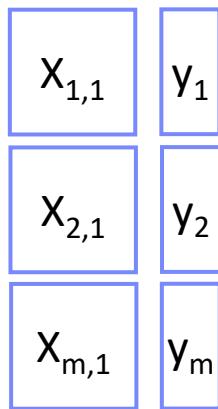
# Example Physical Operators

- Example Linear Regression Direct Solve

- Transpose-self for  $t(X) \%*% X$
- Broadcast-based for  $t(X) \%*% y$
- Logical and physical rewrites
- E.g., Samsara, SystemML

```
A = t(X) %*% X
b = t(X) %*% y
w = solve(A, b)
```

Input Matrices



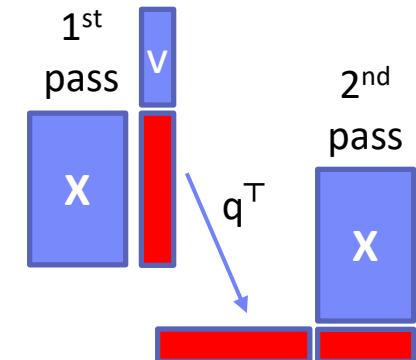
# Fused Operators

## ▪ Motivation

- Problem: **Memory-bandwidth-bound operations** (I/O)
- Goal: Reduce number of **scans and intermediates**

## ▪ Matrix-Vector Chains: $t(X) \%*\% (X \%*\% v)$

- Fused single-pass operator: **mmchain** [PPoPP'15]
- Row-aligned creation/consumption

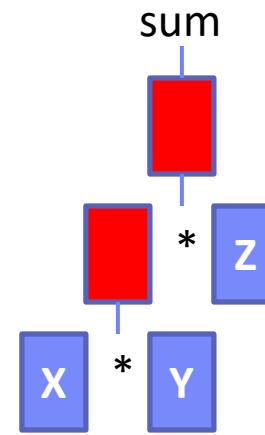


## ▪ Ternary Aggregates: $\text{sum}(X * Y * Z)$

- Fused aggregation operator
- Avoid materialized intermediates

## ▪ Other ML-Specific Operators

- Sample proportion:  $X * (1-X)$
- Sigmoid:  $1 / (1 + \exp(-X))$
- Axpy:  $X + s * Y, X - s * Y$

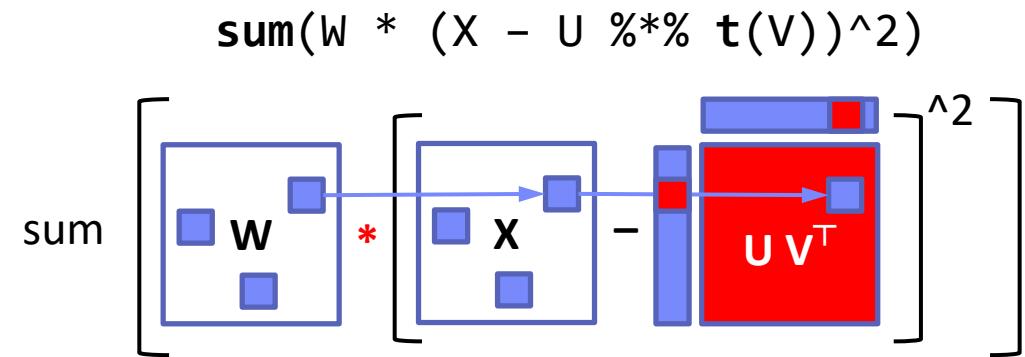


# Sparsity-Exploiting Fused Operators

- **Goal:** Avoid dense intermediates and unnecessary computation

- **#1 Fused Physical Operators**

- E.g., SystemML [PVLDB'16]  
wsloss, wcemm, wdivmm
- Selective computation over non-zeros of  
**“sparse driver”**

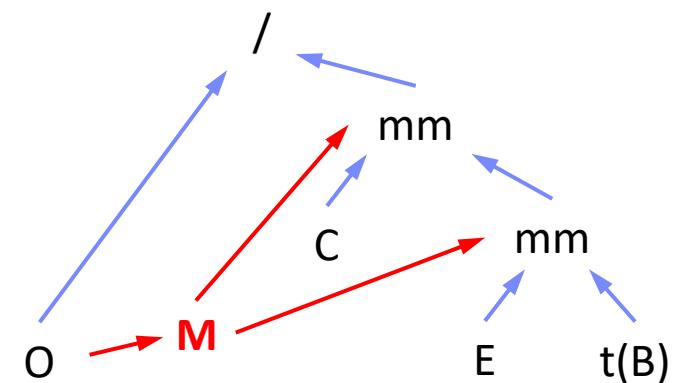


- **#2 Masked Physical Operators**

- E.g., Cumulon MaskMult [SIGMOD'13]
- Create mask of **“sparse driver”**
- Pass mask to single masked matrix multiply operator

➔ Open questions: NaN handling, automatic operator fusion (codegen)

$$O / (C \%*% E \%*% t(B))$$



# Automatic Operator Fusion

## ▪ Motivation

- Large development effort for hand-coded fused operators
- UDF-centric systems w/o pre-defined operators

## ▪ General Approach: Fuse by Access Pattern

- #1 Loop fusion (OptiML, Tupleware, Weld, TensorFlow XLA [github'17])
- #2 Templates (Kasen, SPOOF [CIDR'17])
- Scope: expression or program compilation

## ▪ Additional Techniques

- Tupleware: Micro optimizations (tile-at-a-time, predicates, result allocation)
- Weld: Cross-library optimizations (via common IR of basic operations)
- SystemML-SPOOF: sparsity-exploiting fused operators

→ Open question: Optimization of fusion plans for DAGs  
 (redundant compute vs materialization, access patterns)

```
R = (A + s*B) * C
for( i in 1:n )
    tmp[i] = s*B[i]
for( i in 1:n )
    tmp[i] = A[i]+tmp[i]
for( i in 1:n )
    tmp[i] = tmp[i]*C[i]
for( i in 1:n )
    tmp[i] = (A[i]+s*B[i]) * C[i]
```



# Runtime Adaptation (see AQP)

- **Problem of Unknown/Changing Size Information**
    - Dimensions/sparsity required for **cost comparisons/valid plans**
    - Unknowns → **conservative fallback plans**
  - **Challenges**
    - Conditional control flow, function call graphs, UDFs
    - Data-dependent ops (e.g., sampling, group by classes, output sparsity)
    - Computed size expressions, changing dimensions/sparsity
  - **Approaches**
    - **#1 Lazy expression optimization** (RIOT, OptiML, Emma, Weld, Samsara)
      - Optimize on triggering actions (unconditional scope)
    - **#2 Dynamic inter-DAG recompilation** (SystemML)
      - Split/mark DAGs, recompile DAGs/functions w/ exact stats
- **Open questions:**
- **Estimating the size and sparsity of intermediates**
  - **Adaptive query processing and storage**

# References for Part 4

- T. C. Hu and M. T. Shing: Computation of Matrix Chain Products. Part II. SIAM J. Comput. 13(2): 228-251, 1984.
- T. H. Cormen, et al. Introduction to Algorithms, Third Edition, The MIT Press, pages 370-377, 2009.
- Y. Zhang et al. RIOT: I/O-Ecient Numerical Computing without SQL. In CIDR, 2009.
- A. K. Sujeth et al. OptiML: An Implicitly Parallel Domain-Specific Language for Machine Learning. In ICML, 2011.
- S. Ewen et al. Spinning Fast Iterative Data Flows. PVLDB, 5(11), 2012.
- B. Huang et al. Cumulon: Optimizing Statistical Data Analysis in the Cloud. In SIGMOD, 2013.
- M. Nikolic et al. LINVIEW: Incremental View Maintenance for Complex Analytical Queries. In SIGMOD, 2014.
- M. Boehm et al. SystemML's Optimizer: Plan Generation for Large-Scale Machine Learning Programs. IEEE Data Eng. Bull., 37(3), 2014.
- D. Kernert et al. SpMacho - Optimizing Sparse Linear Algebra Expressions with Probabilistic Density Estimation. In EDBT, 2015.
- A. Ashari et al.: On optimizing machine learning workloads via kernel fusion. In PPoPP, 2015.
- A. Alexandrov et al. Implicit Parallelism through Deep Language Embedding. In SIGMOD, 2015.
- Lele Yu et al. Exploiting Matrix Dependency for Efficient Distributed Matrix Computation. In SIGMOD, 2015.
- M. Abo Khamis et al.: FAQ: Questions Asked Frequently. In PODS, 2016.
- A. Crotty et al. An Architecture for Compiling UDF-centric Workflows. PVLDB, 8(12), 2015.
- M. Boehm et al. SystemML: Declarative Machine Learning on Spark. PVLDB, 9(13), 2016.
- M. Zhang et al. Measuring and Optimizing Distributed Array Programs. PVLDB, 9(12), 2016.
- S. Schelter et al. Samsara: Declarative Machine Learning on Distributed Dataflow Systems. NIPS Workshop MLSystems, 2016.
- T. Elgamal et al. SPOOF: Sum-Product Optimization and Operator Fusion for Large-Scale Machine Learning. In CIDR, 2017.
- S. Palkar et al. Weld: A Common Runtime for High Performance Data Analysis. In CIDR 2017.
- TensorFlow XLA, <https://www.tensorflow.org/performance/xla/>, 2017.
- M. Nikolic and D. Olteanu: Incremental Maintenance of Regression Models over Joins, CoRR, 2017.
- L. Chen et al. Towards Linear Algebra over Normalized Data. PVLDB, to appear, 2017.

# Part 5: Compression, Scan Sharing, and Index Structures

**Matthias Boehm**

IBM Research – Almaden  
San Jose, CA, USA

**SIGMOD 2017**

# Motivation: Workload Characteristics

## ▪ Memory-Bandwidth-Bound Operations

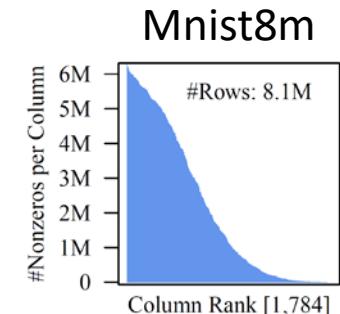
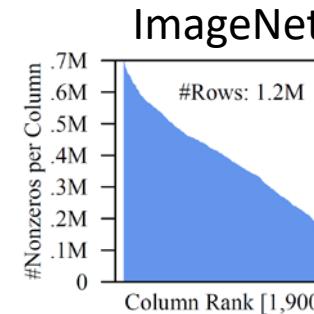
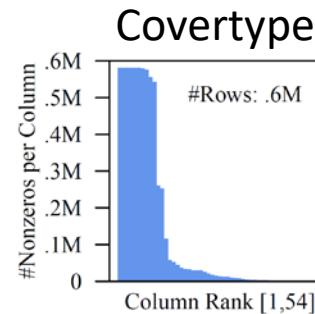
- Iterative ML algorithms w/ read-only data access
- #1: I/O-bound matrix vector products
  - ➔ Crucial to fit matrix into memory  
(single node, distributed, GPU)
  - ➔ Avoid unnecessary scans
- #2: Matrix and vector intermediates
  - ➔ Reduce number of reads and writes



```
while(!converged) {
    ... q = X %*% v ...
}
```

## ▪ Common Data Characteristics

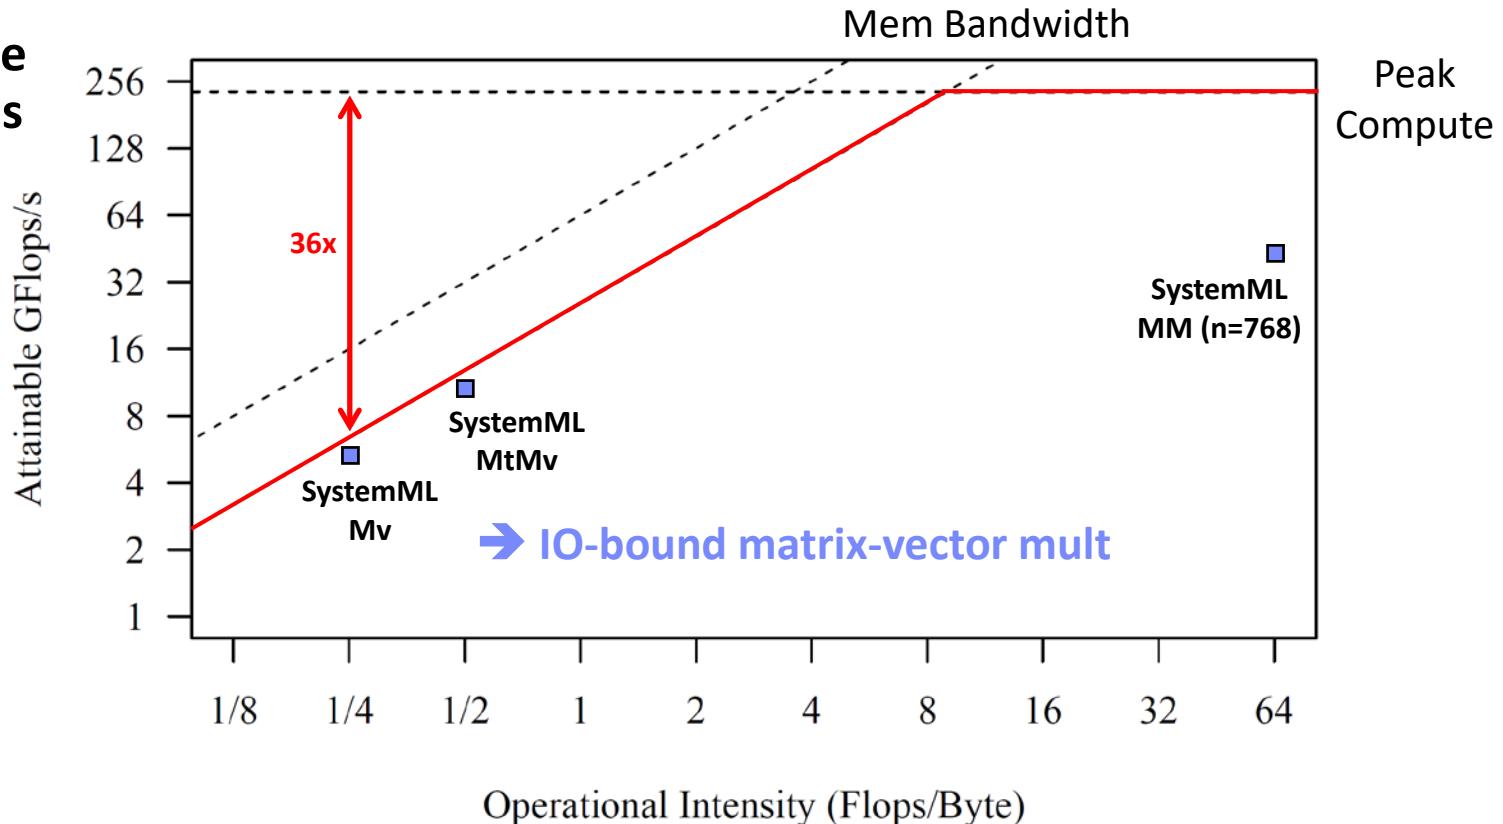
- Tall & skinny matrices (#row >> #columns)
- Non-uniform sparsity
- Low column cardinality
- Column correlations



# Motivation: Workload Characteristics

- **Single Node:** 2x6 E5-2440 @2.4GHz–2.9GHz, DDR3 RAM @1.3GHz (ECC)
  - Peak memory bandwidth: **2 x 32GB/s** (local), **2 x 12.8GB/s** (remote QPI)
  - Peak compute bandwidth: **2 x 115.2GFlops/s**

- **Roofline Analysis**



# Background: Block Partitioning and Layouts

## ■ Blocked Matrix Representations

- Blocks, a.k.a. “tiles”, “chunks”, or “pages”
- #1 **Logical (fixed-size) blocking** ( $\rightarrow$  var. physical size)
- #2 **Physical blocking** ( $\rightarrow$  fixed physical size)
- Blocks encoded independently (dense/sparse)
- Local matrices  $\rightarrow$  single block

Logical blocking  
3,400x2,700 matrix  
(w/  $B_c=1,000$ )

(1,1)	(1,2)	(1,3)
(2,1)	(2,2)	(2,3)
(3,1)	(3,2)	(3,3)
(4,1)	(4,2)	(4,3)

## ■ Common Block Representations

- **Dense** (linearized arrays)
- **CSR** (compressed sparse rows)
- **CSC** (compressed sparse columns)
- **MCSR** (modified CSR)
- **COO** (Coordinate matrix)
- ...

**Dense** (row-major)

.7	0	.1	.2	.4	0	0	.3	0
----	---	----	----	----	---	---	----	---

Example  
3x3 Matrix

.7		.1
.2	.4	
	.3	

**CSR**

0	0	.7
2	2	.1
4	0	.2
5	1	.4
	1	.3

**MCSR**

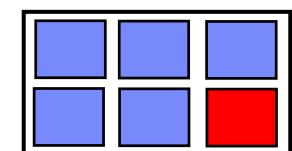
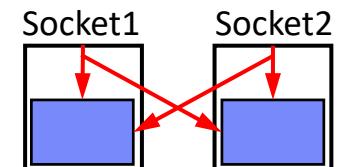
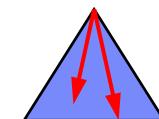
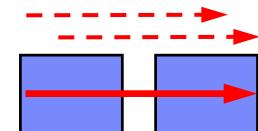
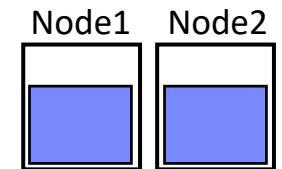
0	2	
.7	.1	
0	1	
.2	.4	
1		.3

**COO**

0	0	.7
0	2	.1
1	0	.2
1	1	.4
2	1	.3

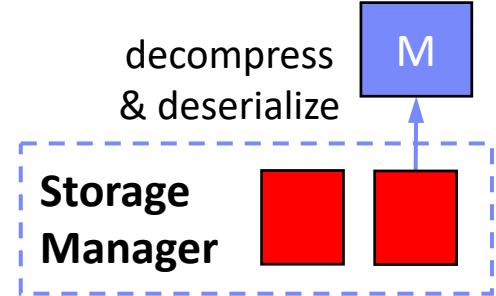
# Overview Techniques for Data-Intensive Machine Learning

- #1 **(Distributed) Caching**
  - Keep read only feature matrix in (distributed) memory
  
- #2 **Compression**
  - Fit larger datasets into available memory
  
- #3 **Scan Sharing (and operator fusion)**
  - Reduce the number of scans as well as read/writes
  
- #4 **Index Structures**
  - Out-of-core data, I/O-aware ops, updates
  
- #5 **NUMA-Aware Partitioning and Replication**
  - Matrix partitioning / replication → data locality
  
- #6 **Buffer Pool Management**
  - Graceful eviction of intermediates, out-of-core ops

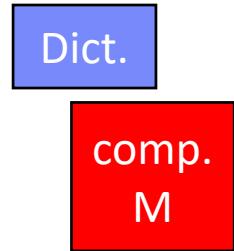


# Compression Techniques

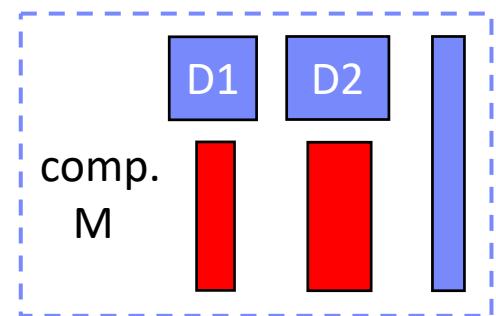
- **#1 Block-Level General-Purpose Compression**
  - Heavyweight or lightweight compression schemes
  - Decompress matrices block-wise for each operation
  - E.g.: Spark RDD compression (Snappy/LZ4),  
[SciDB](#) SM [SSDBM'11], [TileDB](#) SM [PVLDB'16]



- **#2 Block-Level Matrix Compression**
  - Compress matrix block with common encoding scheme
  - Perform LA ops over compressed representation
  - E.g.: [CSR-VI](#) (dict) [CF'08], [cPLS](#) (grammar) [KDD'16],  
[TOC](#) (LZW w/ trie) [CoRR'17]



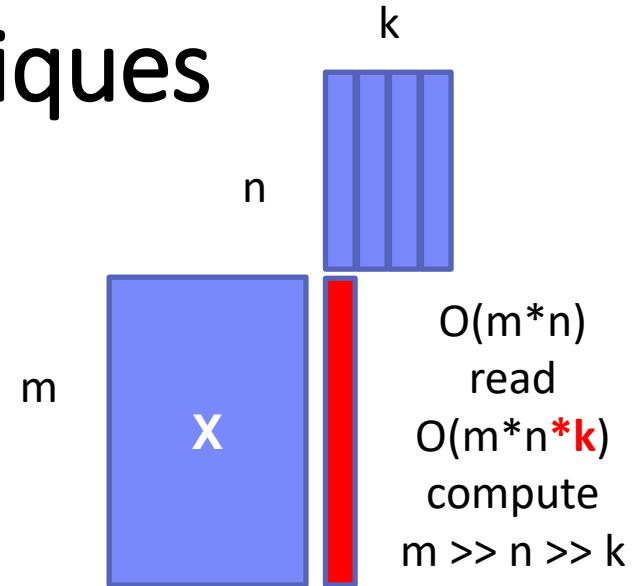
- **#3 Column-Group-Level Matrix Compression**
  - Compress column groups w/ heterogenous schemes
  - Perform LA ops over compressed representation
  - E.g.: [SystemML CLA](#) (RLE, OLE, DDC, UC) [PVLDB'16]



# Scan Sharing Techniques

## ■ #1 Batching

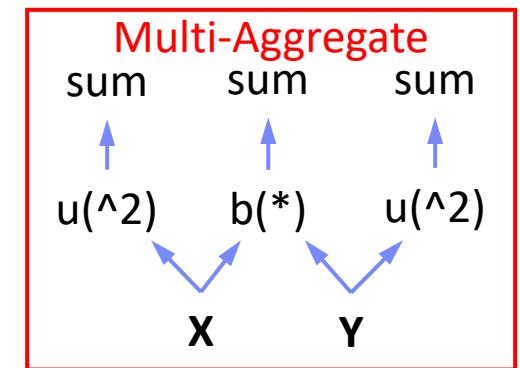
- One-pass evaluation of multiple configurations
- Use cases: EL, CV, feature selection, hyper parameter tuning
- E.g.: [TUPAQ](#) [SoCC'16], [Columbus](#) [SIGMOD'14]



## ■ #2 Fused Operator DAGs

- Avoid unnecessary scans, (e.g., part 4 mmchain)
- Avoid unnecessary writes / reads
- Multi-aggregates, redundancy
- E.g.: [SystemML codegen](#)

```
a = sum(X^2)
b = sum(X*Y)
c = sum(Y^2)
```



## ■ #3 Runtime Piggybacking

- Merge concurrent data-parallel jobs
- “Wait-Merge-Submit-Return”-loop
- E.g.: [SystemML parfor](#) [VLDB'14]

```
parfor( i in 1:numModels )
  while( !converged )
    q = X %*% v; ...
```

# Index Structures and NUMA Awareness

- **Goals:** Out-of-core operations and data placement
- **Index Structures**
  - Tree structures of blocks w/ user-defined/fixed linearization functions
  - **LAB-Tree** (Linearized Array B-tree, RIOT) [PVLDB'11]
    - Leaf-splitting strategies, and update batching via flushing policies
  - **TileDB Storage Manager** [PVLDB'16]
    - Two-level blocking and update batching via fragments
  - **AT MATRIX** (Adaptive Tile Matrix, SAP HANA) [ICDE'16]
    - Two-level blocking and NUMA-aware range partitioning
- **NUMA-Aware Model/Data Replication**
  - **DimmWitted**: HW vs statistical efficiency [PVLDB'14]
  - Model: PerCore, PerNode, PerMachine
  - Data: partitioning (sharding), full replication

→ Open questions:  
Heterogenous hardware,  
cache coherence, etc

# References for Part 5

- K. Kourtis et al. Optimizing Sparse Matrix-Vector Multiplication Using Index and Value Compression. In CF, 2008.
- S. Williams et al.: Roofline: An Insightful Visual Performance Model for Multicore Architectures. Comm. ACM 52(4) 2009.
- Y. Zhang et al. RIOT: I/O-Ecient Numerical Computing without SQL. In CIDR, 2009.
- M. Stonebraker et al. The Architecture of SciDB. In SSDBM, 2011.
- Y. Zhang et al. Storing Matrices on Disk: Theory and Practice Revisited. PVLDB, 4(11), 2011.
- T. Kraska et al. MLbase: A Distributed Machine-learning System. In CIDR, 2013.
- C. Zhang and C. Re. Towards High-throughput Gibbs Sampling at Scale: A Study Across Storage Managers. In SIGMOD, 2013.
- C. Zhang et al. Materialization Optimizations for Feature Selection Workloads. In SIGMOD, 2014.
- M. Boehm et al. Hybrid Parallelization Strategies for Large-Scale Machine Learning in SystemML. PVLDB, 7(7), 2014.
- C. Zhang and C. Re. DimmWitted: A Study of Main-Memory Statistical Analytics. PVLDB, 7(12), 2014.
- E. R. Sparks et al. Automating Model Search for Large Scale Machine Learning. In SoCC, 2015.
- D. Kernert et al. Topology-Aware Optimization of Big Sparse Matrices and Matrix Multiplications on Main-Memory Systems. In ICDE, 2016.
- A. Elgohary et al. Compressed Linear Algebra for Large-Scale Machine Learning. PVLDB, 9(12), 2016.
- M. Boehm et al. SystemML: Declarative Machine Learning on Spark. PVLDB, 9(13), 2016.
- Stavros Papadopoulos et al. The TileDB Array Data Storage Manager. PVLDB 10(4), 2016.
- T. Elgamal et al. SPOOF: Sum-Product Optimization and Operator Fusion for Large-Scale Machine Learning. CIDR, 2017.
- F. Li et al. When Lempel-Ziv-Welch Meets Machine Learning: A Case Study of Accelerating Machine Learning using Coding. CoRR, 2017.

# Backup: Compressed Linear Algebra (CLA)

[PVLDB 2016]

## ▪ Overview compression framework

- Column-wise matrix compression (values + offset lists / references)
- Column co-coding (column groups encoded as single unit)
- Heterogeneous column encoding formats (OLE, RLE, DDC, UC)

Automatic Planning

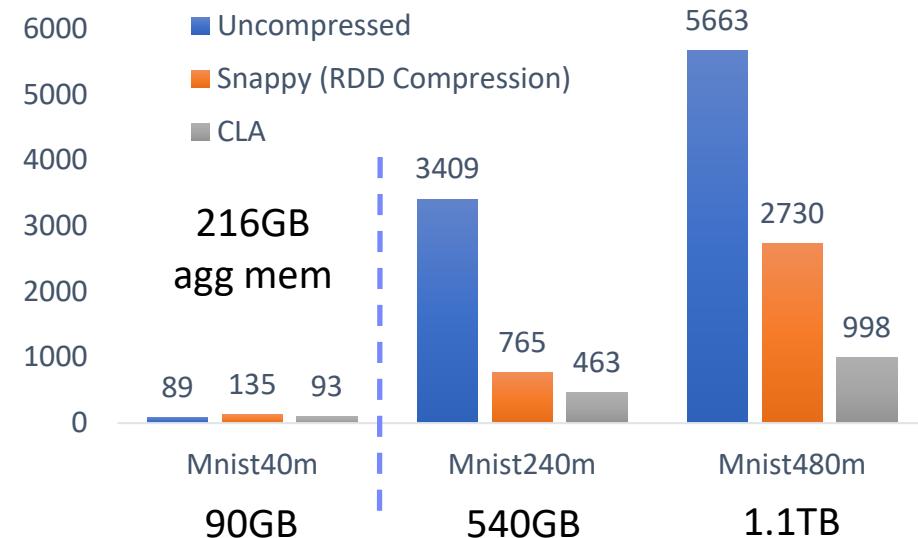
## ▪ Experiments

- LinregCG, 10 iterations, SystemML 0.14
- 1+6 node cluster, Spark 2.1

**Compression Ratios**

Dataset	Gzip	Snappy	CLA
Higgs	1.93	1.38	<b>2.17</b>
Census	17.11	6.04	<b>35.69</b>
Covtype	10.40	6.13	<b>18.19</b>
ImageNet	5.54	3.35	<b>7.34</b>
Mnist8m	4.12	2.60	<b>7.32</b>
Airline78	7.07	4.28	<b>7.44</b>

**End-to-End Performance [sec]**



# Backup: Index Structures

- **Overview Common Indexing Techniques**

- Physical blocking w/ leaf splitting strategies
  - Dense and sparse leaf blocks w/ contiguous ranges of cells
  - Batching of updates (deferred insertion)

- **LAB-Tree (Linearized Array B-tree, RIOT)** [PVLDB 2011]

- Operations: get, scan (iterator w/ given order), left/right indexing (on disk)
  - B-tree w/ physical blocking (sparse/dense), leaves have assigned ranges
  - Array linearization via UDFs (e.g., row/column major, Z-order, etc)
  - **Leaf splitting strategies:** split-in-middle, split-aligned, split-off-dense, split-defer-next, split-balanced-ratio
  - **Flushing policies for update batching:** flush-all, least-recently-used, smallest-page, largest-page, largest-page-probabilistically, largest-group

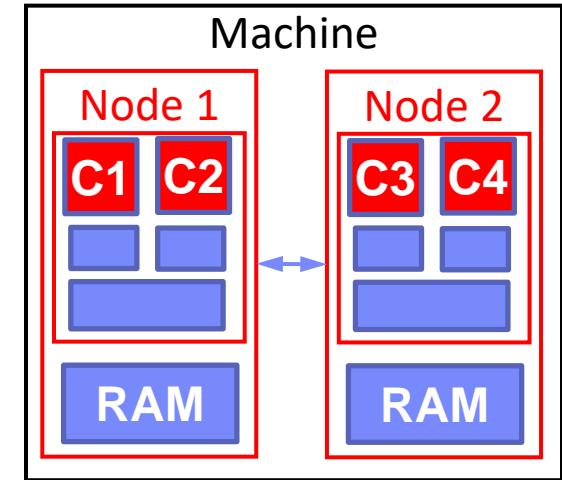
# Backup: Index Structures, cont.

- **AT MATRIX (Adaptive Tile Matrix, SAP HANA)** [ICDE 2016]
  - Operations: matrix multiplication ATMult (in-memory)
  - Two-level blocking: **Adaptive variable-sized tiles** (dense or sparse w/ CSR), composed of atomic squared blocks
  - Two-dimensional quad-tree, w/ Z-order as linearization function
  - **Horizontal partitioning across NUMA nodes**
  
- **TileDB Storage Manager** [PVLDB 2016]
  - Operations: init, write, read, consolidate, finalize (on disk)
  - Two-level blocking: space tiles (fixed size), data tiles (variable size for sparse)
  - Two-level linearization: cell order and tile order (row/column major)
  - **Fragments for update batching**  
("a timestamped snapshot of a batch of array updates")

# Backup: NUMA-Aware Partitioning and Replication

## ■ AT MATRIX (Adaptive Tile Matrix)

- Recursive NUMA-aware partitioning into dense/sparse tiles
- Inter-tile (worker teams) and intra-tile (threads in team) parallelization
- Job scheduling framework from SAP HANA (horizontal range partitioning, socket-local queues with task-stealing)



## ■ NUMA-Aware Model and Data Replication

- **DimmWitted: HW vs statistical efficiency**
- Model Replication
  - PerCore, PerMachine
  - PerNode (hybrid)
- Data Replication
  - Partitioning (sharding)
  - Full replication

➔ Open questions: Heterogenous hardware, cache coherence, etc.

# Backup: Buffer Pool Management

- **#1 Intermediates of LA Programs**
  - Hybrid runtime plans of in-memory and distributed operations
  - **Graceful eviction of intermediates at granularity of variables**
  - Example: SystemML
    - Soft references for in-memory matrices and broadcasts
    - LRU, FIFO buffer replacement strategies
- **#2 Operation/Algorithm-Specific Buffer Management**
  - Operations/algorithms over out-of-core matrices and factor graphs
  - **Page-level storage layout and buffer replacement policies**
  - Example #2a: RIOT
    - Chains of matrix multiplications
    - Operation-aware I/O schedules
  - Example #2b: Elementary
    - LR, CRF, LDA over out-of-core factor graphs
    - Materialization strategies and MRU/LFU buffer replacement

# Part 6: Resource Elasticity

*“Intelligence is the ability to adapt to change.”*

*Stephen Hawking (?)*



**Jun Yang**

Duke University  
Durham, NC, USA

**SIGMOD 2017**

# Rise of Cloud

- **Cluster computing for big data is easier than ever**
  - Clouds allow you to get a cluster on demand, and pay as you go
  - There is a growing ecosystem of platforms and tools for data analysis

## Challenges

- **Maddening array of “knobs”**
  - Hardware provisioning, software configuration, program tuning
- **“Elastic” environment**
  - Multi-tenant clusters, fluctuating markets, failures
  - Particularly hard for large-scale, long-running ML workloads

# Roadmap

- **Provisioning (& scheduling): what do I need (& when)?**
- **Recovery: what do I do when what I need fails?**
- **Working with markets**

☞ These problems are not limited to DB & ML workloads, but we shall see how DB & ML add twists

# Provisioning: Example Decisions

- **Given an ML program, what types of machines to acquire, and how many** *Cumulon [SIGMOD'13+follow-up]*
  - A bigger cluster may get results faster, but cost more
  - No perfect speedup, so big clusters may not give good cost/time trade-off
- **Given a cluster, how to configure the execution of an ML program**
  - What's the appropriate degree of parallelism *SystemML [DEBull'14,PVLDB'16]* for an execution step? *ScalOps [DeBull'12]*
    - Overhead of parallelism isn't always justified
  - How much memory do we allocate to master and work processes? *SystemML [SIGMOD'15]*
    - Optimal allocation depends on computation and data access patterns

👉 **Decisions interact with optimizations discussed earlier**

- Cluster configuration affects degree of parallelism and memory allocation, as well as optimal execution strategies

# Provisioning/Scheduling: Techniques

Depend on the level of abstraction:

- Program is a black box
  - First observe, and then decide; can leverage past execution profiles
- Program is broken down into a workflow with clear input/output for each unit, e.g., *MapReduce*, *Spark*
  - More effective profiling and optimization on a per-unit basis
- Program is specified declaratively, DB-style
  - Reusable and composable cost models
  - Bigger search space through rewrites
  - Cost-based what-if analysis
- Program follows a specific template
  - Even more opportunities arise; e.g., scheduling parameter updates/synchronization in *parameter servers* [**VLDB'10, OSDI'14**] + resource provisioning in *Dolphin* [**MLSys'16**] + adapting learning rate by update staleness in *DynSGD* [**SIGMOD'17**]

*SystemML* [**ICDE'11+follow-up**]

*Cumulon* [**SIGMOD'13+follow-up**]

☞ Adaptation is always key, regardless of abstraction level

# Recovery: General Techniques

Depend on the level of abstraction:

- **Program is a black box**
  - Checkpointing VM state in reliable/redundant storage
- **Program is a workflow with clear input/output for each unit**
  - Write input/output to reliable storage + rerun failed units, e.g.,  
*Hadoop/MapReduce*
  - Intermediate results can be in memory and lost + recover using lineage  
*Spark RDD* [NSDI'12]
- **Program is specified declaratively, DB-style**
  - Finer-grained lineage-based recovery using knowledge of operators + intelligent selective checkpointing  
*Cümülon* [VLDB'15]

# Recovery: Algorithm-Specific

- **Many ML algorithms can tolerate missing input or errors by design**
  - Instead of recovering to a state where as if failures never occurred, convert failures into “soft” ones that algorithms can handle themselves
- **Example: distributed batch gradient descent**

**Narayananmurthy+ (REEF) [BigLearn'13]**

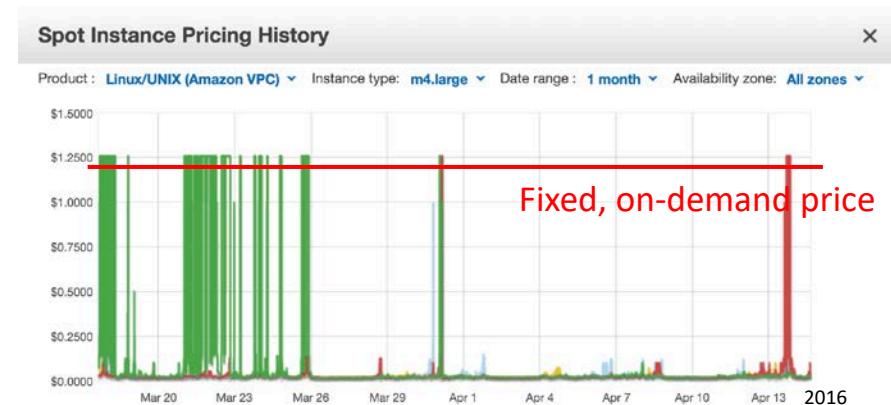
  - In an iteration, if a task fails to calculate the contribution from one partition of data, simply use an approximation (from the previous iteration)
  - Algorithm still converges

☞ **Generalized to user-defined, algorithm-specific “compensations”**

**Schelter+ [CIKM'13]**

# Working with Markets

- “On-demand” (regular) instances: fixed price, guaranteed
- “Spot” instances: availability/price vary over time; e.g; on Amazon:
  - You set a bid price, and get instances if bid price  $\geq$  market price
  - You pay market price (@hour start), by hours
  - You lose the instances if market price rises above your bid, but your last hour will be free
- Price can depend on machine type, region, and time



👉 How do we leverage markets effectively?

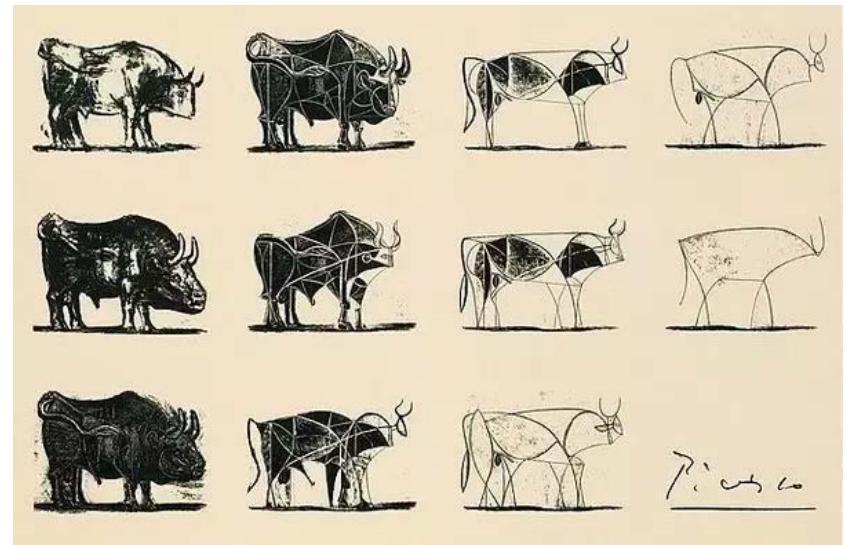
- Pop quiz: would you ever bid higher than the fixed price?
  - Yes! Less chance of losing them, yet still lower cost on average

# Working with Markets: Techniques

- **Diversify your portfolio: consider instances with different types, across regions**
  - If one market is too expensive, turn to others, e.g., *Dyna* [TCC'16]
  - A heterogeneous cluster may be best for mixed workloads, e.g., *Zhang+* [PER'15]
- **Minimizing expected cost is often not enough; need to control risk**
  - Model the market to quantify uncertainty, e.g., *Cümülon(-D)* [PVLDB'15,'17]
- **Zafer+ [Cloud'12] squeezes entire execution on spots in an hour; retries with a higher bid price if you lose them**
  - Losing spots within an hour incurs no cost with Amazon
- ***Dyna* [TCC'16] tries faster spots before falling back to on-demand**
  - But only if doing so improves the execution time distribution
- ***Cümülon* [PVLDB'15] picks the optimal mix of spot/on-demand instances**
  - To minimize expected cost while meeting deadline/budget with high probability
  - Recovers and re-optimizes if spots are lost
- ***Cümülon-D* [PVLDB'17] adapts proactively dynamically and proactively**
  - Bids/terminates as needed, based on execution progress and market condition
  - Solves the optimization problem as a Markov Decision Process (MDP) and pre-compiles a “cookbook” to apply at run time

# Summary

- Large-scale ML is increasingly being done in a cloud
- Challenges of elasticity are not unique to DB & ML
- Lots of uncertainty, but adaption & stochastic modeling come to rescue
- Different levels of abstraction lead to different opportunities—declarative (DB-style) ML enables smarter, more effective solutions



# References for Part 6: Resource Elasticity

- **Cumulon [SIGMOD'13]** Huang et al. "Cumulon: optimizing statistical data analysis in the cloud." SIGMOD 2013
- **Cümülon [PVLDB'15]** Huang et al. "Cümülon: matrix-based data analytics in the cloud with spot instances." PVLDB 2015
- **Cümülon-D [PVLDB'17]** Huang & Yang. "Cümülon-D: data analytics in a dynamic spot market." PVLDB 2017
- **Dolphin [MLSys'16]** Zhou et al. "Dolphin: Runtime Optimization for Distributed Machine Learning." ML Systems Workshop, 2016
- **Dyna [TCC'16]** Zhou et al. "Monetary cost optimizations for hosting workflow-as-a-service in IaaS clouds." TCC 4(1), 2016
- **DynSGD [SIGMOD'17]** Jiang et al. "Heterogeneity-aware Distributed Parameter Servers." SIGMOD 2017
- **Narayananmurthy+ (REEF) [BigLearn'13]** Narayananmurthy et al. "Towards Resource-Elastic Machine Learning." BigLearn 2013
- **Parameter Server [VLDB'10]** Smola & Narayananmurthy. "An architecture for parallel topic models." VLDB 2010
- **Parameter Server [OSDI'14]** Li et al. "Scaling Distributed Machine Learning with the Parameter Server." OSDI 2014
- **Schelter+ [CIKM'13]** Schelter et al. "All Roads Lead to Rome: Optimistic Recovery for Distributed Iterative Data Processing." CIKM 2013
- **ScalOps [DeBull'12]** Borkar et al. "Declarative systems for large-scale machine learning." IEEE Data Eng. Bulletin, 35(2), 2012
- **Spark RDD [NSDI'12]** Zaharia et al. "Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing ." NSDI 2017
- **SystemML [ICDE'11]** Ghoting et al. "SystemML: Declarative machine learning on MapReduce." ICDE 2011
- **SystemML [SIGMOD'15]** Huang et al. "Resource elasticity for large-scale machine learning." SIGMOD 2015
- **SystemML [VLDB'16]** Boehm et al. "SystemML: Declarative machine learning on Spark." PVLDB 9(13), 2016
- **Zafer+ [Cloud'12]** Zafer et al. "Optimal bids for spot VMs in a cloud for deadline constrained jobs." Cloud Computing, 2012
- **Zhang+ [PER'15]** Zhang et al. "Exploiting Cloud Heterogeneity to Optimize Performance and Cost of MapReduce Processing." Performance Evaluation Review, 42(4), 2015

# Part 7: ML Lifecycle Systems

Arun Kumar  
UC San Diego  
La Jolla, CA, USA

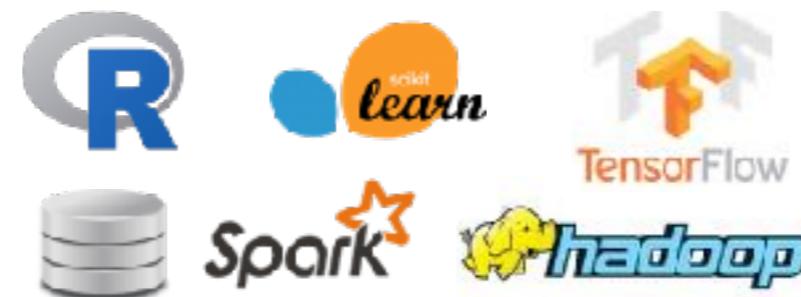
SIGMOD 2017

# Overview: ML Lifecycle Issues



Data Scientist/  
ML Engineer

Build → Deploy



Data sourcing  
**Feature engineering**  
and model selection

Model serving  
Tighter loop between  
inference and learning

Model management

# Feature Engineering

*Q: What is feature engineering (FE)?*

The process of obtaining a formal representation of the data-generating process as structured signals (features) for an ML model

*Q: Why is it important from a data management perspective?*

High-quality features are the “secret sauce” of applied ML

FE operations are basically data transformations!

Often “brushed under the carpet” by ML community

*Q: What sort of operations constitute feature engineering?*

Depends on the data type!

Structured data: Whitening, feature selection/ranking, joins, PCA, etc.

Text: Bag-of-words, Parsing-based, Domain-specific, Word2Vec, etc.

Deep CNNs and RNNs for images, audio, video, time series, etc.

# Feature Engineering Systems

## Feature selection:

Obtain a subset of features to improve accuracy and/or interpretability

## Columbus [SIGMOD'14]:

Often not a single algorithm but a human-in-the-loop dialogue process

Data scientist explores multiple subsets based on domain insights

Understanding  
customer churn

CustID	Churn?	Age	Income	Gender	City	...
...	...	...	...	...	...	...



*Evaluate error with all features in chosen set  
Drop demographic features and re-evaluate  
Add Gender back in and so on ...*

A few such common steps encoded as “declarative” ops in DSL

Impl. on top of R/Python; optimizing code-gen middleware

Batching/materialization; QR decomposition; coresets; warm start

# Feature Engineering Systems

Treating FE as a dataflow-oriented process; DB-style optimizations:

**Brainwash** [CIDR'13] / **DeepDive** [DataEng'14]

Workflows of UDFs; feature recommendations

**KeystoneML** [ICDE'17]

Alternative phy. impl. of solvers; cost-based op. selection

Reducing amount of work for feature coding/evaluation:

**Zombie** [ICDE'16]

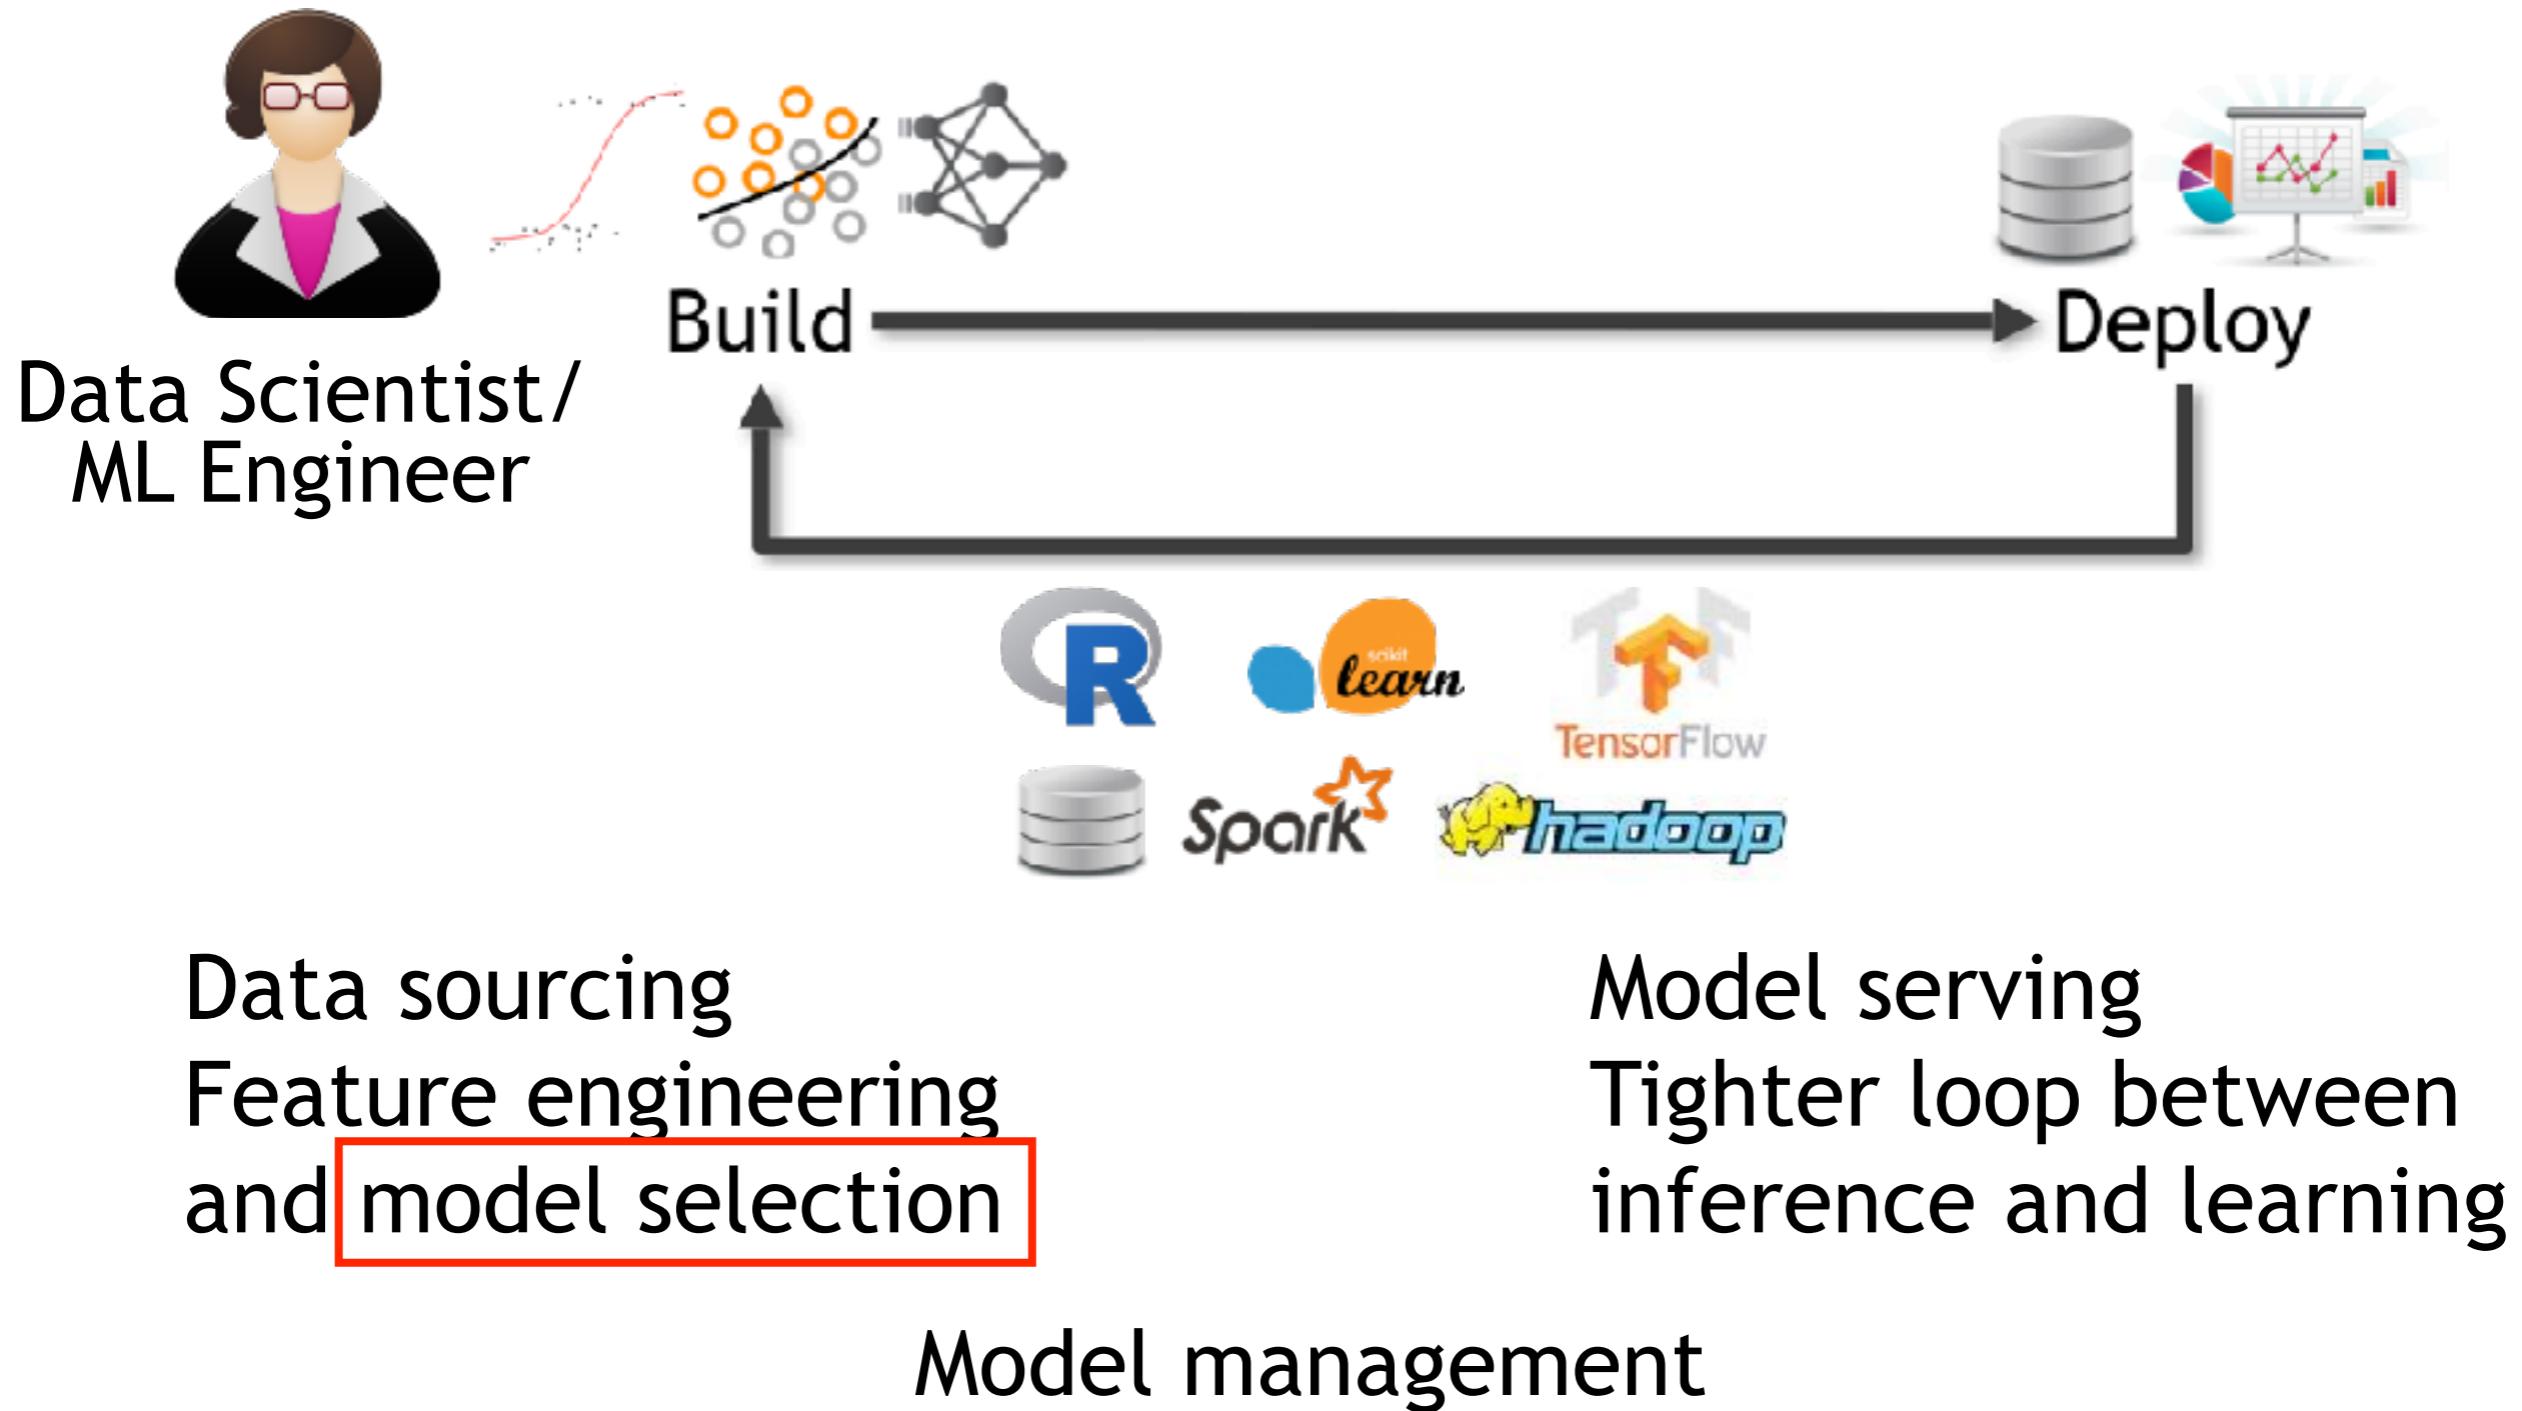
Index structure to sub select relevant data; bandit techniques

Applying learning theory to skip features and help with sourcing tables:

**Hamlet** [SIGMOD'16]

More open questions remain in *systematizing* feature engineering

# Overview: ML Lifecycle Issues



# Model Selection

*Q: What is model selection (MS)?*

The process of obtaining a prediction function to capture a data-generating process using data generated by that process

**MSMS [SIGMODRec'15]**  
Model Selection Triple (MST)  
(FE, AS, PT)

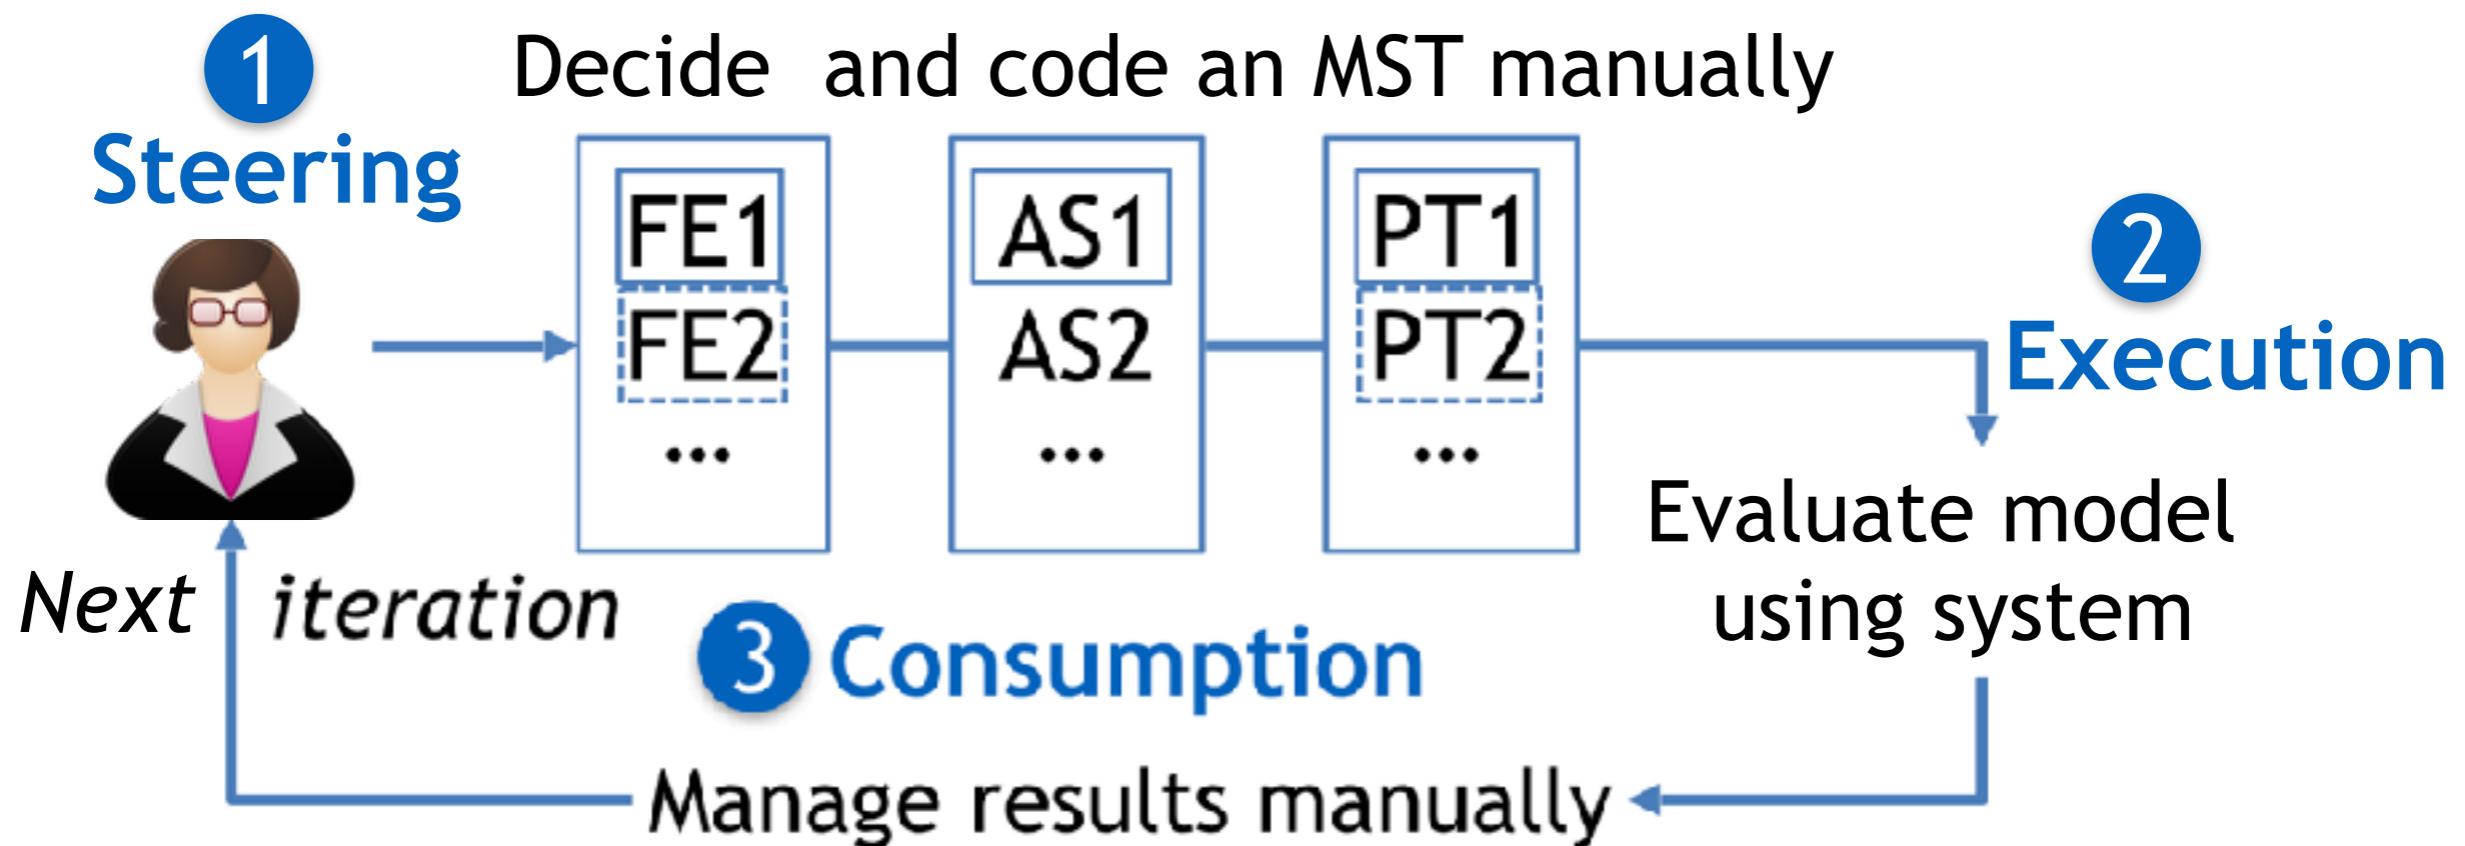
FE: Feature Engineering  
AS: Algorithm Selection  
PT: (Hyper-)Parameter Tuning

*Q: Why is it important from a data management perspective?*

FE, AS, and PT often access the dataset (or subsets) repeatedly  
A lot of opportunities to improve efficiency with DB-style opt.  
FE, AS, and PT are *inter-dependent* and together constitute MS

# Model Selection Process

MSMS [SIGMODRec'15] Model Selection Triple (MST): (FE, AS, PT)



Data scientists typically think at a higher level of abstraction

Automation essentially groups MSTs en masse

*MS abstractions can help capture intermediate points*

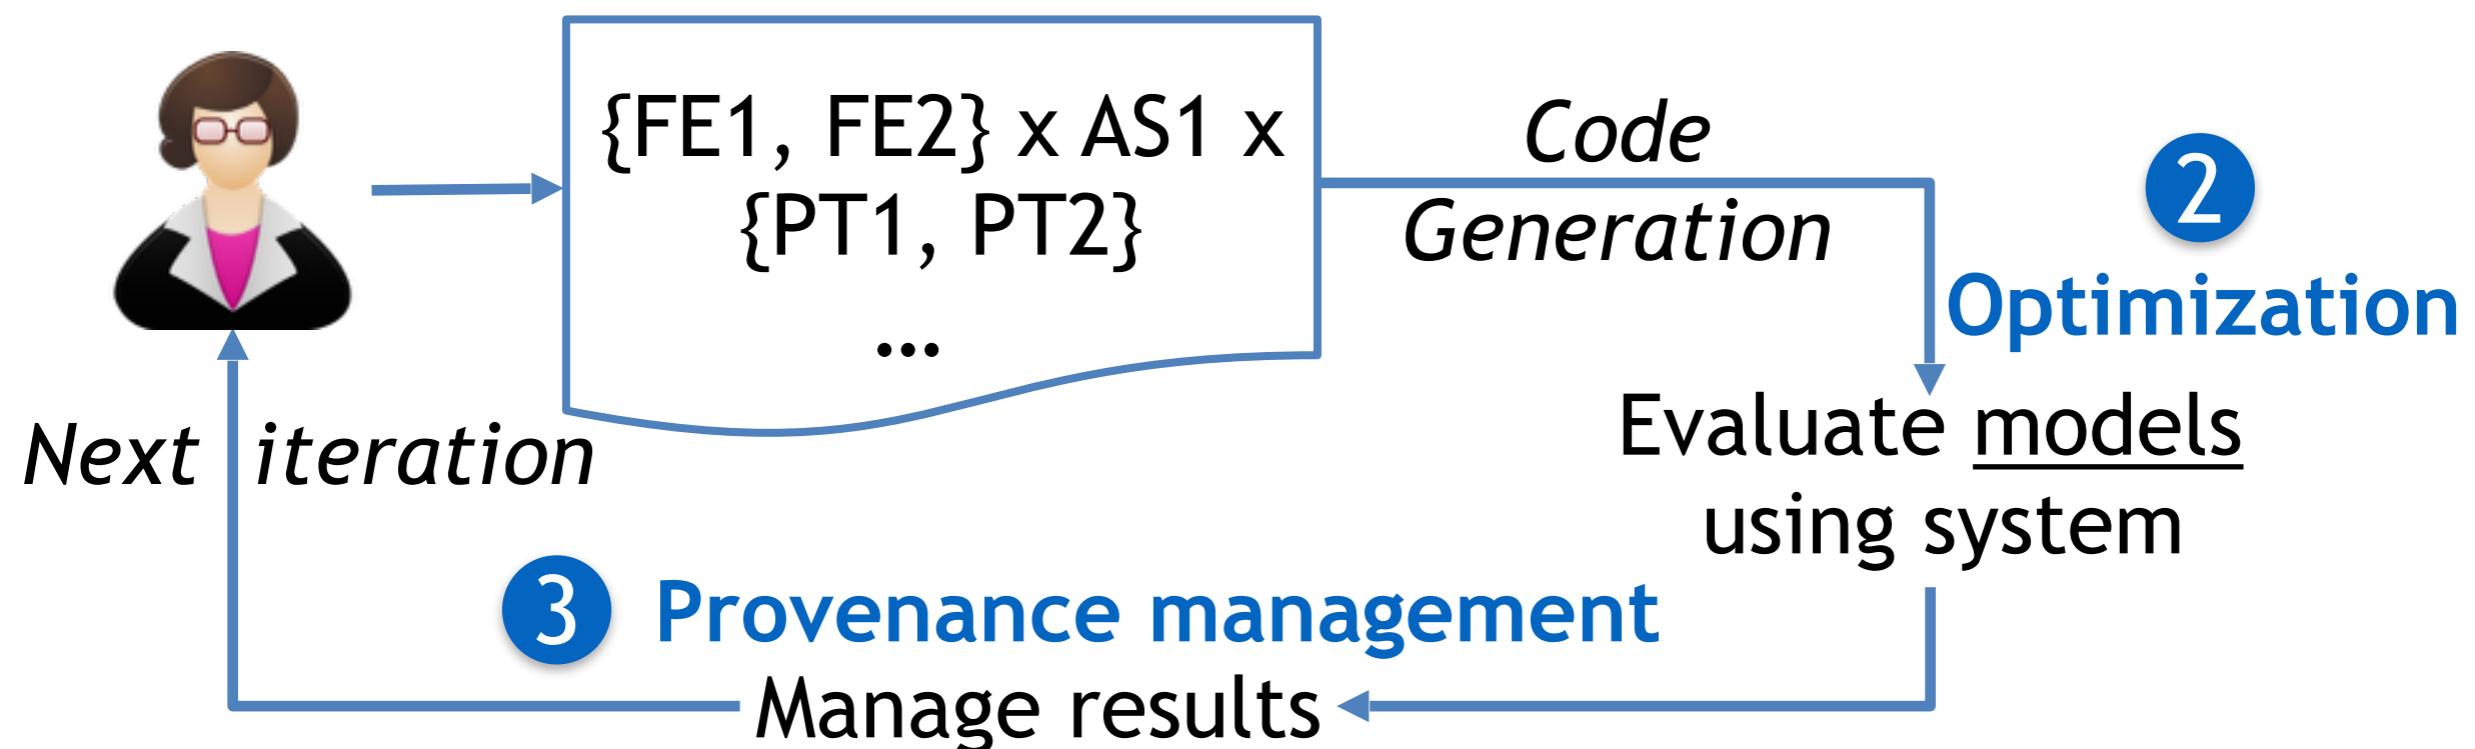
# Model Selection Process

MSMS [SIGMODRec'15]

Model Selection Triple (MST): (FE, AS, PT)

## 1 “Declarative” interfaces

Group a set of “logically related” MSTs



Many old and recent MS abstractions can be “retro-fitted”

Several new MS abstractions can be introduced to co-exist

# Model Selection Management Systems (MSMS)

MSMS [SIGMODRec'15]

The Higher Layers: Declarative Interfaces (some in hindsight!)

Autotuned  
functions

E.g., `glmnet()` in R  
 $\text{FE} \times \text{AS} \times \{\text{PT}\}$

Columbus

E.g., `StepAdd()`  
 $\{\text{FE}\} \times \text{AS} \times \text{PT}$

MLBase

E.g., `doClassify()`  
 $\text{FE} \times \{\text{AS1} \times \{\text{PT}\},$   
 $\text{AS2} \times \{\text{PT}\}\}$

New Abstractions

...  
 $\{\text{FE}\} \times \{\text{AS} \times \text{PT}\}, ...$

$\{ \{\text{FE}\} \times \{\text{AS}\} \times \{\text{PT}\} \}$

The Narrow Waist:  
A set of logically related  
Model Selection Triples (MST)



In-memory



In-RDBMS



Others

The Lower Layers: Optimized Implementations

# Model Selection Systems

Automation of AS and PT search with pre-defined search space:

**MLbase** [CIDR'13] / **TuPAQ** [SoCC'15]

Declarative ML tasks (e.g., “DoClassify”); fixed set of algorithms

Data batching; bandit techniques for explore-exploit search

**Hemingway** [MLSys'16]

Joint AS and cluster sizing for optimization algorithms

Observe-and-adapt approach for convergence properties

DB-style optimizations for PT and general meta-learning:

**SystemML** [ICDE'15]; **GLADE** [DanaC'12]

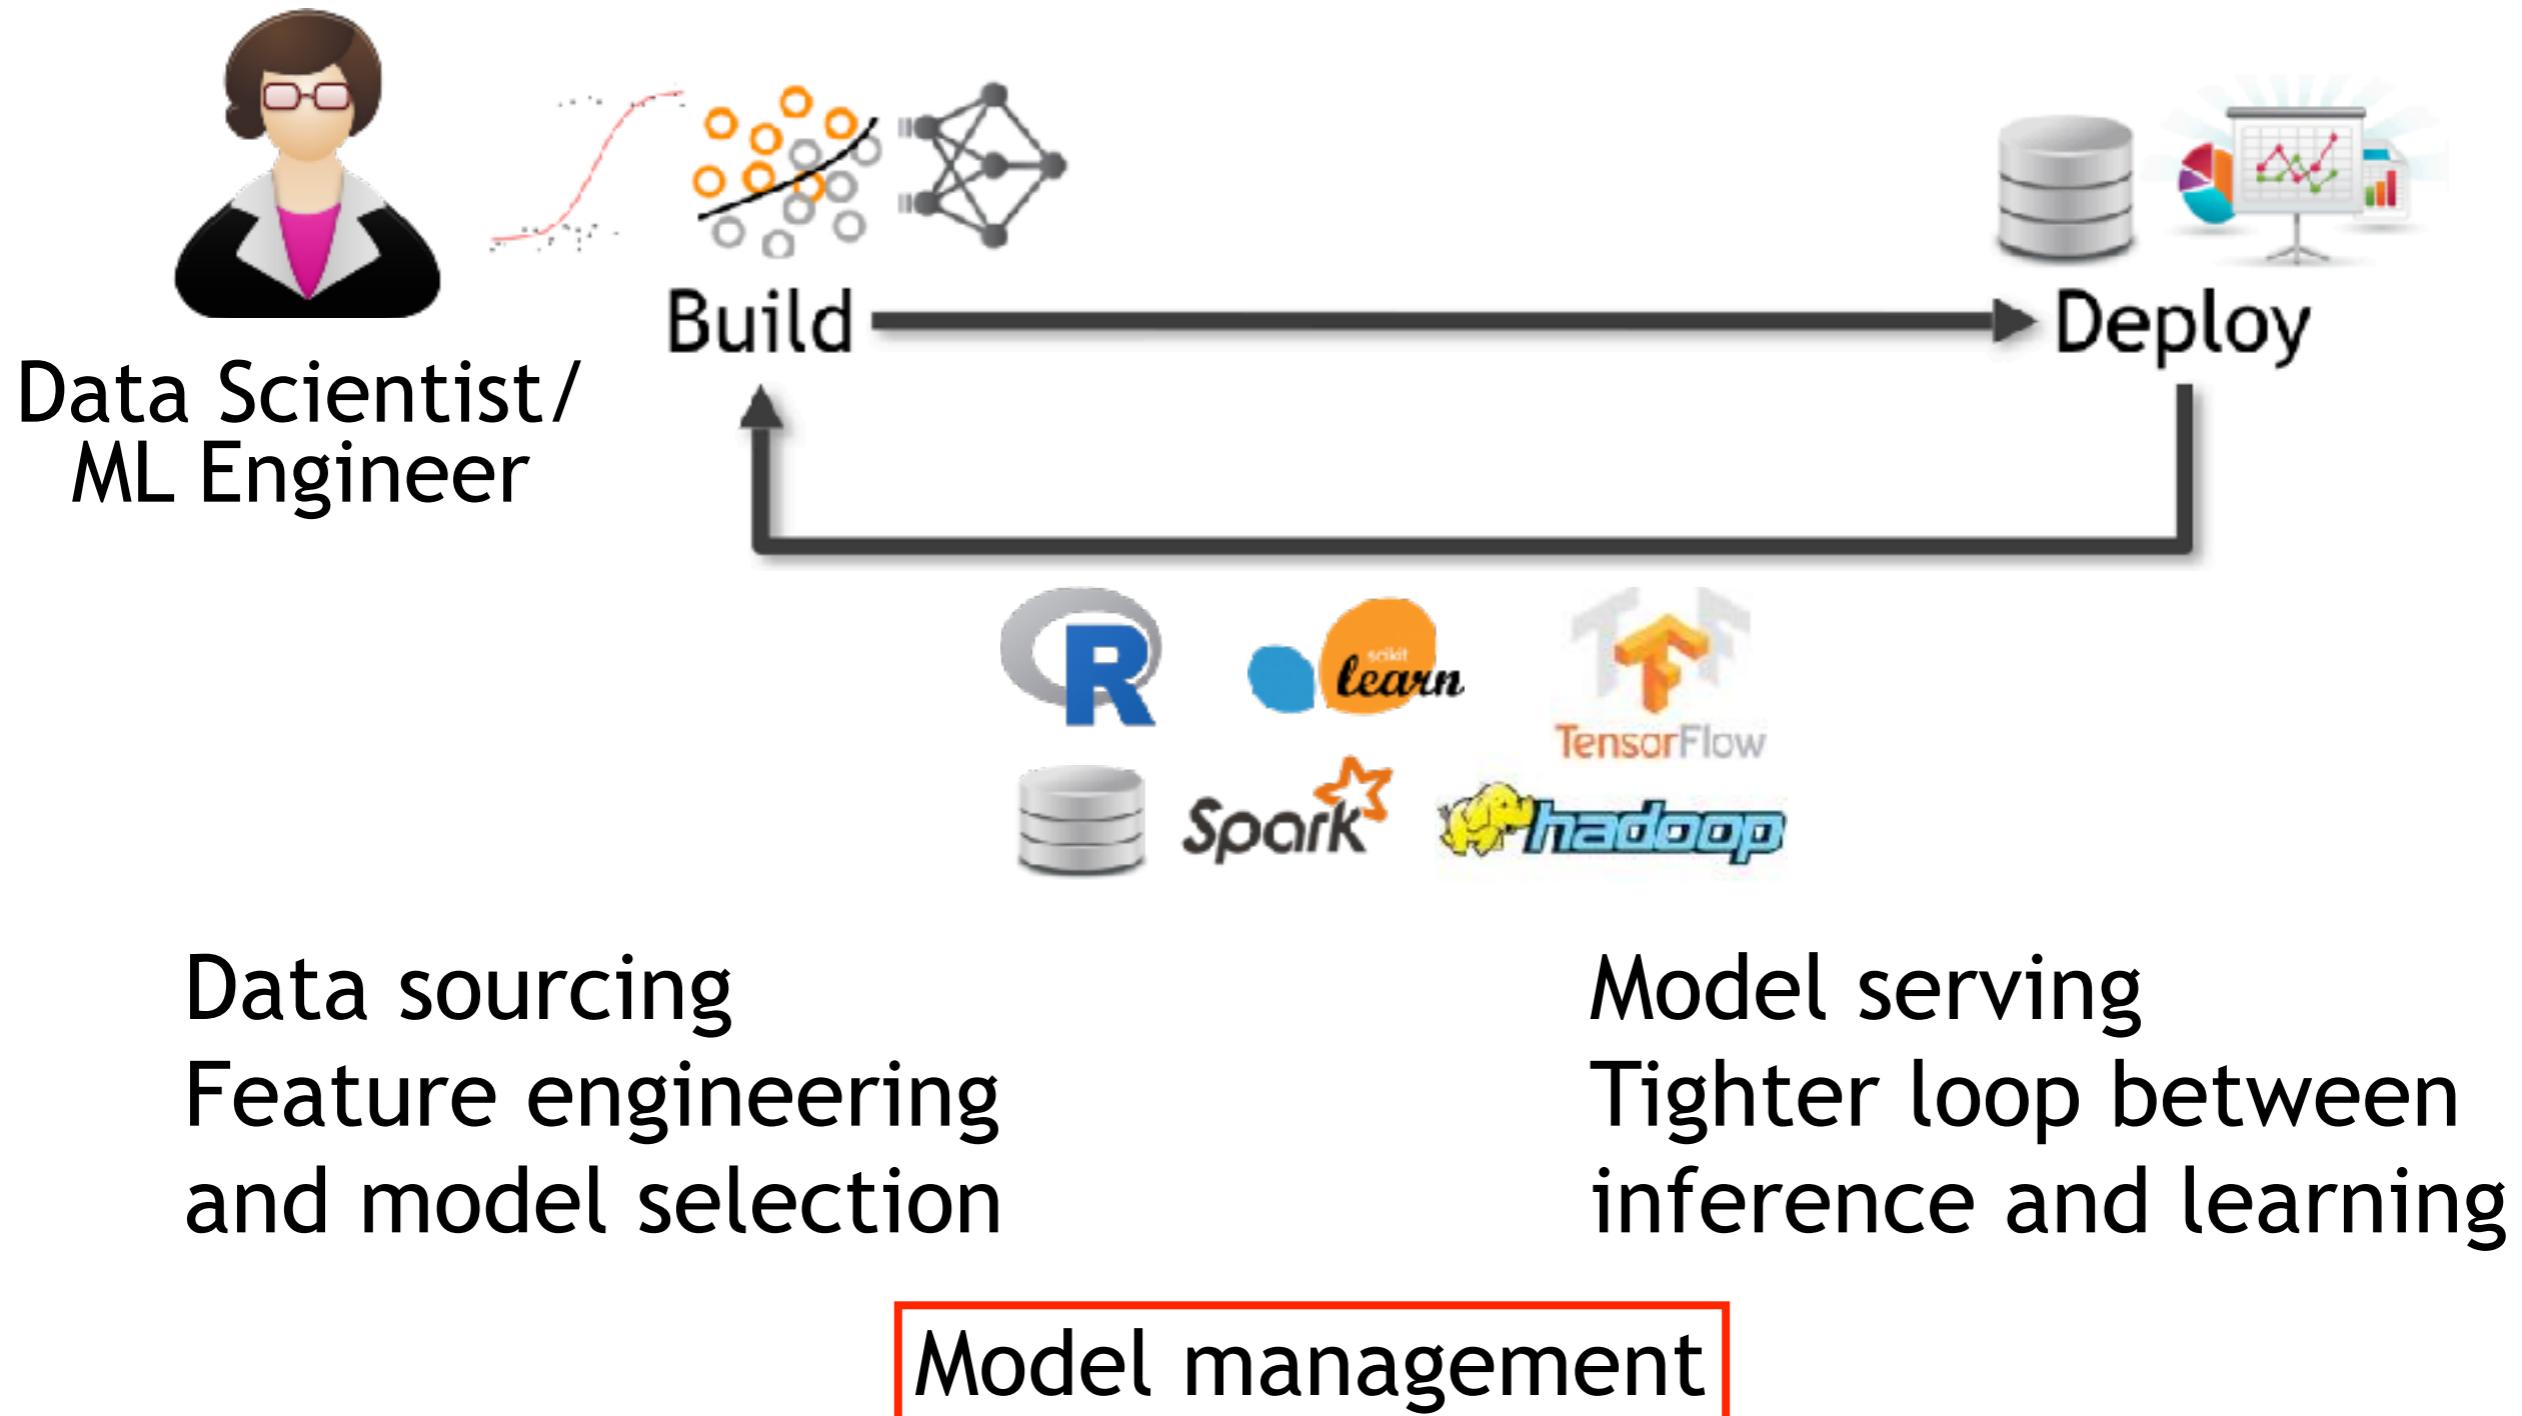
Many open questions remain on optimizing/improving model selection

Interactions of PT with AS and FE

Exploiting redundancy across and within MSTs; cost models

Incorporating constraint/approximations and visualizations, etc.

# Overview: ML Lifecycle Issues



# Model Management Systems

*Q: What is model management?*

Treating trained models as data themselves (store, query, debug, etc.)

Integrating ML models with SQL querying: **LongView** [CIDR'11]

Iterative ML debugging: **MindTagger** [VLDB'15], **PALM** [HILDA'17]

Specialized storage engines and custom optimizations:

**ModelHub** [ICDE'17]

Versioned storage/retrieval of CNNs (sets of float matrices)

Optimizations for reducing storage footprint

Many open questions on managing large space of MSTs, especially for large models (DNNs/trees); ML provenance and debugging

# Other ML Lifecycle Issues

**Model Serving:** High-throughput/low-latency inference/re-learning

**MacroBase** [SIGMOD'17]

**Clipper** [NSDI'17] / **Velox** [CIDR'15]

Integrating data-driven applications with reinforcement learning

**Data Sourcing:** Modeling labeling process; ML+cleaning; ML+pricing

**Snorkel** [NIPS'16]

**ActiveClean** [VLDB'16]

**Model-Based Pricing** [DEEM'17]

**Interactive Model Building:** Human-in-the-loop interfaces

**Ava** [CIDR'17]

**Vizdom** [VLDB'15]

# References: Part 7

- ActiveClean [VLDB'16]: ActiveClean: Interactive Data Cleaning For Statistical Modeling
- Ava [CIDR'17]: Ava: From Data to Insights Through Conversation
- Brainwash [CIDR'13]: Brainwash: A Data System for Feature Engineering
- Clipper [NSDI'17]: Clipper: A Low-Latency Online Prediction Serving System
- Hamlet [SIGMOD'16]: To Join or Not to Join? Thinking Twice about Joins before Feature Selection
- Hemingway [MLSys'16]: Hemingway: Modeling Distributed Optimization Algorithms
- KeystoneML [ICDE'17]: KeystoneML: Optimizing Pipelines for Large-Scale Advanced Analytics
- Longview [CIDR'11]: The Case for Predictive Database Systems: Opportunities and Challenges
- MacroBase [SIGMOD'17]: MacroBase: Prioritizing Attention in Fast Data
- MindTagger [VLDB'15]: MindTagger: A Demonstration of Data Labeling in Knowledge Base Construction
- MLbase [CIDR'13]: MLbase: A Distributed Machine-learning System
- Model-Based Pricing [DEEM'17]: Model-based Pricing: Do Not Pay for More than What You Learn!
- ModelDB [HILDA'16]: MODELDB: A System for Machine Learning Model Management
- ModelHub [ICDE'17]: ModelHub: Towards Unified Data and Lifecycle Management for Deep Learning
- MSMS [SIGMODRec'15]: Model Selection Management Systems: The Next Frontier of Advanced Analytics
- PALM [HILDA'17]: PALM: Machine Learning Explanations For Iterative Debugging
- Snorkel [NIPS'16]: Data Programming: Creating Large Training Sets, Quickly
- SystemML [ICDE'15]: Efficient Sample Generation for Scalable Meta Learning
- TuPAQ [SoCC'15]: Automating Model Search for Large Scale Machine Learning
- Velox [CIDR'15]: The Missing Piece in Complex Analytics: Low Latency, Scalable Model Management and Serving with Velox
- Vizdom [VLDB'15]: Vizdom: Interactive Analytics through Pen and Touch
- Zombie [ICDE'16]: Input Selection for Fast Feature Engineering

# Part 8: Open Problems and Conclusions

**Arun Kumar**

UC San Diego  
San Diego, CA, USA

**Matthias Boehm**

IBM Research – Almaden  
San Jose, CA, USA

**Jun Yang**

Duke University  
Durham, NC, USA

**SIGMOD 2017**

# Open Problems: Optimizer and Runtime

- #1 **Size and Sparsity Estimation**
  - Fundamental building block for cost comparisons / valid plan generation
  - Issues: function calls, UDFs, data-dependent operators, changing sizes
- #2 **Convergence Estimation**
  - Number of iterations until convergence unknown
  - Required for cost comparisons and progress estimation
- #3 **Adaptive Query Processing and Storage**
  - Unknown or changing workloads → adaptive query processing
  - Currently limited to inter-DAG recompilation and expression optimization
- #4 **Automatic Rewrites and Operator Fusion**
  - Huge potential for simplification rewrites and operator fusion
  - Challenging in presence of new access methods, compression, etc.
- #5 **Special Value Handling**
  - Special values such as NaN, INF, -0 ignored by most systems → **incorrect results**
  - Support these special values w/o sacrificing performance

# Open Problems: End-to-End Lifecycle

- #6 Integrating Relational and Linear Algebra
  - Seamless optimizer / runtime integration in holistic framework
  - Including data transformations, training and prediction
- #7 Seamless Feature Engineering and Model Selection
  - (Semi-)automating feature engineering and model selection
  - Including abstractions, meta-algorithms, and system architectures
- #8 ML System Benchmarks
  - Existing benchmarks limited to ML tasks in terms of reference implementations of large-scale ML libraries or SQL-centric workloads
  - Broader range of benchmarks at various abstraction levels

# Conclusions

- **Summary**
  - Compelling arguments for integrating **ML → DB** and **DB → ML**
  - ML in data systems, DB-inspired ML systems, ML lifecycle systems
- **#1 Existing Work to Build Upon**
  - Awareness of existing systems and techniques
  - Survey of effective optimization and runtime techniques
- **#2 Where the Data Management Community Can Help**
  - Integrating ML into existing data systems
  - Optimizer and runtime techniques for large-scale ML systems
  - Tools and systems to simplify/improve the end-to-end ML lifecycle

**➔ Many open technical problems**