

Distributed Query Processing (Chapters 7, 8, 9)

DISTRIBUTED QUERY PROCESSING PROBLEM

Besides the choice of ordering relational algebra operations, **we must also select the best sites to process data.**

Example: `SELECT *`
`FROM E, G`
`WHERE E.ENO=G.ENO AND`
`RESP='Manager'`

`G` or `ASG (ENO, PNO, RESP, DUR)`
`E` or `EMP(ENO, ENAME, TITLE)`

Bad algebra:

$$\sigma_{RESP='Manager' \wedge E.ENO=G.ENO} (E \times G)$$

Good algebra:

$$E \bowtie_{ENO} (\sigma_{RESP='Manager'}(G))$$

$$EMP \triangleright \triangleleft_{ENO} (\sigma_{RESP='Manager'}(ASG))$$


Fragmentation Schema

$$\begin{aligned} EMP_1 &= \sigma_{ENO \leq "E3"}(EMP) \\ EMP_2 &= \sigma_{ENO > "E3"}(EMP) \\ ASG_1 &= \sigma_{ENO \leq "E3"}(EMP_1) \\ ASG_2 &= \sigma_{ENO > "E3"}(EMP_2) \end{aligned}$$


Allocation Schema

Site 1: ASG_1
 Site 2: ASG_2
 Site 3: EMP_1
 Site 4: EMP_2

Query Plans




(a) Strategy A




(b) Strategy B

Regular Join



Semijoin: $R \triangleright \triangleleft_A S \Leftrightarrow (R_{SJ_A} S) \triangleright \triangleleft_A S$



Typically $|S'| + |R'| \ll |R|$

More processing overhead to do a join using semi-join

Operation	Complexity
Select Project (without duplicate elimination)	$O(n)$
Project (with duplicate elimination) Group	$O(n \log n)$
Join Semijoin Division Set Operators	$O(n \log n)$
Cartesian Product	$O(n^2)$

Figure 7.2. Complexity of Relational Algebra Operations

Assume that tuples of each relation must be sorted.

n denotes the relation cardinality.

PRINCIPLES:

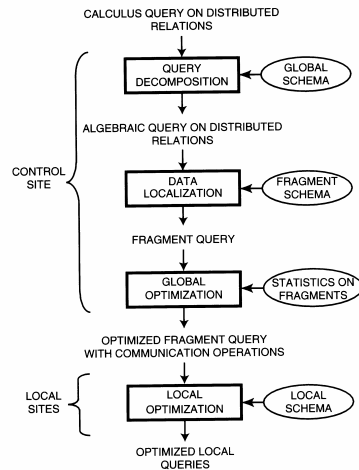
- The most selective operations that reduce cardinalities should be performed first.
- Operations should be ordered by increasing complexity so that expensive operations can be avoided or delayed.

OBJECTIVES OF QUERY OPTIMIZATION

1. Good measurement parameters:
 1. **Total cost**: The sum of all time incurred in processing the operations at various sites and in intersite communication.
 2. **Response Time**: The time elapsed for executing the query.
2. **Cost = f(CPU_cost, IO_cost, Communication_Cost)**
 - Minimize CPU_cost: use less expensive strategies
 - Minimize IO_cost: use fast access methods and good buffer management strategies
 - Minimize communication_cost: choose the execution sites intelligently.

Very slow communication networks:
Minimize communication costs generally at the expense of local processing.

Distributed Query Processing Steps



Query Decomposition:

- Normalization
- Semantically analyze the normalized query to eliminate incorrect queries.
- Simplify the correct query by removing redundant predicates.
- Restructure the algebraic query into a "better" algebraic specification.

This step is the same as standalone DBMS.

Distributed Query Processing Steps

Data Localization:

- Determine which fragments are involved and transform the global query into fragment queries.

Global Query Optimization:

- Find the "best" ordering of fragment queries and specifies the communication operations.

Local Query Optimization:

- Each site determines the access methods for the local fragment queries using the local schema.

Query Decomposition and Data Localization

Query Decomposition

- Normalization
- Semantic Analysis
- Simplification
- Rewriting

Data Localization: Reduction techniques for different types of fragmentation.

- Horizontal
- Vertical
- Derived
- Hybrid

Normalization

- The input query may be arbitrary complex.
- Normalization transforms the query into a normalized form to facilitate further processing.

Two possible normal forms:

Conjunctive Normal Form:

$$(p_{11} \vee p_{12} \vee \dots \vee p_{1n}) \wedge (p_{m1} \vee p_{m2} \vee \dots \vee p_{mn})$$

Disjunctive Normal Form:

$$(p_{11} \wedge p_{12} \wedge \dots \wedge p_{1n}) \vee (p_{m1} \wedge p_{m2} \wedge \dots \wedge p_{mn})$$

where p_{ij} is a simple predicate.

Simple Predicates

Given a relation $R(A_1, A_2, \dots, A_n)$ where A_i has domain D_i , a simple predicate p_j defined on R has the form

$p_j: A_i \theta \text{ Value}$
where $\theta \in \{=, <, \neq, \leq, >, \geq\}$ and $\text{Value} \in D_i$

Example:

J	JNO	JNAME	BUDGET	LOC
	J1	Instrumental	150,000	Montreal
	J2	Database Dev.	135,000	New York
	J3	CAD/CAM	250,000	New York
	J4	Maintenance	350,000	Orlando

Simple predicates:

$p_1: \text{JNAME} = \text{'Maintenance'}$

$p_2: \text{BUDGET} = \text{'200,000'}$

Examples:

$E(\underline{\text{ENO}}, \text{ENAME}, \text{TITLE})$
 $G(\underline{\text{ENO}}, \underline{\text{JNO}}, \text{RESP}, \text{DUR})$

Query: SELECT ENAME
FROM E, G
WHERE E.ENO=G.ENO AND G.JNO='J1' AND
(DUR=12 OR DUR=24) } Query qualification

Qualification in conjunctive normal form:

$(E.ENO = G.ENO) \wedge (G.JNO = \text{'J1'}) \wedge (DUR = 12 \vee DUR = 24)$

Qualification in disjunctive normal form:

$(E.ENO = G.ENO \wedge G.JNO = \text{'J1'} \wedge DUR = 12) \vee$
 $(E.ENO = G.ENO \wedge G.JNO = \text{'J1'} \wedge DUR = 24)$

↑
Replicated join

↑
Replicated select

Disjunctive normal form
can lead to replicated join
and select predicates.

Semantic Analysis

Semantic analysis enables rejection of incorrect queries.

- Type checking detects type incorrect problems.

```
SELECT E#
FROM E
WHERE ENAME>200;
```

← Undefined attribute!



← Incompatible operation!

- Query graph:** Determine the semantic correctness of a conjunctive multivariable (multi-table) query without negation.

```
SELECT ENAME, RESP
FROM EMP,ASG,PROJ
WHERE EMP.ENO=ASG.ENO AND PNAME='CAD/CAM' AND
DUR>=36 AND TITLE='Programmer'
```

```
SELECT ENAME, RESP
FROM EMP,ASG,PROJ
WHERE EMP.ENO=ASG.ENO AND
PNAME='CAD/CAM' AND DUR>=36 AND
TITLE='Programmer'
```

```
EMP(ENO, ENAME, TITLE)
AGG(ENO, JNO, RESP, DUR)
PROJ(PNO, PNAME, LOC)
```



Systems should reject queries with the unconnected join graphs.




Figure 8.2. Disconnected Query Graph

Elimination of Redundancy

Redundant predicates are likely to arise when a query is the result of system transformations applied to the user query.

Such redundancy and thus redundant work may be eliminated by simplifying the qualification using well-known idempotency rules.

Example:

$$\begin{aligned} & (\neg p_1 \wedge (p_1 \vee p_2) \wedge \neg p_2) \vee p_3 \\ &= (\neg p_1 \wedge p_1 \wedge \neg p_2) \vee (\neg p_1 \wedge p_2 \wedge \neg p_2) \vee p_3 \\ &= (F \wedge \neg p_2) \vee (\neg p_1 \wedge F) \vee p_3 \\ &= F \vee F \vee p_3 = p_3 \end{aligned}$$

Rewriting

Two Steps:

1. Transforming the query into an algebraic relational query tree.
2. Restructuring the algebraic tree to improve performance