

Tensor programs

Part I' + Part II'

Eugene Golikov

April 9, 2021

École Polytechnique Fédérale de Lausanne, Switzerland

Former researcher at DeepPavlov.ai, Moscow Institute of Physics and Technology, Russia

Define:

width of a network = the minimal number of nodes in its hidden representations.

Talk subject: neural nets in the limit of infinite width.

Reason:

1. Sufficiently wide nets \approx infinitely wide nets;
2. Infinitely wide nets are much easier to study theoretically;
3. Infinitely wide nets enjoy a number of cool properties (below).

Given He initialization (standard) and certain parameterization, a fully-connected feedforward network w/o shared weights enjoy the following properties:

1. It converges to a **Gaussian process** at initialization as width $\rightarrow \infty$ [Matthews et al., 2018];
2. Its GD training dynamics converges to a **kernel GD with a constant kernel** as width $\rightarrow \infty$ [Jacot et al., 2018];
3. The spectrum of its **input-output jacobian** can be computed in the limit of infinite width using a **free independence principle** [Pennington et al., 2017].

Tensor programs series [Yang, 2019, Yang, 2020a, Yang, 2020b] prove these properties for a wide class of models as follows:

1. Introduce a wide class of models called *tensor programs*;
2. Prove a *Master theorem* about their limit behavior;
3. Deduce the properties above from the Master theorem.

Overall talk construction strategy:

1. Take one of the properties discussed above;
2. Illustrate it on a simple model;
3. Introduce a class of tensor programs sufficient to express this property;
4. Prove the corresponding Master theorem;
5. Deduce the property from the Master theorem;
6. Proceed with another property.

Convergence to Gaussian processes

Consider a forward pass:

$$\underbrace{h^l = W^l x^{l-1}}_{\text{pre-activations, } \in \mathbb{R}^{n_l}}, \quad \underbrace{x^{l-1} = \phi(h^{l-1})}_{\text{activations, } \in \mathbb{R}^{n_{l-1}}}; \quad W^l \sim \mathcal{N}\left(0, \frac{\sigma_W^2}{n_{l-1}}\right). \quad (1)$$

In a matrix form:

$$\begin{pmatrix} h_1^l \\ h_2^l \\ \dots \\ h_{n_l}^l \end{pmatrix} = \underbrace{\begin{pmatrix} W_{11}^l & W_{12}^l & \dots & W_{1n_{l-1}}^l \\ W_{21}^l & W_{22}^l & \dots & W_{2n_{l-1}}^l \\ \dots & \dots & \dots & \dots \\ W_{n_l1}^l & W_{n_l2}^l & \dots & W_{n_l n_{l-1}}^l \end{pmatrix}}_{\text{all iid } \sim \mathcal{N}(0, \sigma_W^2/n_{l-1})} \times \begin{pmatrix} x_1^{l-1} \\ x_2^{l-1} \\ \dots \\ x_{n_{l-1}}^{l-1} \end{pmatrix}.$$

$$\underbrace{\begin{pmatrix} h_1^I \\ h_2^I \\ \dots \\ h_{n_I}^I \end{pmatrix}}_{\substack{\text{components tend to Gaussians} \\ \text{as } n_{I-1} \rightarrow \infty \text{ by CLT}}} = \underbrace{\begin{pmatrix} W_{11}^I & W_{12}^I & \dots & W_{1n_{I-1}}^I \\ W_{21}^I & W_{22}^I & \dots & W_{2n_{I-1}}^I \\ \dots & \dots & \dots & \dots \\ W_{n_I1}^I & W_{n_I2}^I & \dots & W_{n_In_{I-1}}^I \end{pmatrix}}_{\text{all iid } \sim \mathcal{N}(0, \sigma_W^2/n_{I-1})} \times \underbrace{\begin{pmatrix} x_1^{I-1} \\ x_2^{I-1} \\ \dots \\ x_{n_{I-1}}^{I-1} \end{pmatrix}}_{\text{suppose all components are iid}}.$$

$$\underbrace{\begin{pmatrix} h_1^I \\ h_2^I \\ \vdots \\ h_{n_I}^I \end{pmatrix}}_{\text{different components are uncorrelated}} = \underbrace{\begin{pmatrix} W_{11}^I & W_{12}^I & \cdots & W_{1n_I-1}^I \\ W_{21}^I & W_{22}^I & \cdots & W_{2n_I-1}^I \\ \vdots & \vdots & \ddots & \vdots \\ W_{n_I1}^I & W_{n_I2}^I & \cdots & W_{n_In_I-1}^I \end{pmatrix}}_{\text{rows are iid}} \times \underbrace{\begin{pmatrix} x_1^{I-1} \\ x_2^{I-1} \\ \vdots \\ x_{n_I-1}^{I-1} \end{pmatrix}}_{\text{suppose all components are iid}} .$$

$$\underbrace{\begin{pmatrix} h_1^l \\ h_2^l \\ \dots \\ h_{n_l}^l \end{pmatrix}}_{\text{tends to a Gaussian vector with iid components as } n_{l-1} \rightarrow \infty \text{ by CLT}} = \underbrace{\begin{pmatrix} W_{11}^l & W_{12}^l & \dots & W_{1n_{l-1}}^l \\ W_{21}^l & W_{22}^l & \dots & W_{2n_{l-1}}^l \\ \dots & \dots & \dots & \dots \\ W_{n_l1}^l & W_{n_l2}^l & \dots & W_{n_ln_{l-1}}^l \end{pmatrix}}_{\text{all iid } \sim \mathcal{N}(0, \sigma_W^2/n_{l-1})} \times \underbrace{\begin{pmatrix} x_1^{l-1} \\ x_2^{l-1} \\ \dots \\ x_{n_{l-1}}^{l-1} \end{pmatrix}}_{\text{suppose all components are iid}}.$$

$$\underbrace{\begin{pmatrix} h_1^l & \bar{h}_1^l \\ h_2^l & \bar{h}_2^l \\ \dots & \dots \\ h_{n_l}^l & \bar{h}_{n_l}^l \end{pmatrix}}_{\substack{\text{each row converges to} \\ \text{a multivariate Gaussian vector} \\ \text{as } n_{l-1} \rightarrow \infty \text{ by the vector CLT}}} = \underbrace{\begin{pmatrix} W_{11}^l & W_{12}^l & \dots & W_{1n_{l-1}}^l \\ W_{21}^l & W_{22}^l & \dots & W_{2n_{l-1}}^l \\ \dots & \dots & \dots & \dots \\ W_{n_l1}^l & W_{n_l2}^l & \dots & W_{n_ln_{l-1}}^l \end{pmatrix}}_{\text{all iid } \sim \mathcal{N}(0, \sigma_W^2/n_{l-1})} \times \underbrace{\begin{pmatrix} x_1^{l-1} & \bar{x}_1^{l-1} \\ x_2^{l-1} & \bar{x}_2^{l-1} \\ \dots & \dots \\ x_{n_{l-1}}^{l-1} & \bar{x}_{n_{l-1}}^{l-1} \end{pmatrix}}_{\substack{\text{suppose all components} \\ \text{of each column are iid}}}.$$

Let $\xi_{1:M}$ be a batch of M inputs.

$$\begin{pmatrix} h_1^l(\xi_1) & h_1^l(\xi_2) & \dots & h_1^l(\xi_M) \\ h_2^l(\xi_1) & h_2^l(\xi_2) & \dots & h_2^l(\xi_M) \\ \dots & \dots & \dots & \dots \\ h_{n_l}^l(\xi_1) & h_{n_l}^l(\xi_2) & \dots & h_{n_l}^l(\xi_M) \end{pmatrix}$$

- Batch dimension — tends to a multivariate zero-mean Gaussian as $n_{l-1} \rightarrow \infty$;
- Neuron dimension — components tend to iid Gaussians.

Corollary (informal)

A neural net converges to a GP at initialization as $n_{1:L} \rightarrow \infty$ sequentially.

Off-topic remark: what happens during training?

1. When quadratic loss is optimized with gradient flow, a neural net remains a GP $\forall t > 0$;
2. The result of training corresponds to the result of GP inference iff only the readout layer is trained.

Does the previous result hold if some **weights are shared**?

Suppose we want to compute the limit distribution of WWx :

$$\underbrace{\begin{pmatrix} W_{11} & W_{12} & \cdots & W_{1n} \\ W_{21} & W_{22} & \cdots & W_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ W_{n1} & W_{n2} & \cdots & W_{nn} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} \times \underbrace{\begin{pmatrix} W_{11} & W_{12} & \cdots & W_{1n} \\ W_{21} & W_{22} & \cdots & W_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ W_{n1} & W_{n2} & \cdots & W_{nn} \end{pmatrix}}_{\text{the same matrix}} \times \underbrace{\begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}}_{\text{vector with iid components}}.$$

We can directly apply CLT if WW **has iid components**.

Let us compute the limit distribution of $(WW)_\alpha$:

$$\underbrace{\begin{pmatrix} W_{\alpha 1} & \color{red}{W_{\alpha\alpha}} & \cdots & W_{\alpha n} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} \underbrace{\begin{pmatrix} W_{11} & W_{1\alpha} & \cdots & W_{1n} \\ W_{\alpha 1} & \color{red}{W_{\alpha\alpha}} & \cdots & W_{\alpha n} \\ \cdots & \cdots & \cdots & \cdots \\ W_{n1} & W_{n\alpha} & \cdots & W_{nn} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} = \underbrace{\begin{pmatrix} \sum(\text{iid w. mean} = 0) \\ \color{red}{W_{\alpha\alpha}^2} + \sum(\text{iid w. mean} = 0) \\ \cdots \\ \sum(\text{iid w. mean} = 0) \end{pmatrix}}_{\text{all iid except for } \alpha\text{'s term}}^T$$

Fortunately, since $W_{\alpha\alpha} \sim \mathcal{N}(0, \sigma_W^2/n)$, $\color{red}{W_{\alpha\alpha}^2} x_\alpha \propto 1/n \rightarrow 0$.

Conclusion:

1. **CLT is still applicable** even if the same matrix W is used several times;
2. $(WW_X)_\alpha$ and $(W\tilde{W}_X)_\alpha$ with $\tilde{W} \stackrel{d}{=} W$ have **the same distribution** in the limit.

However, we **cannot substitute** W with its iid copy \tilde{W} :

- WW_X and W_X are correlated in the limit;
- $W\tilde{W}_X$ and \tilde{W}_X are not.

Consider $h = Wx$ and $\bar{h} = W\bar{x}$ where x and \bar{x} **can depend on** W .

$$\underbrace{\begin{pmatrix} h_1 & \bar{h}_1 \\ h_2 & \bar{h}_2 \\ \dots & \dots \\ h_n & \bar{h}_n \end{pmatrix}}_{\substack{\text{each row converges to a zero-mean Gaussian} \\ \text{with covariance } \sigma_W^2 \mathbb{E}(x^\top \bar{x})/n \\ \text{as } n \rightarrow \infty \text{ by the vector CLT}}} = \underbrace{\begin{pmatrix} W_{11} & W_{12} & \dots & W_{1n} \\ W_{21} & W_{22} & \dots & W_{2n} \\ \dots & \dots & \dots & \dots \\ W_{n1} & W_{n2} & \dots & W_{nn} \end{pmatrix}}_{\text{all iid } \sim \mathcal{N}(0, \sigma_W^2/n)} \times \underbrace{\begin{pmatrix} x_1 & \bar{x}_1 \\ x_2 & \bar{x}_2 \\ \dots & \dots \\ x_n & \bar{x}_n \end{pmatrix}}_{\substack{\text{suppose all components} \\ \text{of each column are iid}}}.$$

A NETSOR program = (a set of input vars, a sequence of commands),

where variables are of three different **types**:

1. A-vars: matrices with iid Gaussian entries;
2. G-vars: vectors with *asymptotically* iid Gaussian entries;
3. H-vars: images of G-vars by coordinatewise nonlinearities.

Each command generates a new variable from the previous ones using one of the following **ops**:

1. MatMul: $(W : A, x : H) \rightarrow Wx : G;$
2. LinComb: $(\{x_i : G, a_i \in \mathbb{R}\}_{i=1}^k) \rightarrow \sum_{i=1}^k a_i x_i : G;$
3. Nonlin: $(\{x_i : G\}_{i=1}^k, \phi : \mathbb{R}^k \rightarrow \mathbb{R}) \rightarrow \phi(x_{1:k}) : H.$

Initialization assumption:

1. All hidden dimensions are equal to n ;
2. $\forall W : A$ we sample $W_{\alpha\beta} \sim \mathcal{N}(0, \sigma_W^2/n)$ iid;

Our goal: compute the distributions of all G-vars in the program in the limit of $n \rightarrow \infty$.

Claim:

Let $g^{1:M}$ be a set of all G-vars in the program.

$$\begin{pmatrix} g_1^1 & g_1^2 & \dots & g_1^M \\ g_2^1 & g_2^2 & \dots & g_2^M \\ \dots & \dots & \dots & \dots \\ g_n^1 & g_n^2 & \dots & g_n^M \end{pmatrix}$$

- "Batch" dimension — converges to $\mathcal{N}(\mu, \Sigma)$ with $\mu = \{\mu(g^i)\}_{i=1}^M$ and $\Sigma = \{\Sigma(g^i, g^j)\}_{i,j=1}^M$ defined below;
- Neuron dimension — components tend to be iid Gaussians.

$(g_\alpha^1, \dots, g_\alpha^M)$ becomes jointly Gaussian with mean and covariance **defined by CLT**¹:

$$\mu(g) = \begin{cases} 0 & \text{if } g = Wy. \end{cases} \quad (2)$$

$$\Sigma(g, \bar{g}) = \begin{cases} \sigma_W^2 \mathbb{E}_Z \phi(Z) \bar{\phi}(Z) & \text{if } g = W\phi(Z) \text{ and } \bar{g} = W\bar{\phi}(Z); \\ 0 & \text{else.} \end{cases} \quad (3)$$

Here $Z \sim \mathcal{N}(\mu, \Sigma)$ is a set of all previous G-vars.

¹we have suppressed the LinComb op for brevity.

Theorem (Netsor Master Theorem, [Yang, 2019], informal)

Let the `NETSOR` program satisfy the initialization assumption and let all nonlinearities do not grow too fast. Let $g^{1:M}$ be a set of all G -vars in the program. Then, for any well-behaved ψ ,

$$\underbrace{\frac{1}{n} \sum_{\alpha=1}^n \psi(g_{\alpha}^1, \dots, g_{\alpha}^M)}_{\text{empirical mean over rows}} \rightarrow \underbrace{\mathbb{E}_{Z \sim \mathcal{N}(\mu, \Sigma)} \psi(Z)}_{\substack{\text{expectation} \\ \text{over the corresponding Gaussian}}} \quad (4)$$

a.s. as $n \rightarrow \infty$, where $\mu = \{\mu(g^i)\}_{i=1}^M$ and $\Sigma = \{\Sigma(g^i, g^j)\}_{i,j=1}^M$.

A NETSOR program

- is able to express the **first forward pass** of a wide class of neural nets (i.e. with shared/structured weights, with BNs etc.);
- reveals its limiting Gaussian process behavior.

Questions:

1. Can we express a **backward pass** as a NETSOR program?
2. What is its limiting behavior?

Intermedia: a Neural Tangent Kernel

Let $f(\cdot; \theta)$ be a parametric model.

Define a **neural tangent kernel** as

$$\Theta_t(\xi, \bar{\xi}) = \underbrace{\nabla_{\theta}^T f(\xi; \theta_t) \nabla_{\theta} f(\bar{\xi}; \theta_t)}_{\text{"gradient similarity"}}. \quad (5)$$

It drives evolution of **model predictions**; e.g. for square loss:

$$\dot{f}(\bar{\xi}; \theta_t) = (\vec{y} - f(\vec{\xi}; \theta_t))^{\top} \Theta_t(\vec{\xi}, \bar{\xi}), \quad \text{where } (\vec{\xi}, \vec{y}) \text{ is a train dataset.}$$

Assuming $\Theta_t(\cdot, \cdot) \approx \Theta_0(\cdot, \cdot)$ makes the dynamics **analytically tractable**.

Theorem ([Jacot et al., 2018], informal)

Suppose we have a feedforward neural net of width n parameterized in a certain way. Then, as $n \rightarrow \infty$,

1. $\Theta_0(\xi, \bar{\xi})$ converges to a deterministic $\mathring{\Theta}(\xi, \bar{\xi})$;
2. Moreover, $\Theta_t(\xi, \bar{\xi})$ converges to the same $\mathring{\Theta}(\xi, \bar{\xi})$.

Parameterize W^l as $\omega^l/\sqrt{n_{l-1}}$:

$$f = \underbrace{\frac{1}{\sqrt{n_L}} \omega^{L+1}}_{\text{the network output}} x^L, \quad \underbrace{x^l = \phi(h^l)}_{\text{activations}}, \quad \underbrace{h^l = \frac{1}{\sqrt{n_{l-1}}} \omega^l x^{l-1}}_{\text{preactivations}}, \quad l \leq L; \quad \omega_{ij}^l \sim \mathcal{N}(0, \sigma_W^2) \text{ iid.}$$

NTK is defined as

$$\Theta(\xi, \bar{\xi}) = \nabla_{\theta}^T f(\xi; \theta) \nabla_{\theta} f(\bar{\xi}; \theta) = \sum_{l=1}^{L+1} \underbrace{\text{tr}(\nabla_{\omega^l}^T f(\xi) \nabla_{\omega^l} f(\bar{\xi}))}_{\text{"layer-wise gradient similarity"}}. \quad (6)$$

$$\Theta(\xi, \bar{\xi}) = \nabla_{\theta}^T f(\xi; \theta) \nabla_{\theta} f(\bar{\xi}; \theta) = \sum_{l=1}^{L+1} \underbrace{\text{tr}(\nabla_{\omega^l}^T f(\xi) \nabla_{\omega^l} f(\bar{\xi}))}_{\text{"layer-wise gradient similarity"}} . \quad (7)$$

Weight gradient can be expressed as

$$\nabla_{\omega^l} f = \frac{1}{\sqrt{n_{l-1} n_l}} \underbrace{dh^l}_{\substack{\text{backward pass} \\ \text{up to the layer } l}} \times \underbrace{x^{l-1, \top}}_{\substack{\text{forward pass} \\ \text{up to the layer } l-1}} , \quad \text{where } dh^l \propto \nabla_{h_l} f. \quad (8)$$

Plug (8) into (7):

$$\Theta(\xi, \bar{\xi}) = \sum_{l=1}^{L+1} \underbrace{\left(\frac{dh^{l, \top} d\bar{h}^l}{n_l} \right)}_{\text{"backward pass similarity"}} \times \underbrace{\left(\frac{x^{l-1, \top} \bar{x}^{l-1}}{n_{l-1}} \right)}_{\text{"forward pass similarity"}} . \quad (9)$$

Consider the second multiplier:

$$\frac{x^{l-1, \top} \bar{x}^{l-1}}{n_{l-1}} = \frac{1}{n_{l-1}} \sum_{\alpha=1}^{n_{l-1}} \phi(h_{\alpha}^{l-1}) \phi(\bar{h}_{\alpha}^{l-1}) = \underbrace{\frac{1}{n_{l-1}} \sum_{\alpha=1}^{n_{l-1}} \psi(h_{\alpha}^{l-1}, \bar{h}_{\alpha}^{l-1})}_{\text{the limit is given by the Master theorem!}} \quad \text{for } \psi(x, y) = \phi(x)\phi(y).$$

Can we compute the limit of the first multiplier in the same way?

For simplicity, assume $n_1 = \dots = n_L = n$. Recall $W^l = \omega^l / \sqrt{n}$.

Relations between forward and backward passes:

Forward pass:	Backward pass:
$x^l = \phi(h^l) : \text{Nonlin}$	$dh^l = dx^l \odot \phi'(h^l) : \text{Nonlin}$
$h^l = W^l x^{l-1} : \text{MatMul}$	$dx^{l-1} = W^{l,\top} dh^l : \text{MatMul?}$

Problems:

1. W and W^\top cannot be both input variables since they are dependent;
2. A NETSOR program does not allow for multiplying by a transposed A-var.

A Netsor program cannot express the backward pass!

A **NETSORT** program = (a set of input vars, a sequence of commands),

where variables are of three different **types**:

1. A-vars: matrices with iid Gaussian entries;
2. G-vars: vectors with *asymptotically* iid Gaussian entries;
3. H-vars: images of G-vars by coordinatewise nonlinearities.

Each command generates a new variable from the previous ones using one of the following **ops**:

1. **Trsp**: $W : A \rightarrow W^T : A$;
2. **MatMul**: $(W : A, x : H) \rightarrow Wx : G$;
3. **LinComb**: $(\{x_i : G, a_i \in \mathbb{R}\}_{i=1}^k) \rightarrow \sum_{i=1}^k a_i x_i : G$;
4. **Nonlin**: $(\{x_i : G\}_{i=1}^k, \phi : \mathbb{R}^k \rightarrow \mathbb{R}) \rightarrow \phi(x_{1:k}) : H$.

Can we keep the same symbolic rules for mean and covariance of G-vars?

$$\mu(g) = \begin{cases} 0 & \text{if } g = Wy; \\ 0? & \text{if } g = W^\top y. \end{cases} \quad (10)$$

$$\Sigma(g, \bar{g}) = \begin{cases} \sigma_W^2 \mathbb{E}_Z \phi(Z) \bar{\phi}(Z) & \text{if } g = W\phi(Z) \text{ and } \bar{g} = W\bar{\phi}(Z); \\ \text{some other rule?} & \text{if } g = W\phi(Z) \text{ and } \bar{g} = W^\top \bar{\phi}(Z); \\ 0 & \text{else.} \end{cases} \quad (11)$$

Here $Z \sim \mathcal{N}(\mu, \Sigma)$ is a set of all previous G-vars.

Let us check that $\mu(WW_X) = 0$:

$$\underbrace{\begin{pmatrix} W_{\alpha 1} & \color{red}{W_{\alpha\alpha}} & \cdots & W_{\alpha n} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} \underbrace{\begin{pmatrix} W_{11} & W_{1\alpha} & \cdots & W_{1n} \\ W_{\alpha 1} & \color{red}{W_{\alpha\alpha}} & \cdots & W_{\alpha n} \\ \cdots & \cdots & \cdots & \cdots \\ W_{n1} & W_{n\alpha} & \cdots & W_{nn} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} = \underbrace{\begin{pmatrix} \sum(\text{iid w. mean} = 0) \\ \color{red}{W_{\alpha\alpha}^2} + \sum(\text{iid w. mean} = 0) \\ \cdots \\ \sum(\text{iid w. mean} = 0) \end{pmatrix}}_{\text{all iid except for } \alpha\text{'s term}}^T$$

$$\mu(WW_X) = \mathbb{E}((WW_X)_\alpha) = \underbrace{\mathbb{E}\left(\sum_{\beta \neq \alpha} \sum_{\gamma} W_{\alpha\beta} W_{\beta\gamma} x_\gamma\right)}_{\mathbb{E}(\sum \text{iid w. mean}=0)=0} + \underbrace{\mathbb{E}\left(\sum_{\gamma} W_{\alpha\alpha} W_{\alpha\gamma} x_\gamma\right)}_{=\mathbb{E} W_{\alpha\alpha}^2 x_\alpha \propto 1/n \rightarrow 0}. \quad (12)$$

Do we have $\mu(W \color{red}{W}^T X) = 0$?

Let us compute $\mu(WW^T x)$:

$$\underbrace{\begin{pmatrix} W_{\alpha 1} & W_{\alpha \alpha} & \dots & W_{\alpha n} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} \underbrace{\begin{pmatrix} W_{11} & W_{\alpha 1} & \dots & W_{n1} \\ W_{1\alpha} & W_{\alpha \alpha} & \dots & W_{n\alpha} \\ \dots & \dots & \dots & \dots \\ W_{1n} & W_{\alpha n} & \dots & W_{nn} \end{pmatrix}}_{\text{all iid} \sim \mathcal{N}(0, \sigma_W^2/n)} = \underbrace{\begin{pmatrix} \sum(\text{iid w. mean} = 0) \\ \sum_{\beta} W_{\alpha\beta}^2 \\ \dots \\ \sum(\text{iid w. mean} = 0) \end{pmatrix}}_{\text{all iid except for } \alpha\text{'s term}}^T$$

$$\mu(WW^T x) = \mathbb{E}((WW^T x)_{\alpha}) = \underbrace{\mathbb{E}\left(\sum_{\beta} \sum_{\gamma \neq \alpha} W_{\alpha\beta} W_{\gamma\beta} x_{\gamma}\right)}_{\mathbb{E}(\sum \text{iid w. mean}=0)=0} + \underbrace{\mathbb{E}\left(\sum_{\beta} W_{\alpha\beta}^2 x_{\beta}\right)}_{\text{converges to } \sigma_W^2 \mu(x) \neq 0}. \quad (13)$$

The previous symbolic rules are not applicable for **general** `NETSORT` programs,
but
they are applicable to `NETSORT` programs expressing backpropagation.

Claim: the rule

$$\mu(g) = 0 \quad \text{if } g = Wy. \quad (14)$$

works for `NETSORT` programs expressing backpropagation.

Evidence: consider $dx^{l-1} = \textcolor{red}{W}^{l,T}(dx^l \odot \phi'(h^l))$. Let $\phi(z) = z^2/2$:

$$\mu(dx^{l-1}) = \mathbb{E}(dx_\alpha^{l-1}) = \underbrace{\mathbb{E}\left(\sum_{\beta} W_{\beta\alpha}^l dx_{\beta}^l \sum_{\gamma \neq \beta} W_{\beta\gamma}^l x_{\gamma}^{l-1}\right)}_{\mathbb{E}(\sum \text{iid w. mean}=0)=0} + \underbrace{\mathbb{E}\left(x_{\alpha}^{l-1} \sum_{\beta} (\textcolor{red}{W}_{\beta\alpha}^l)^2 dx_{\beta}^l\right)}_{\text{converges to } \mu(x^{l-1})\sigma_W^2 \mu(dx^l)}. \quad (15)$$

Hence $\mu(dx^{l-1}) \propto \mu(dx^l)$ which by induction implies $\mu(dx^{l-1}) \propto \mu(dx^L)$.

But $dx^L = \omega^{L+1}!$ Hence $\mu(dx^{l-1}) = \mu(\omega^{L+1}) = 0$.

Proposition (²)

Consider a neural network and a NETSORT program expressing its backward pass.

The symbolic rules for μ and Σ are valid, if

- 1. The output layer has zero mean;*
- 2. It is sampled independently from other parameters;*
- 3. It is not used anywhere else in the program.*

²There is a more general condition called "BP-likeness" which we do not show here.

Theorem (Netsor[⊤] Master Theorem, [Yang, 2020a]; informal)

Let a NETSOR[⊤] program express a forward or a backward pass in a neural network, satisfy the initialization assumption, and let all nonlinearities do not grow too fast. Let $g^{1:M}$ be a set of all G-vars in the program. Then, for any well-behaved $\psi : \mathbb{R}^M \rightarrow \mathbb{R}$,

$$\frac{1}{n} \sum_{\alpha=1}^n \psi(g_{\alpha}^1, \dots, g_{\alpha}^M) \rightarrow \mathbb{E}_{Z \sim \mathcal{N}(\mu, \Sigma)} \psi(Z)$$

a.s. as $n \rightarrow \infty$, where $\mu = \{\mu(g^i)\}_{i=1}^M$ and $\Sigma = \{\Sigma(g^i, g^j)\}_{i,j=1}^M$.

The result is (almost) the same as for Netsor programs!

Back to NTK computation:

$$\Theta(\xi, \bar{\xi}) = \sum_{l=1}^{L+1} \nabla_{\omega^l}^T f(\xi) \nabla_{\omega^l} f(\bar{\xi}) = \sum_{l=1}^{L+1} \left(\frac{dh^l, \top d\bar{h}^l}{n_l} \right) \left(\frac{x^{l-1, \top} \bar{x}^{l-1}}{n_{l-1}} \right). \quad (16)$$






Consider the first multiplier:

$$\frac{dh^l, \top d\bar{h}^l}{n_l} = \frac{1}{n_l} \sum_{\alpha=1}^{n_l} dx_{\alpha}^l d\bar{x}_{\alpha}^l \phi'(h_{\alpha}^l) \phi'(\bar{h}_{\alpha}^l) = \underbrace{\frac{1}{n_l} \sum_{\alpha=1}^{n_l} \psi(dx_{\alpha}^l, d\bar{x}_{\alpha}^l, h_{\alpha}^l, \bar{h}_{\alpha}^l)}_{\text{the limit is given by the Master theorem!}}$$

for $\psi(x, y, z, w) = xy\phi'(z)\phi'(w)$. (17)

A BP-like NETSORT program

- is able to express the **first forward and backward passes** of a wide class of neural nets (i.e. with shared/structured weights, with BNs etc.);
- reveals their limiting Gaussian process behavior;
- can be applied to initial NTK computation.

-  Jacot, A., Gabriel, F., and Hongler, C. (2018).
Neural tangent kernel: Convergence and generalization in neural networks.
In Advances in neural information processing systems, pages 8571–8580.
-  Matthews, A. G. d. G., Hron, J., Rowland, M., Turner, R. E., and Ghahramani, Z. (2018).
Gaussian process behaviour in wide deep neural networks.
In International Conference on Learning Representations.
-  Pennington, J., Schoenholz, S., and Ganguli, S. (2017).
Resurrecting the sigmoid in deep learning through dynamical isometry: theory and practice.
In Advances in neural information processing systems, pages 4785–4795.
-  Yang, G. (2019).
Tensor programs i: Wide feedforward or recurrent neural networks of any architecture are gaussian processes.
arXiv preprint arXiv:1910.12478.
-  Yang, G. (2020a).

Tensor programs ii: Neural tangent kernel for any architecture.

arXiv preprint arXiv:2006.14548.



Yang, G. (2020b).

Tensor programs iii: Neural matrix laws.

arXiv preprint arXiv:2009.10685.