

02/22 ML Meeting

ICLR'21 - Deployment-Efficient Reinforcement Learning via Model-Based Offline Optimization

Tatsuya Matsushima, Hiroki Furuta, Yutaka Matsuo, Ofir Nachum, Shixiang
Shane Gu

The University of Tokyo, Google Research

Citation: 45

Problem Formulation

- In real world application, it's **costly and risky** to update the policy deployed on a cluster.
 - Costly: Updating the policy of the RL deployment **costs high communication cost between the clusters.**
 - Risky: The **new policy may act unexpectedly.**
 - **It's impractical to update the RL policy frequently**, so we need to reduce the times of the deployment as many as possible

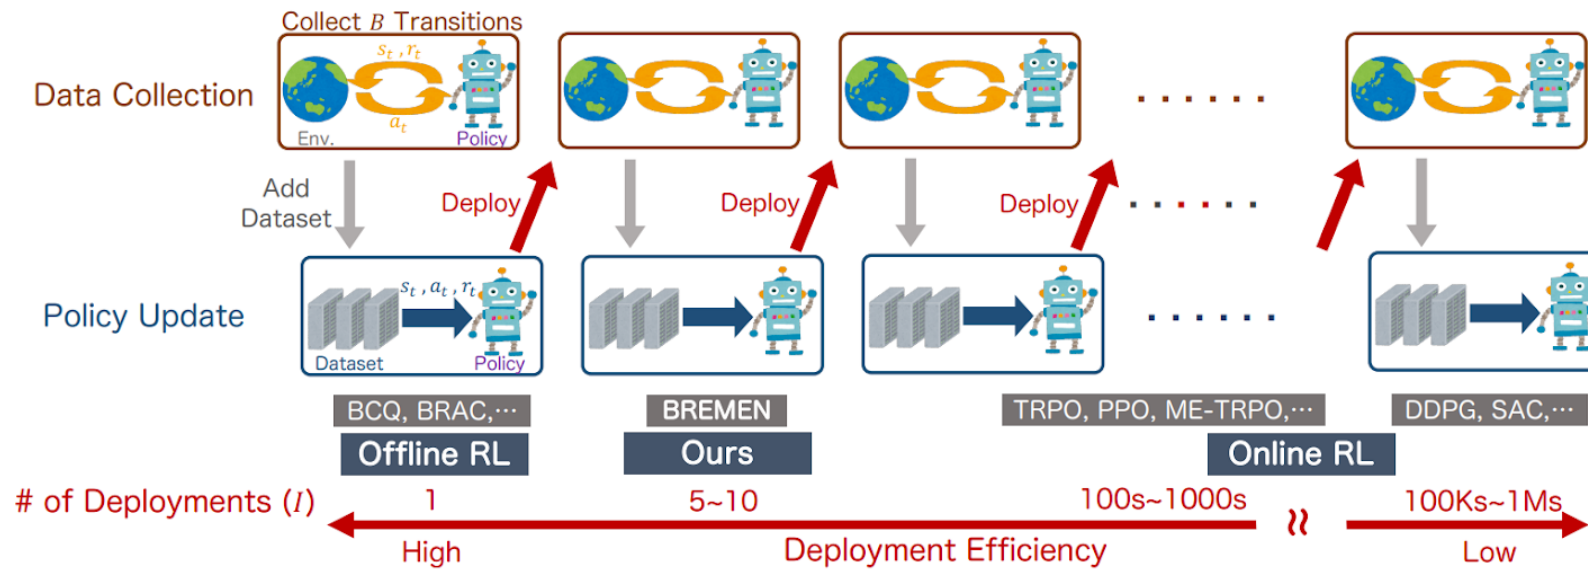


Figure 1: *Deployment efficiency* is defined as the number of changes in the data-collection policy (I), which is vital for managing costs and risks of new policy deployment. Online RL algorithms typically require many iterations of policy deployment and data collection, which leads to extremely low deployment efficiency. In contrast, most pure offline algorithms consider updating a policy from a fixed dataset without additional deployment and often fail to learn from a randomly initialized data-collection policy. Interestingly, most state-of-the-art off-policy algorithms are still evaluated in heavily online settings. For example, SAC (Haarnoja et al., 2018) collects one sample per policy update, amounting to 100,000 to 1 million deployments for learning standard benchmark domains.

- Formulate a new problem: Deployment-Efficient RL
- Definition of “Deployment-Efficient”: Deployment efficiency is defined as the number of changes in the data-collection policy.

Challenges

- **Online-RL** update policy frequently -> **Low deployment efficiency**
- **Offline-RL** collect dataset only once -> **Can only get sub-optimal policy**
- **Get we get the optimal policy with high deployment efficiency?**

Idea

- Reduce the update frequency of the data-collecting policy and only update it when it has collected a large batch of trajectories
- Use **importance sampling and trust-region optimization** to correct the distribution shift between data-collecting policy π_θ and training(target) policy π_{θ_k}

$$\begin{aligned} \theta_{k+1} = \arg \max_{\theta} \mathbb{E}_{s, a \sim \theta_k} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\theta_k}(a|s)} A^{\pi_{\theta_k}}(s, a) \right] \\ \text{s.t.} \quad \mathbb{E} [D_{KL}(\pi_{\theta}(\cdot|s) || \pi_{\theta_k}(\cdot|s))] \leq \delta \end{aligned}$$

- Named 'BREMEN'

Disadvantages:

- Importance sampling suffers from high variance while the data-collecting policy π_θ is largely different from training(target) policy π_{θ_k} . It would cause low sample-efficiency.
- Trade-off between sample efficiency and deployment efficiency

Following Works

arXiv'22 - Deployment Constrained Reinforcement Learning with Model-based Uncertainty Regularized Batch Optimization

- Improve the sample efficiency of previous work 'BREMEN'
- Encourage the agent to explore low-data regions.
- Train a world model $f(s_t, a_t) = s_{t+1}$ If the prediction error of the next state given a state-action pair is high, then the uncertainty is high. We take this prediction error as an uncertainty quantifier.

Conclusion

It seems that for a practical RL algorithm, it had better satisfy 3 requirements:

- **Sample efficient during policy optimization**
- **Deployment efficient during data collection**
- **Safe exploration during data collection**

Idea

- Explore the state / action which is **high expectation & high variance of the Q-value to collect diverse trajectories**:
- high variance & high expectation > low variance & high expectation > low variance & low expectation > high variance & low expectation

$$\arg \max_{\theta} \sum_{t=0}^T K(\mathbb{E}[Q_{offline}(s_t, \pi_{\theta}(s_t); \mathcal{D})], \text{var}[Q_{offline}(s_t, \pi_{\theta}(s_t); \mathcal{D})])$$

Suppose the offline Q-function $Q_{offline}(s, a; \mathcal{D})$ is trained in dataset \mathcal{D} and normalized to $0 \sim 1$. Function $K(a, b)$ would give higher value if the pair (a, b) is in the higher order as the second bullet point

- The variance of the Q-function can be replaced by a quantifier of the diversification of the gradient of Q-function because diverse gradients can provide diverse state-action pair / trajectories.