

# Uncertainty-Based Offline Reinforcement Learning with Diversified Q-Ensemble

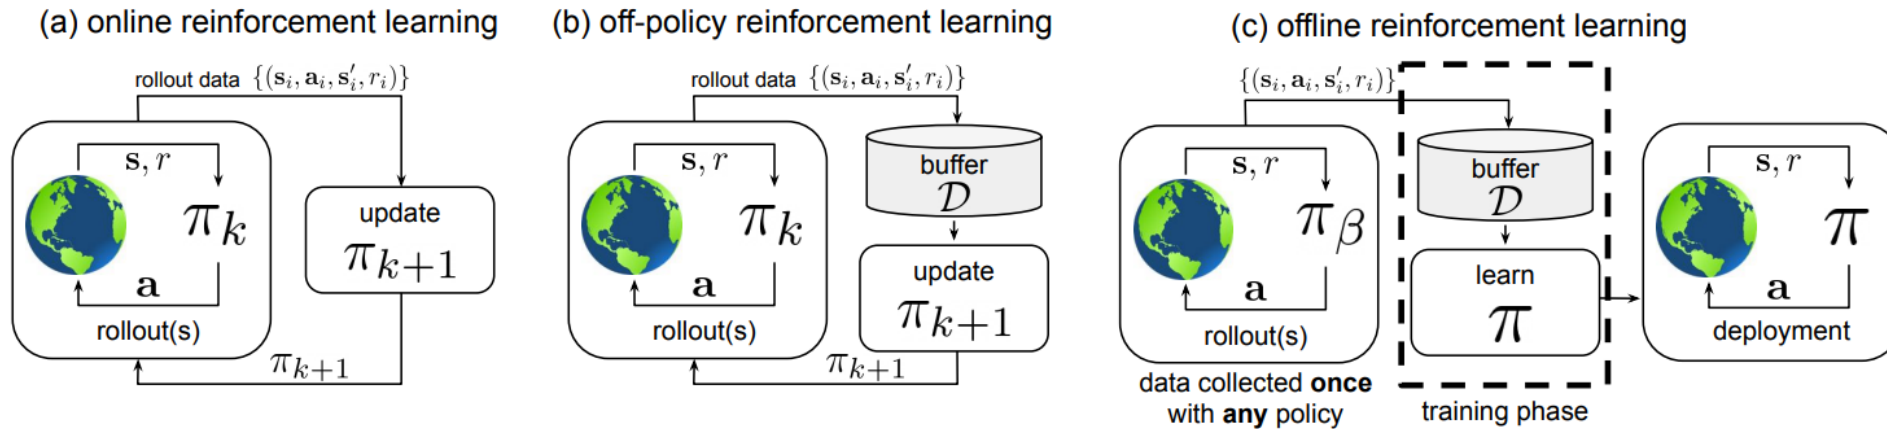
ICML'21

Citation Count: 1

Gaon An, Seungyong Moon, Jang-Hyun Kim, Hyun Oh Song

Seoul National University, Neural Processing Research Center, DeepMetrics

# Offline RL



We aim to learn a policy  $\pi$  from the history of trajectories  $\mathcal{D} = \{s_t, a_t, s'_t, r_t\}$  generated by behavioral policy  $\pi_\beta$  s.t. the performance  $\pi \geq \pi_\beta$  (Which means we want to train an agent's policy  $\pi$  only from the history of trajectories  $\mathcal{D}$ )

# Challenge

- Extrapolation Error
  - The agent may overestimate unseen  $Q^\pi(s, a)$ , which gives a higher  $Q^\pi(s, a)$  value than the optimal  $Q^{\pi^*}(s, a)$  value .
  - In the offline setting, the policy cannot correct such over-optimistic Q-values.

- Solution
  - If we add a regularizer to the equation in order to make the agent (1) underestimate the  $Q^\pi(s, a)$  value of unseen action  $a$  given a state  $s$  or (2) choose the action that closes to the action already in the history trajectory  $\mathcal{D}$  given a state (in practice, choose the action that are higher than the a threshold probability from the action distribution), we can avoid the crazy actions that the agent may do.
  - But a trade-off between **Optimality** and **Conservativeness**.

$$\hat{Q}_{k+1}^\pi \leftarrow \arg \min_Q \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} \left[ \left( Q(\mathbf{s}, \mathbf{a}) - \left( r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{a}' \sim \pi_k(\mathbf{a}'|\mathbf{s}')} [\hat{Q}_k^\pi(\mathbf{s}', \mathbf{a}')] \right) \right)^2 \right]$$

$$\pi_{k+1} \leftarrow \arg \max_{\pi} \mathbb{E}_{\mathbf{s} \sim \mathcal{D}} \left[ \mathbb{E}_{\mathbf{a} \sim \pi(\mathbf{a}|\mathbf{s})} [\hat{Q}_{k+1}^\pi(\mathbf{s}, \mathbf{a})] \right] \text{ s.t. } D(\pi, \pi_\beta) \leq \epsilon.$$

# Idea

- Penalize the Q-function with the most pessimistic Q-network of the ensemble Q-network.

Modify the following SAC objective

$$\min_{\phi} \mathbb{E}_{(\mathbf{s}, \mathbf{a}, \mathbf{s}') \sim \mathcal{D}} \left[ \left( Q_{\phi}(\mathbf{s}, \mathbf{a}) - \left( r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{a}' \sim \pi_{\theta}(\cdot | \mathbf{s}')} [Q_{\phi'}(\mathbf{s}', \mathbf{a}')] - \beta \log \pi_{\theta}(\mathbf{a}' | \mathbf{s}') \right) \right)^2 \right]$$
$$\max_{\theta} \mathbb{E}_{\mathbf{s} \sim \mathcal{D}, \mathbf{a} \sim \pi_{\theta}(\cdot | \mathbf{s})} [Q_{\phi}(\mathbf{s}, \mathbf{a}) - \beta \log \pi_{\theta}(\mathbf{a} | \mathbf{s})]$$

# Idea

To SAC-N

$$\min_{\phi_i} \mathbb{E}_{s,a,s' \sim \mathcal{D}} \left[ \left( Q_{\phi_i}(s, a) - \left( r(s, a) + \gamma \mathbb{E}_{a' \sim \pi_{\theta}(\cdot | s')} \left[ \min_{j=1, \dots, N} Q_{\phi'_j}(s', a') - \beta \log \pi_{\theta}(a' | s') \right] \right) \right)^2 \right]$$
$$\max_{\theta} \mathbb{E}_{s \sim \mathcal{D}, a \sim \pi_{\theta}(\cdot | s)} \left[ \min_{j=1, \dots, N} Q_{\phi_j}(s, a) - \beta \log \pi_{\theta}(a | s) \right]$$

Where  $\phi$  is the parameters of the Q-network  $Q_{\phi}$ ,  $\theta$  is the parameters of policy network  $\pi_{\theta}$ . The subscript  $j$  means the  $j$ -th network.

# Idea

And the authors surprisingly found that SAC-N will outperform than the SOTA offline-RL algorithm "CQL" when the number of ensemble is large enough.

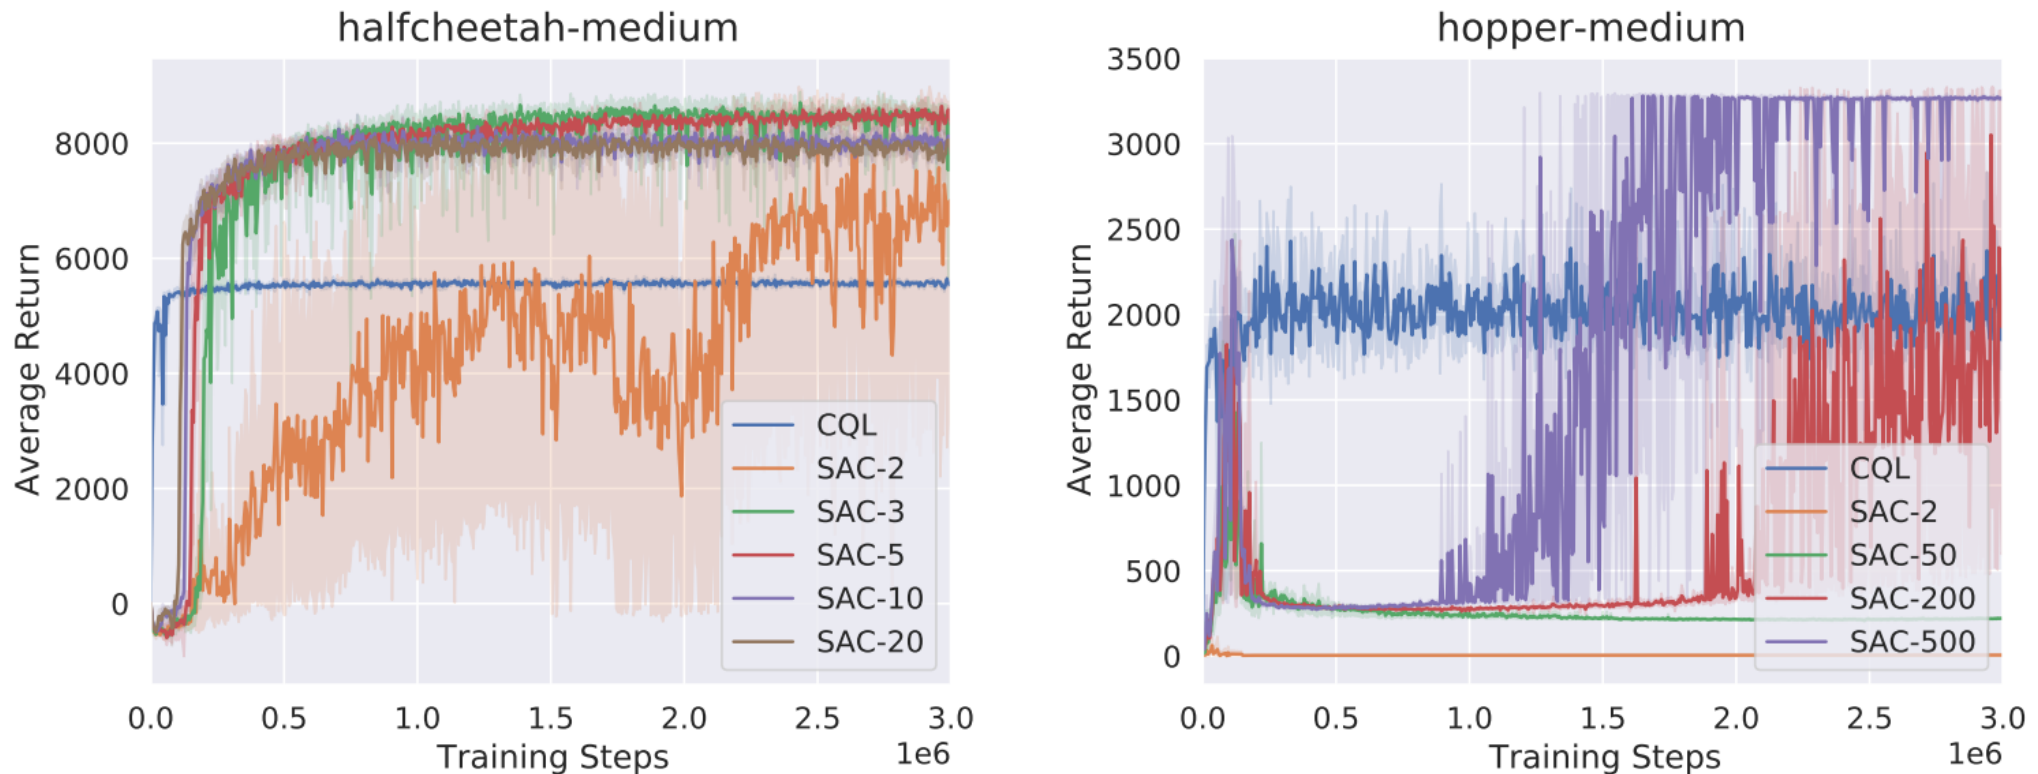


Figure 1: Performance of SAC- $N$  on halfcheetah-medium and hopper-medium datasets while varying  $N$ , compared to CQL. 'Average Return' denotes the undiscounted return of each policies on evaluation. Results averaged over 4 seeds.

# Idea

- Obviously, the redundant Q-networks of SAC-N cost lots of computation. The authors aim to reduce the size of the ensemble Q-network while achieving the same performance.
- The authors found that the performance of SAC-N is negatively correlated with the degree to which the input gradients of Q-functions  $\nabla_a Q_{\phi_j}(s, a)$  are aligned, which increases with  $N$ .
- Note that out-of-distribution state means the probability of the state that appears in the dataset is lower than a given threshold. Similarly, Out-of-distribution action means the action that appears in the dataset is lower than a given threshold and state. In the other hand, in-distribution action means higher than the threshold.

Argument: Agent performance -> Variance of Q-value of OOD action -> Diversification of the gradients of the Q-network



# Evidence 1

The Q-value predictions for the OOD actions have a higher variance.

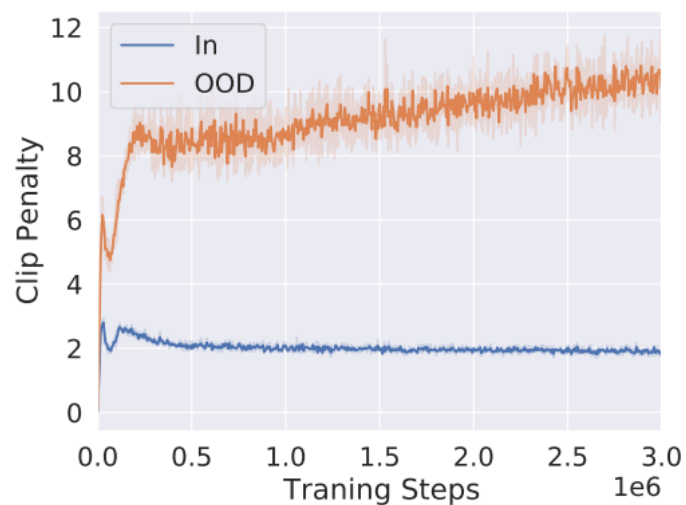
- Here we define the **penalty from the clipping** as

$$\mathbb{E}_{s \sim D, a \sim \pi(\cdot|s)} \left[ \frac{1}{N} \sum_{j=1}^N Q_{\phi_j}(s, a) - \min_{j=1, \dots, N} Q_{\phi_j}(s, a) \right]$$

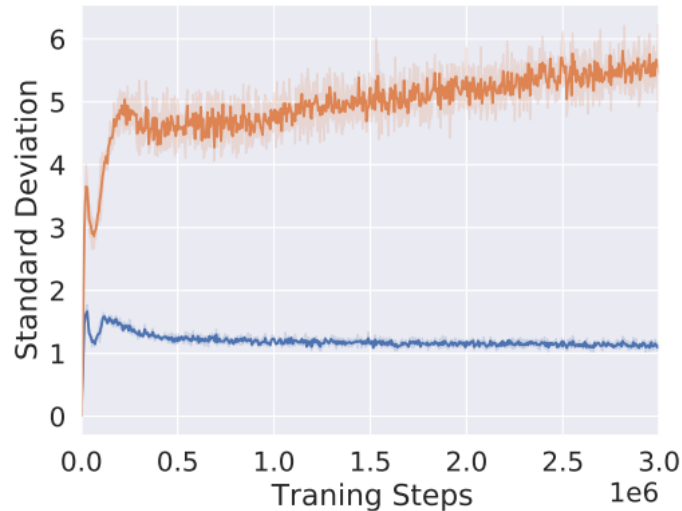
- We also **approximate the expected minimum of the realizations** following the work of Royston

$$\mathbb{E} \left[ \min_{j=1, \dots, N} Q_j(s, a) \right] \approx m(s, a) - \Phi^{-1} \left( \frac{N - \frac{\pi}{8}}{N - \frac{\pi}{4} + 1} \right) \sigma(s, a)$$

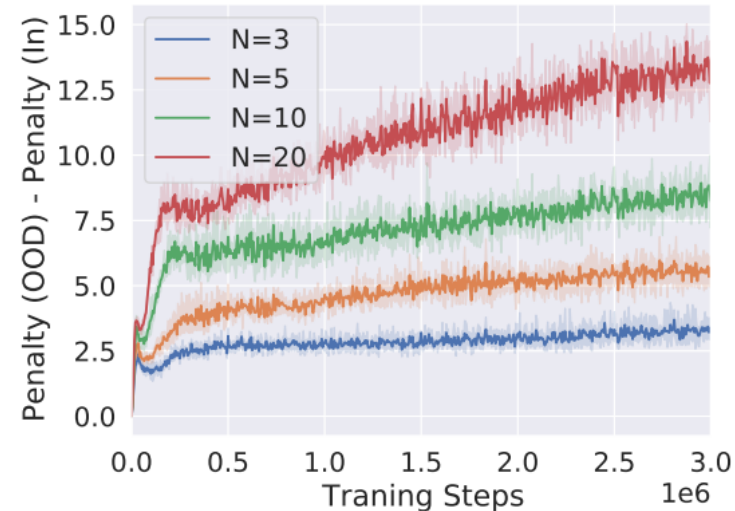
Where we suppose  $Q(s, a) \sim \mathcal{N}(m(s, a), \sigma(s, a))$  and  $\Phi$  is the CDF of the standard Gaussian distribution.



(a) Clip penalty



(b) Standard deviation



(c) Clip penalty gap

Figure 2: (a) and (b) each plots the size of the clip penalty and the standard deviation of the Q-value estimates for in-distribution (behavior) and OOD (random) actions while training SAC-10 on halfcheetah-medium dataset. (c) plots the gap of the clip penalty between the in-distribution and OOD actions while varying  $N$ . Results averaged over 4 seeds.

The Q-value predictions for the **OOD** actions have a higher variance and the size of the penalty and the standard deviation are highly correlated.

## Evidence 2

The performance of the learned policy degrades significantly when the Q-functions share a similar local structure.

The minimum cosine similarity between the gradients of the Q-functions is

$$\min_{i \neq j} \langle \nabla_a Q_{\phi_i}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle$$

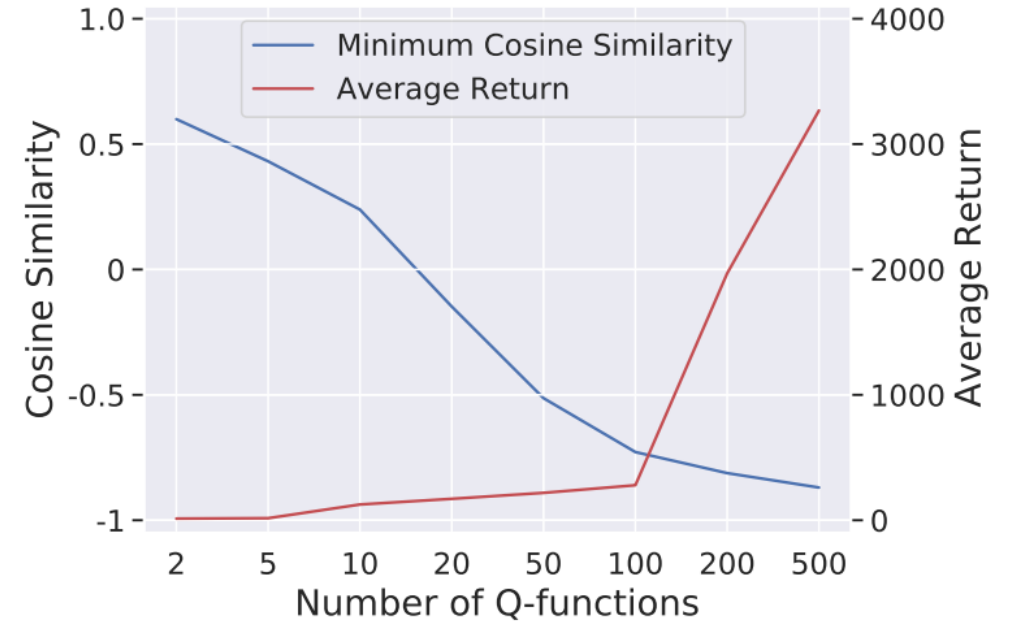


Figure 4: Plot of the minimum cosine similarity between the input gradients of Q-functions and the average return while varying the number of Q-functions.

We use Taylor expansion to expand the Q function  $Q_{\phi_j}(s, a + kw_2)$

$$\begin{aligned}\text{Var}(Q_{\phi_j}(s, a + kw_2)) &\approx \text{Var}(Q_{\phi_j}(s, a) + k\langle w_2, \nabla_a Q_{\phi_j}(s, a) \rangle) \\ &= \text{Var}(Q(s, a) + k\langle w_2, \nabla_a Q_{\phi_j}(s, a) \rangle) \\ &= k^2 \text{Var}(\langle w_2, \nabla_a Q_{\phi_j}(s, a) \rangle) \\ &= k^2 w_2^\top \text{Var}(\nabla_a Q_{\phi_j}(s, a)) w \\ &\geq k^2 w_{\min}^\top \text{Var}(\nabla_a Q_{\phi_j}(s, a)) w_{\min} \\ &= k^2 \lambda_{\min}\end{aligned}$$

Where  $w_{\min}$  and  $\lambda_{\min}$  are the smallest eigenvector and eigenvalue for the matrix  $\text{Var}(\nabla_a Q_{\phi_j}(s, a))$

However, It's hard to compute the smallest eigenvalue. Thus, the authors **decide to maximize the sum of the eigenvalues instead of the smallest eigenvalue.**

$$k^2 \lambda_{\min} \leq \frac{k^2}{|\mathcal{A}|} \sum_{j=1}^{|\mathcal{A}|} \lambda_j$$

Since the total variance is equal to the sum of all eigenvalues([reference](#)), derive

$$= \frac{k^2}{|\mathcal{A}|} \text{tr}(\text{Var}(\nabla_a Q_{\phi_j}(s, a)))$$

**Lemma 1.** *The total variance of the matrix  $\text{Var}(\nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a}))$  is equal to  $1 - \|\bar{q}\|_2^2$ , where  $\bar{q} = \frac{1}{N} \sum_{j=1}^N \nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a})$ .*

With Lemma 1, we can derive

$$= \frac{k^2}{|\mathcal{A}|} (1 - \|\bar{q}\|_2^2) = \frac{k^2}{|\mathcal{A}|} (1 - \langle \frac{1}{N} \sum_{i=1}^N q_i, \frac{1}{N} \sum_{j=1}^N q_j \rangle)$$

Let  $\min_{i \neq j} \langle \nabla_a Q_{\phi_j}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle = 1 - \epsilon$ . With proposition 1, we can derive

$$\leq \frac{1}{|\mathcal{A}|} \frac{N-1}{N} k^2 \epsilon = \frac{1}{|\mathcal{A}|} \frac{N-1}{N} k^2 (1 - \min_{i \neq j} \langle \nabla_a Q_{\phi_j}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle)$$

If  $\text{Var}(Q_{\phi_j}(s, a + kw_2))$  is small, according to the approximation of the minimum Q-network ensemble is  $\mathbb{E} [\min_{j=1,\dots,N} Q_j(s, a)] \approx m(s, a) - \Phi^{-1} \left( \frac{N - \frac{\pi}{8}}{N - \frac{\pi}{4} + 1} \right) \sigma(s, a)$ , the action  $a + kw_2$  is not sufficiently penalized.

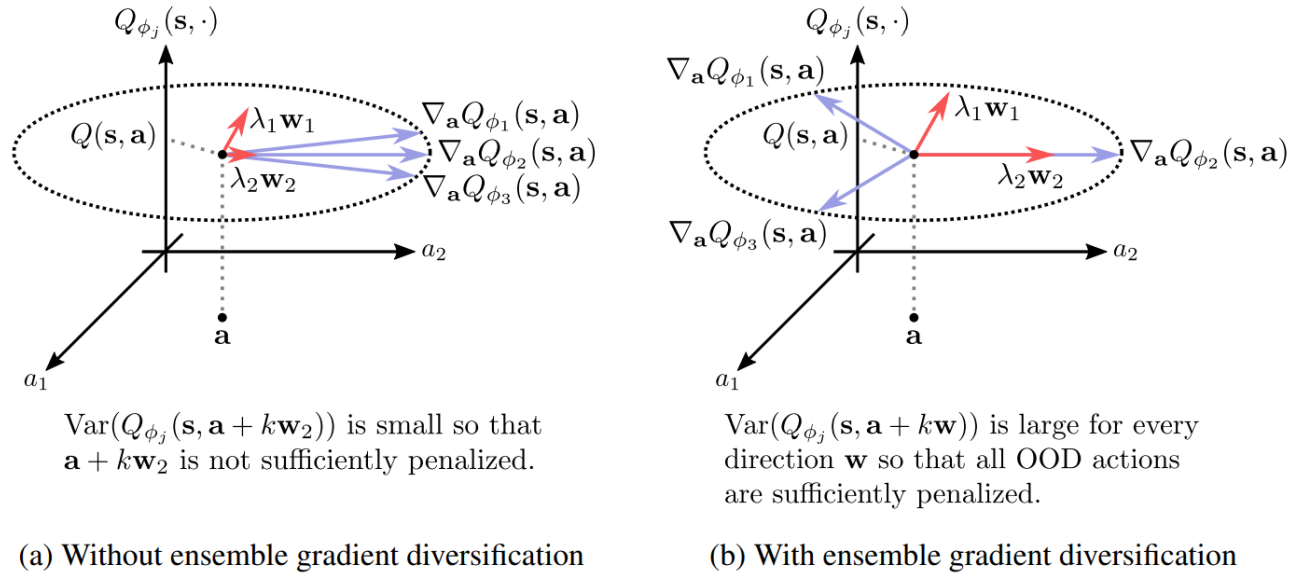


Figure 3: Illustration of the ensemble gradient diversification. The vector  $\lambda_i \mathbf{w}_i$  represents the normalized eigenvector  $\mathbf{w}_i$  of  $\text{Var}(\nabla_{\mathbf{a}} Q_{\phi_j}(s, \mathbf{a}))$  multiplied by its eigenvalue  $\lambda_i$ .

As a result, we now connect the inner product of the gradients of the Q-network(alignment) and the variance of the OOD action  $\text{Var}(Q_{\phi_j}(s, a + kw_2))$ .

Thus, we aim to enlarge the penalty of OOD action. As a result, we aim to diversify the gradients of the Q-network ensemble  $\nabla_a Q_{\phi_i}(s, a)$

$$\min_{\phi} J_{ES}(Q_{\phi}) := \mathbb{E}_{s,a \sim \mathcal{D}} \left[ \left\langle \frac{1}{N} \sum_{i=1}^N \nabla_a Q_{\phi_i}(s, a), \frac{1}{N} \sum_{j=1}^N \nabla_a Q_{\phi_j}(s, a) \right\rangle \right]$$

Then, we can reformulate the equation

$$\min_{\phi} J_{ES}(Q_{\phi}) := \mathbb{E}_{s,a \sim \mathcal{D}} \left[ \frac{1}{N-1} \sum_{1 \leq i \neq j \leq N} \langle \nabla_a Q_{\phi_i}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle \right]$$



# Algorithm

---

**Algorithm 1** Ensemble-Diversified Actor Critic (EDAC)

---

- 1: Initialize policy parameters  $\theta$ , Q-function parameters  $\{\phi_j\}_{j=1}^N$ , target Q-function parameters  $\{\phi'_j\}_{j=1}^N$ , and offline data replay buffer  $\mathcal{D}$
- 2: **repeat**
- 3:     Sample a mini-batch  $B = \{(\mathbf{s}, \mathbf{a}, r, \mathbf{s}')\}$  from  $\mathcal{D}$
- 4:     Compute target Q-values (shared by all Q-functions):

$$y(r, \mathbf{s}') = r + \gamma \left( \min_{j=1, \dots, N} Q_{\phi'_j}(\mathbf{s}', \mathbf{a}') - \beta \log \pi_{\theta}(\mathbf{a}' | \mathbf{s}') \right), \quad \mathbf{a}' \sim \pi_{\theta}(\cdot | \mathbf{s}')$$

- 5:     Update each Q-function  $Q_{\phi_i}$  with gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(\mathbf{s}, \mathbf{a}, r, \mathbf{s}') \in B} \left( \left( Q_{\phi_i}(\mathbf{s}, \mathbf{a}) - y(r, \mathbf{s}') \right)^2 + \frac{\eta}{N-1} \sum_{1 \leq i \neq j \leq N} \text{ES}_{\phi_i, \phi_j}(\mathbf{s}, \mathbf{a}) \right)$$

- 6:     Update policy with gradient ascent using

$$\nabla_{\theta} \frac{1}{|B|} \sum_{\mathbf{s} \in B} \left( \min_{j=1, \dots, N} Q_{\phi_j}(\mathbf{s}, \tilde{\mathbf{a}}_{\theta}(\mathbf{s})) - \beta \log \pi_{\theta}(\tilde{\mathbf{a}}_{\theta}(\mathbf{s}) | \mathbf{s}) \right),$$

where  $\tilde{\mathbf{a}}_{\theta}(\mathbf{s})$  is a sample from  $\pi_{\theta}(\cdot | \mathbf{s})$  which is differentiable w.r.t.  $\theta$  via the reparametrization trick.

- 7:     Update target networks with  $\phi'_i \leftarrow \rho \phi'_i + (1 - \rho) \phi_i$
-

**Lemma 1.** *The total variance of the matrix  $\text{Var}(\nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a}))$  is equal to  $1 - \|\bar{q}\|_2^2$ , where  $\bar{q} = \frac{1}{N} \sum_{j=1}^N \nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a})$ .*

**Proposition 1.** *Suppose  $Q_{\phi_j}(\mathbf{s}, \mathbf{a}) = Q(\mathbf{s}, \mathbf{a})$  and  $Q_{\phi_j}(\mathbf{s}, \cdot)$  is locally linear in the neighborhood of  $\mathbf{a}$  for all  $j \in [N]$ . Let  $\lambda_{\min}$  and  $\mathbf{w}_{\min}$  be the smallest eigenvalue and the corresponding normalized eigenvector of the matrix  $\text{Var}(\nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a}))$  and  $\epsilon > 0$  be the value such that  $\min_{i \neq j} \langle \nabla_{\mathbf{a}} Q_{\phi_i}(\mathbf{s}, \mathbf{a}), \nabla_{\mathbf{a}} Q_{\phi_j}(\mathbf{s}, \mathbf{a}) \rangle = 1 - \epsilon$ . Then, the variance of the  $Q$ -values for an OOD action in the neighborhood along the direction of  $\mathbf{w}_{\min}$  is upper-bounded as follows:*

$$\text{Var}(Q_{\phi_j}(\mathbf{s}, \mathbf{a} + k\mathbf{w}_{\min})) \leq \frac{1}{|\mathcal{A}|} \frac{N-1}{N} k^2 \epsilon,$$

where  $|\mathcal{A}|$  is the action space dimension.

# Proof Sketch

Then, since the smallest eigenvalue is hard to compute, we compute the sum of the all eigenvalues of the covariance matrix  $\text{Var}(\nabla_a Q_{\phi_j}(s, a))$

$$\lambda_{\min} \leq \frac{1}{|\mathcal{A}|} \sum_{j=1}^{|\mathcal{A}|} \lambda_j$$

Since the total variance is equal to the sum of all eigenvalues, we can derive

$$= \frac{1}{|\mathcal{A}|} \text{tr}(\text{Var}(\nabla_a Q_{\phi_j}(s, a)))$$

Let  $q_j = \nabla_a Q_{\phi_j}(s, a)$  be the normalized gradients of Q-network, and  $\bar{q} = \frac{1}{N} \sum_j$

$$\begin{aligned} &= \frac{1}{|\mathcal{A}|} \text{tr}\left(\frac{1}{N} \sum_j (q_j - \bar{q})(q_j - \bar{q})^\top\right) \\ &= \frac{1}{|\mathcal{A}|} \frac{1}{N} \sum_j \text{tr}((q_j - \bar{q})(q_j - \bar{q})^\top) \end{aligned}$$

$$\begin{aligned}
&= \frac{1}{|\mathcal{A}|} \frac{1}{N} \sum_j \text{tr}((q_j - \bar{q})^\top (q_j - \bar{q})) \\
&= \frac{1}{|\mathcal{A}|} \frac{1}{N} \sum_j \text{tr}((q_j - \bar{q})^\top (q_j - \bar{q})) \\
&= \frac{1}{|\mathcal{A}|N} \sum_j (q_j^\top q_j + \bar{q}^\top \bar{q} - 2q_j^\top \bar{q}) \\
&= \frac{1}{|\mathcal{A}|} (1 + \bar{q}^\top \bar{q} - 2(\frac{1}{N} \sum_j q_j^\top) \bar{q}) \\
&= \frac{1}{|\mathcal{A}|} (1 - \|\bar{q}\|_2^2) \\
&= \frac{1}{|\mathcal{A}|} (1 - \langle \frac{1}{N} \sum_{i=1}^N q_i, \frac{1}{N} \sum_{j=1}^N q_j \rangle)
\end{aligned}$$

$$\begin{aligned}
&= \frac{1}{|\mathcal{A}|} \left( 1 - \left( \frac{1}{N^2} \left( \sum_{j=1}^N \langle q_j, q_j \rangle + \sum_{1 \leq i \neq j \leq N} \langle q_i, q_j \rangle \right) \right) \right) \\
&= \frac{1}{|\mathcal{A}|} \left( 1 - \left( \frac{1}{N^2} \left( \sum_{j=1}^N \langle q_j, q_j \rangle + \sum_{1 \leq i \neq j \leq N} \langle q_i, q_j \rangle \right) \right) \right)
\end{aligned}$$

Let  $\min_{i \neq j} \langle \nabla_a Q_{\phi_j}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle = 1 - \epsilon$

$$\begin{aligned}
&\leq \frac{1}{|\mathcal{A}|} (1 - (N + N(N - 1)(1 - \epsilon))) \\
&= \frac{1}{|\mathcal{A}|} \frac{N - 1}{N} \epsilon
\end{aligned}$$

Thus,

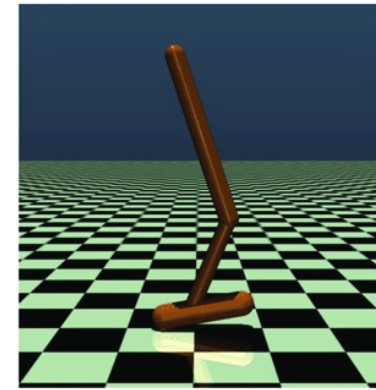
$$\begin{aligned}
& \text{Var}(\nabla_a Q_{\phi_j}(s, a + kw_{\min})) \\
&= k^2 w_{\min}^\top \text{Var}(\nabla_a Q_{\phi_j}(s, a)) w_{\min} \\
&= k^2 \lambda_{\min} \leq \frac{1}{|\mathcal{A}|} (1 - \langle \frac{1}{N} \sum_{i=1}^N q_i, \frac{1}{N} \sum_{j=1}^N q_j \rangle) \\
&\leq \frac{1}{|\mathcal{A}|} \frac{N-1}{N} k^2 \epsilon = \frac{1}{|\mathcal{A}|} \frac{N-1}{N} k^2 (1 - \min_{i \neq j} \langle \nabla_a Q_{\phi_j}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle)
\end{aligned}$$

Thus, instead of maximize the smallest eigenvalue  $\max_{\phi} k^2 \lambda_{\min}$ , **minimizing the cosine similarity of the gradients of the Q-networks is cheaper**

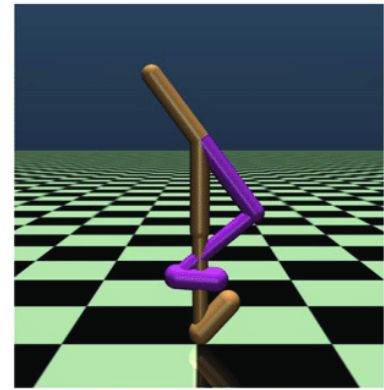
$$\min_{\phi} \mathbb{E}_{s,a \sim \mathcal{D}} \left[ \frac{1}{N-1} \sum_{1 \leq i \neq j \leq N} \langle \nabla_a Q_{\phi_j}(s, a), \nabla_a Q_{\phi_j}(s, a) \rangle \right]$$

# Experiment - D4RL Dataset

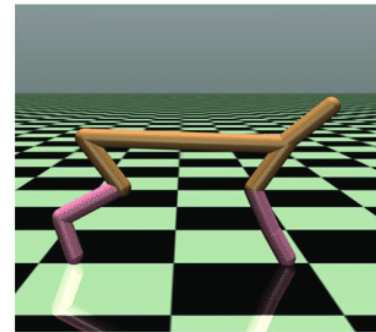
- expert: a fully trained online expert
- medium: a suboptimal policy with approximately 1/3 the performance of the expert
- medium-expert: a mixture of medium and expert policies
- medium-replay: the replay buffer of a policy trained up to the performance of the medium agent
- full-replay: the final replay buffer of the expert policy
- 1M transitions



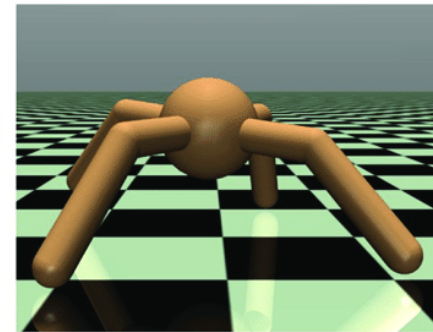
Hopper



Walker2d



Half-Cheetah



Ant



Table 1: Normalized average returns on D4RL Gym tasks, averaged over 4 random seeds. CQL (Paper) denotes the results reported in the original paper.

Task Name	BC	SAC	REM	CQL (Paper)	CQL (Reproduced)	SAC- <i>N</i> (Ours)	EDAC (Ours)
halfcheetah-random	2.2±0.0	29.7±1.4	-0.8±1.1	<b>35.4</b>	31.3±3.5	28.0±0.9	28.4±1.0
halfcheetah-medium	43.2±0.6	55.2±27.8	-0.8±1.3	44.4	46.9±0.4	<b>67.5±1.2</b>	<b>65.9±0.6</b>
halfcheetah-expert	91.8±1.5	-0.8±1.8	4.1±5.7	104.8	97.3±1.1	<b>105.2±2.6</b>	<b>106.8±3.4</b>
halfcheetah-medium-expert	44.0±1.6	28.4±19.4	0.7±3.7	62.4	95.0±1.4	<b>107.1±2.0</b>	<b>106.3±1.9</b>
halfcheetah-medium-replay	37.6±2.1	0.8±1.0	6.6±11.0	46.2	45.3±0.3	<b>63.9±0.8</b>	<b>61.3±1.9</b>
halfcheetah-full-replay	62.9±0.8	<b>86.8±1.0</b>	27.8±35.4	-	76.9±0.9	84.5±1.2	84.6±0.9
hopper-random	3.7±0.6	9.9±1.5	3.4±2.2	10.8	5.3±0.6	<b>31.3±0.0</b>	<b>25.3±10.4</b>
hopper-medium	54.1±3.8	0.8±0.0	0.7±0.0	86.6	61.9±6.4	<b>100.3±0.3</b>	<b>101.6±0.6</b>
hopper-expert	107.7±9.7	0.7±0.0	0.8±0.0	109.9	106.5±9.1	<b>110.3±0.3</b>	<b>110.1±0.1</b>
hopper-medium-expert	53.9±4.7	0.7±0.0	0.8±0.0	<b>111.0</b>	96.9±15.1	110.1±0.3	110.7±0.1
hopper-medium-replay	16.6±4.8	7.4±0.5	27.5±15.2	48.6	86.3±7.3	<b>101.8±0.5</b>	<b>101.0±0.5</b>
hopper-full-replay	19.9±12.9	41.1±17.9	19.7±24.6	-	101.9±0.6	<b>102.9±0.3</b>	<b>105.4±0.7</b>
walker2d-random	1.3±0.1	0.9±0.8	6.9±8.3	7.0	5.4±1.7	<b>21.7±0.0</b>	<b>16.6±7.0</b>
walker2d-medium	70.9±11.0	-0.3±0.2	0.2±0.7	74.5	79.5±3.2	<b>87.9±0.2</b>	<b>92.5±0.8</b>
walker2d-expert	108.7±0.2	0.7±0.3	1.0±2.3	<b>121.6</b>	109.3±0.1	107.4±2.4	115.1±1.9
walker2d-medium-expert	90.1±13.2	1.9±3.9	-0.1±0.0	98.7	109.1±0.2	<b>116.7±0.4</b>	<b>114.7±0.9</b>
walker2d-medium-replay	20.3±9.8	-0.4±0.3	12.5±6.2	32.6	76.8±10.0	<b>78.7±0.7</b>	<b>87.1±2.3</b>
walker2d-full-replay	68.8±17.7	27.9±47.3	-0.2±0.3	-	94.2±1.9	<b>94.6±0.5</b>	<b>99.8±0.7</b>
Average	49.9	16.2	6.2	-	73.7	<b>84.5</b>	<b>85.2</b>

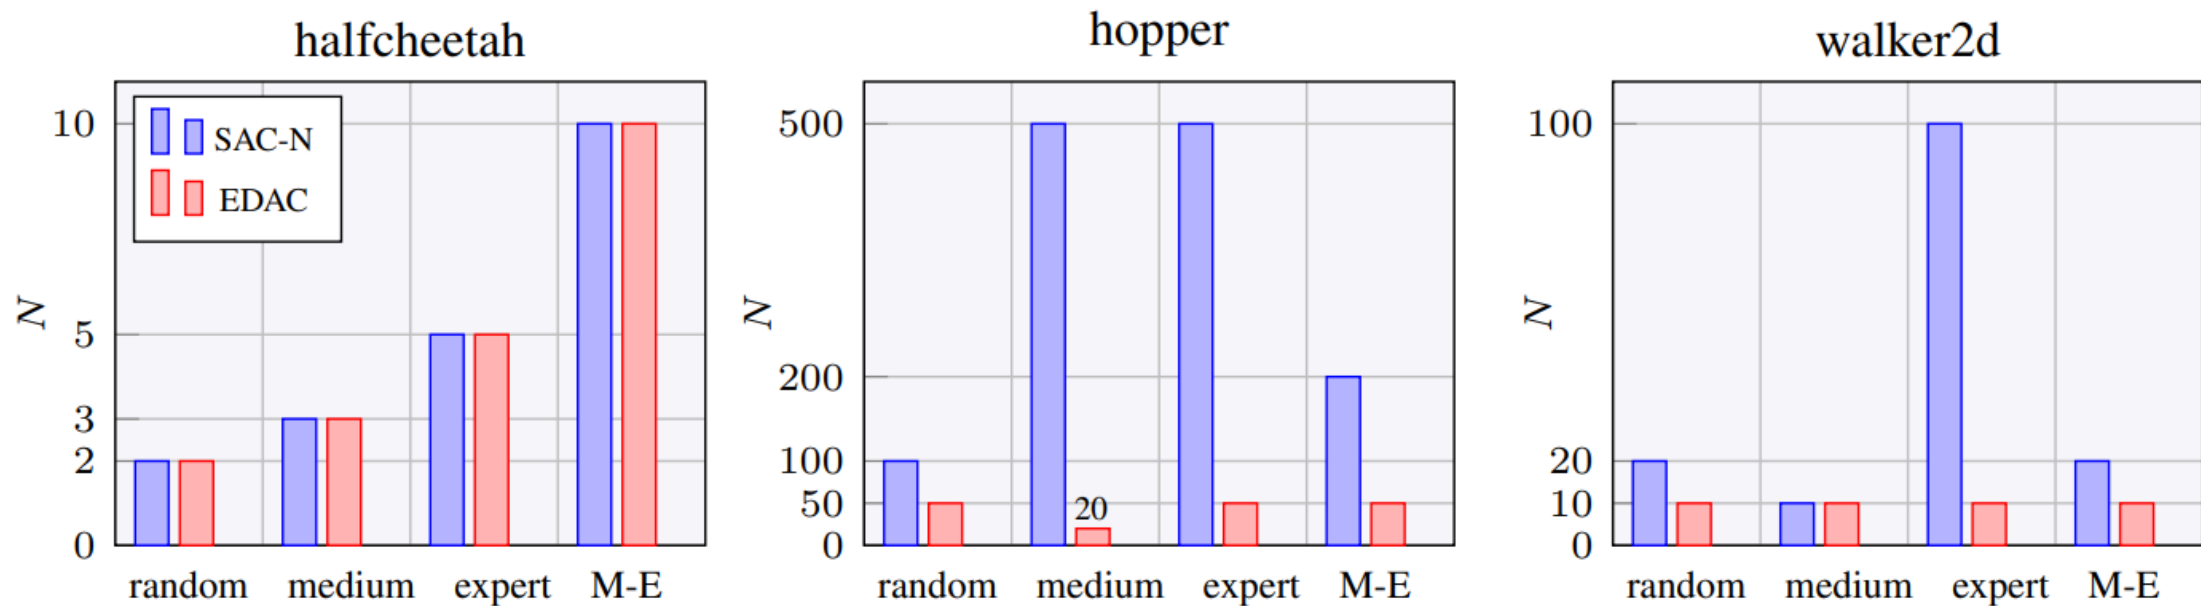


Figure 5: Minimum number of Q-ensembles ( $N$ ) required to achieve the performance reported in Table 1. M-E denotes medium-expert. We omit the results of medium-replay and full-replay as SAC- $N$  already works well with a small number of ensembles (less than or equal to 5). For more details of the experiment, please refer to Appendix C.

Table 3: Computational costs of each method.

	<b>Runtime</b> (s/epoch)	<b>GPU Mem.</b> (GB)
<b>SAC</b>	21.4	1.3
<b>CQL</b>	38.2	1.4
<b>SAC-500</b>	44.1	5.1
<b>EDAC</b>	30.8	1.8

Table 2: Normalized average returns on D4RL Adroit tasks, averaged over 4 random seeds.

Task Name	BC	SAC	REM	CQL (Paper)	CQL (Reproduced)	SAC- $N$ (Ours)	EDAC (Ours)
pen-human	$25.8 \pm 8.8$	$4.3 \pm 3.8$	$5.4 \pm 4.3$	<b>55.8</b>	$35.2 \pm 6.6$	$9.5 \pm 1.1$	$52.1 \pm 8.6$
hammer-human	$3.1 \pm 3.2$	$0.2 \pm 0.0$	$0.3 \pm 0.0$	2.1	$0.6 \pm 0.5$	$0.3 \pm 0.0$	$0.8 \pm 0.4$
door-human	$2.8 \pm 0.7$	$-0.3 \pm 0.0$	$-0.3 \pm 0.0$	9.1	$1.2 \pm 1.8$	$-0.3 \pm 0.0$	<b><math>10.7 \pm 6.8</math></b>
relocate-human	$0.0 \pm 0.0$	$-0.3 \pm 0.0$	$-0.3 \pm 0.0$	0.35	$0.0 \pm 0.0$	$-0.1 \pm 0.1$	$0.1 \pm 0.1$
pen-cloned	$38.3 \pm 11.9$	$-0.8 \pm 3.2$	$-1.0 \pm 0.1$	40.3	$27.2 \pm 11.3$	<b><math>64.1 \pm 8.7</math></b>	<b><math>68.2 \pm 7.3</math></b>
hammer-cloned	$0.7 \pm 0.3$	$0.1 \pm 0.1$	$-0.3 \pm 0.0$	5.7	$1.4 \pm 2.1$	$0.2 \pm 0.2$	$0.3 \pm 0.0$
door-cloned	$0.0 \pm 0.0$	$-0.3 \pm 0.1$	$-0.3 \pm 0.0$	3.5	$2.4 \pm 2.4$	$-0.3 \pm 0.0$	<b><math>9.6 \pm 8.3</math></b>
relocate-cloned	$0.1 \pm 0.0$	$-0.1 \pm 0.1$	$-0.2 \pm 0.2$	-0.1	$0.0 \pm 0.0$	$0.0 \pm 0.0$	$0.0 \pm 0.0$

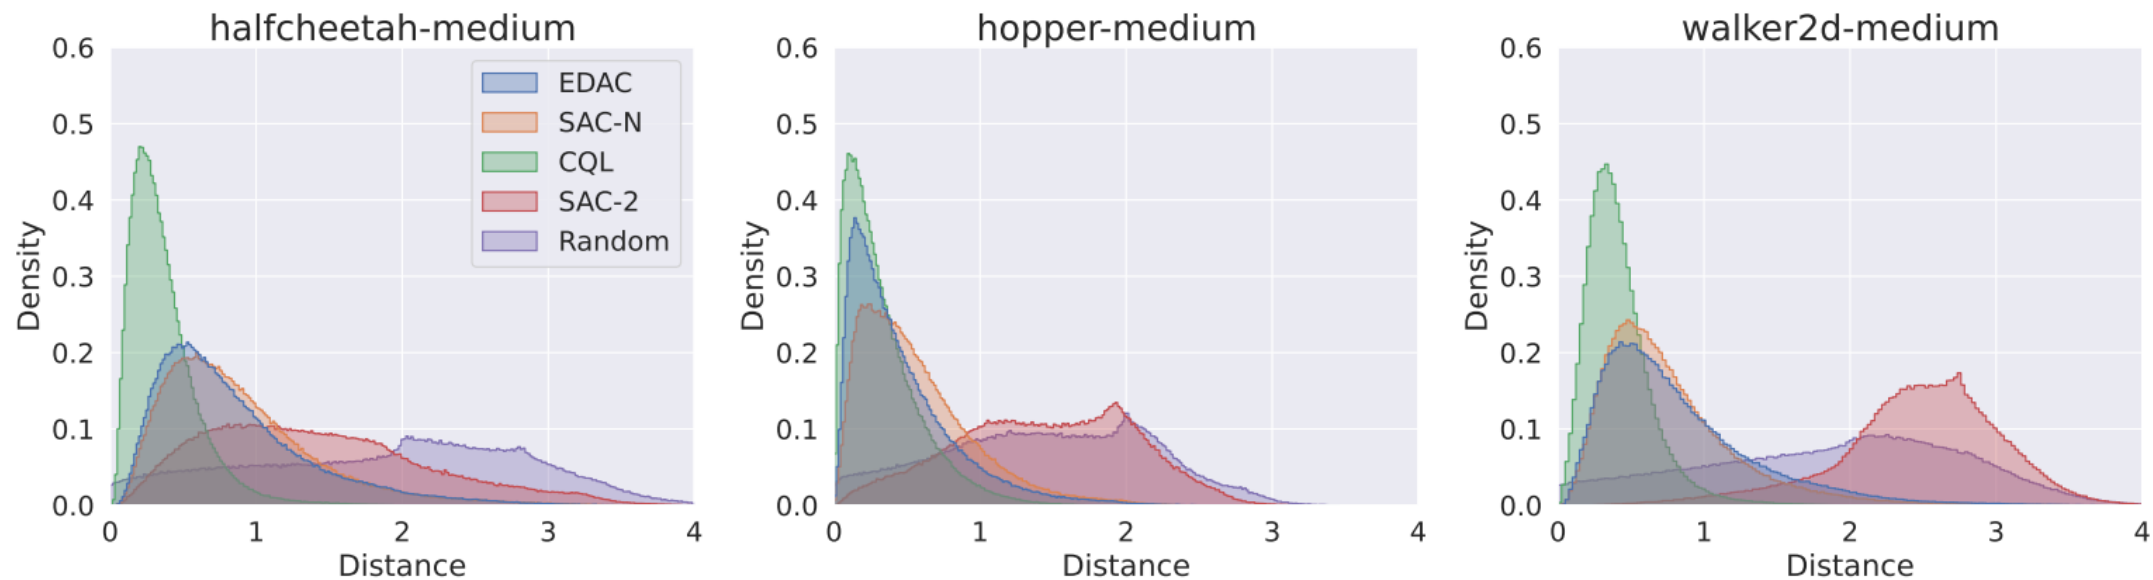


Figure 6: Histograms of the distances between the actions from each methods (EDAC, SAC- $N$ , CQL, SAC-2, and a random policy) and the actions from the dataset. For more details of the experiment, please refer to Appendix C.

# Conclusion

- SAC-N can be efficiently leveraged to construct an uncertainty-based offline RL method that outperforms previous methods on various datasets.
- we proposed Ensemble-Diversifying Actor-Critic (EDAC) that effectively reduces the required number of ensemble networks for quantifying and penalizing the epistemic uncertainty.